Synchronization via Semaphores

General Synchronization Task

- Operations or instructions of concurrent processes (or threads) can interleave.
- Depending on the situation, there are **orders** of operations/instructions that can be considered **VALID** and there are orders that can be considered **INVALID**.
- **Synchronization** is about implementing different mechanisms such that only the **VALID orders** of operation/instructions can occur regardless of how the threads/processes are scheduled by the operating system.

Synchronization Task/Problem 1a: **(Thread) Barrier**

[Problem Description]

Task/Problem 1a: Barrier: a() < Y()

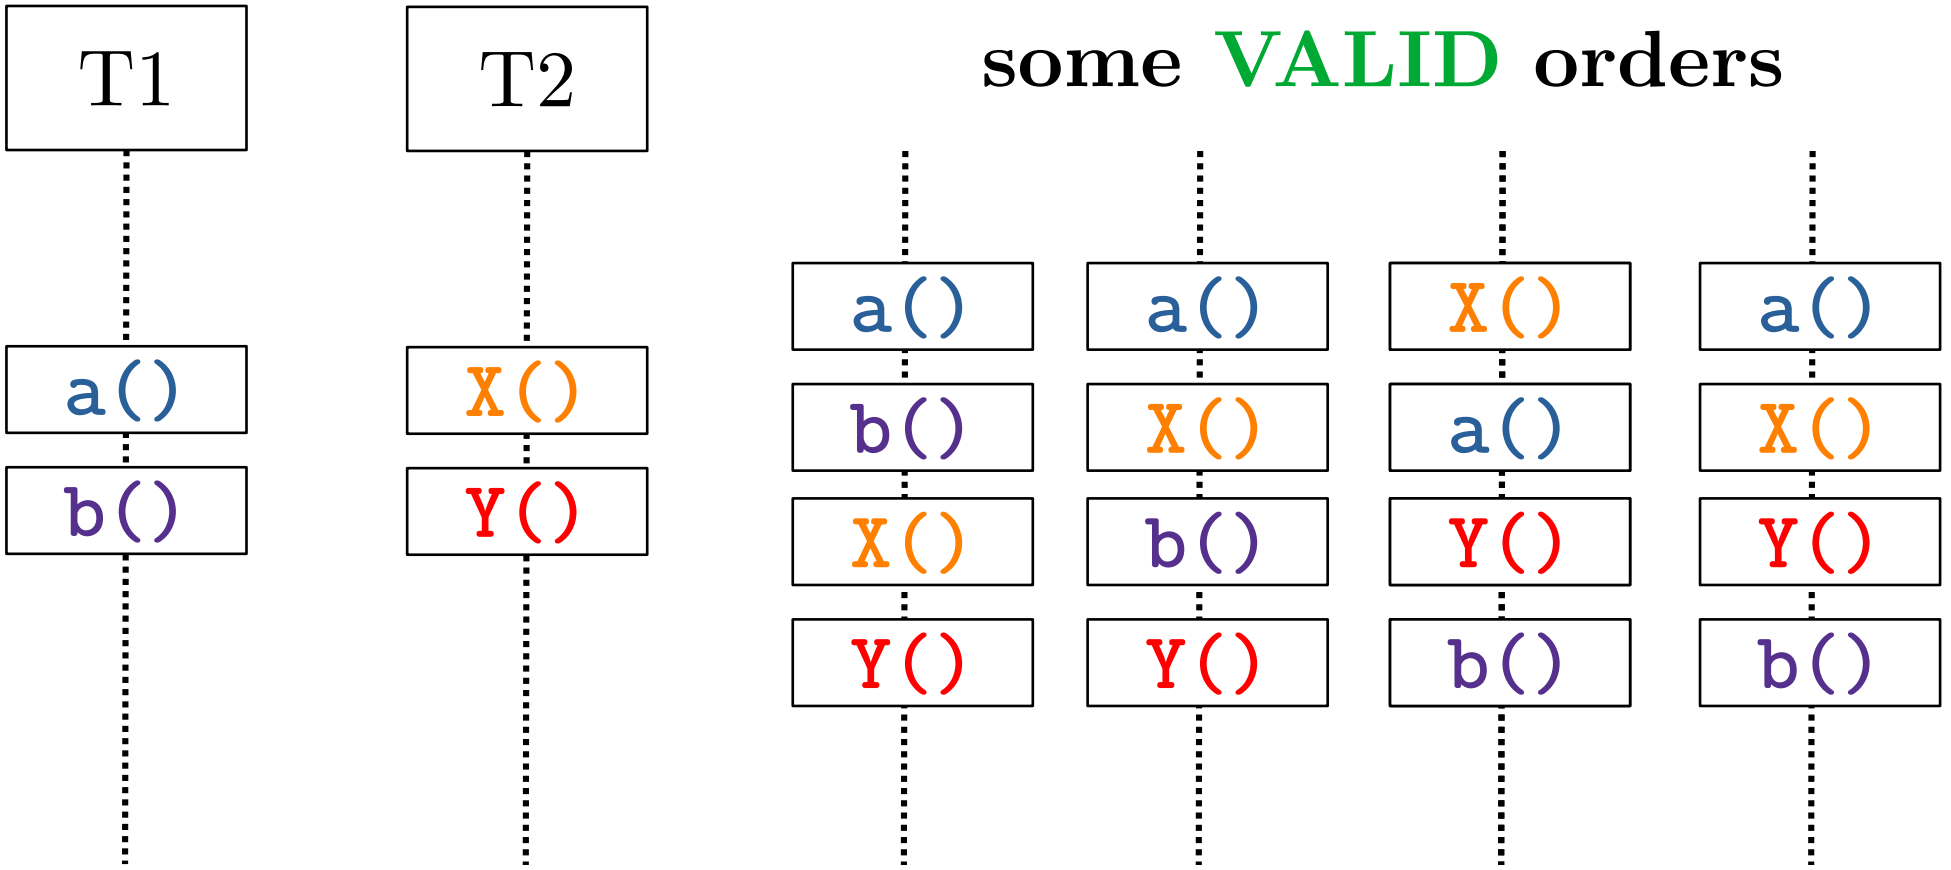

Constraint: We consider only those orders where a() is performed before Y() as VALID.

Task/Problem 1a: Barrier: a() < Y()

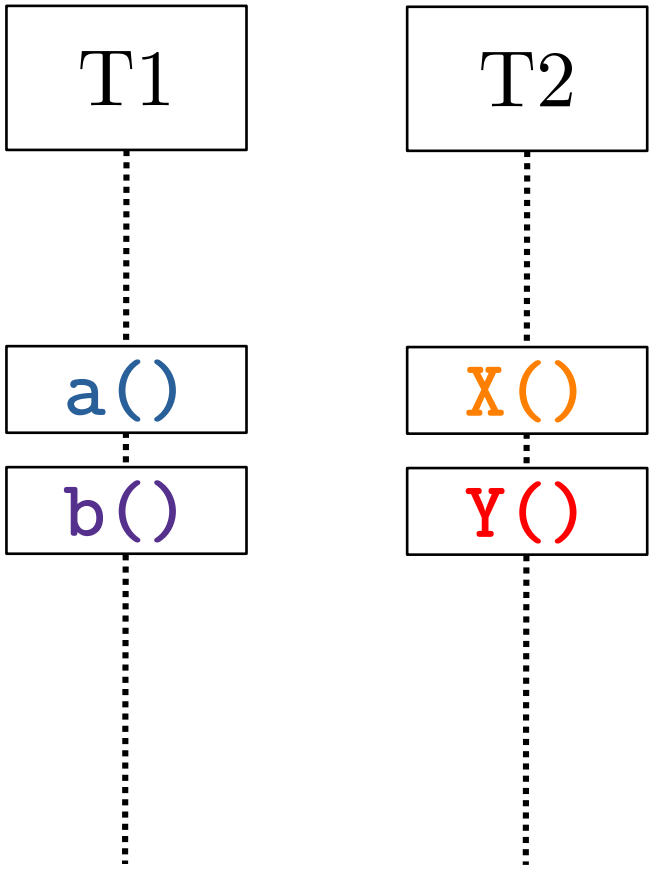

INVALID order

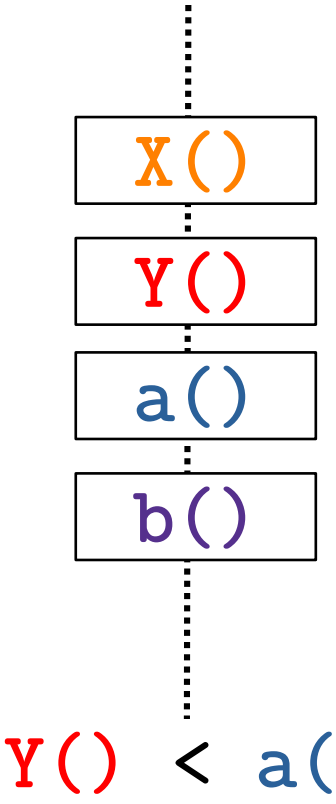

Y() < a()

Constraint: We consider only those orders where a() is performed before d() as VALID.

Task/Problem 1a: Barrier: a() < Y()

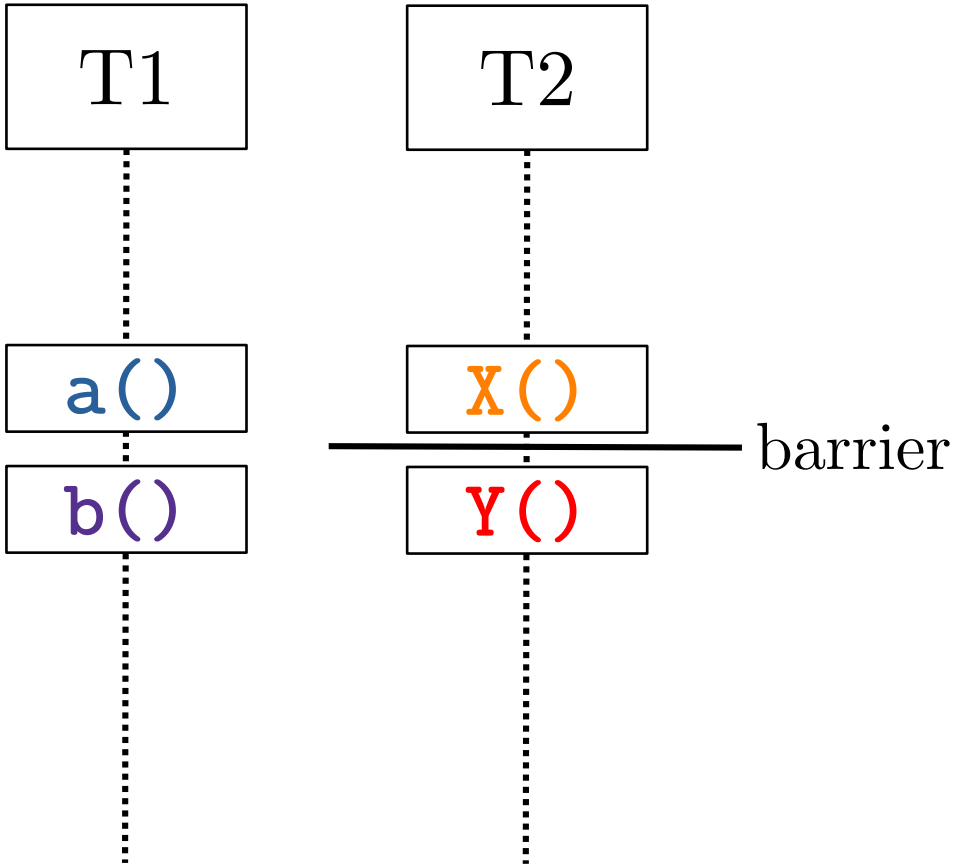

A *barrier* should be placed before the call Y() in T2.

This barrier means T2 has to wait for T1 to complete the a() call before performing Y().

Constraint: We consider only those orders where a() is performed before Y() as VALID.

Synchronization Task/Problem 1b: **(Thread) Barriers**

[Problem Description]

Task/Problem 1b: Barriers: a() < Y(), X() < b()

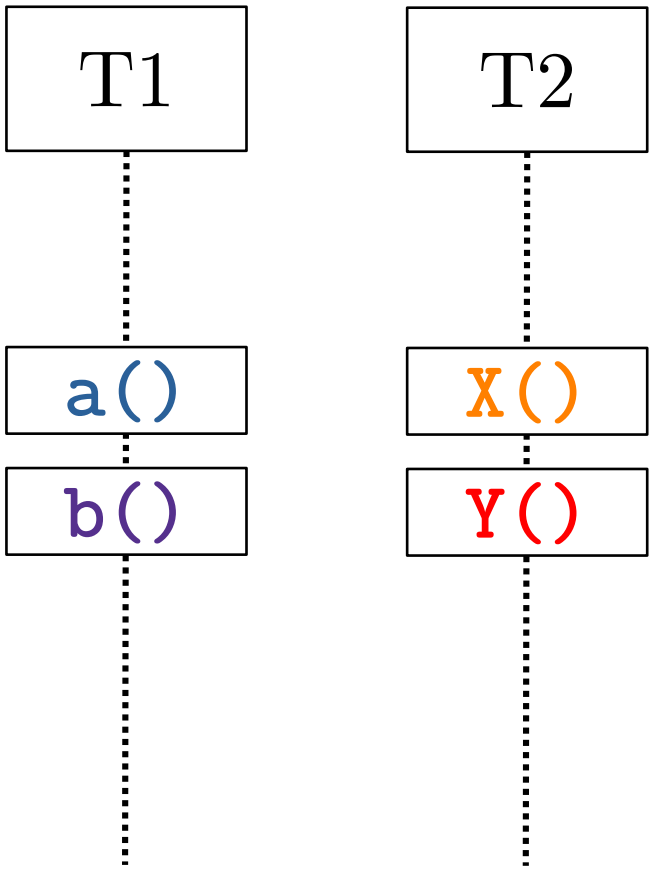

INVALID orders

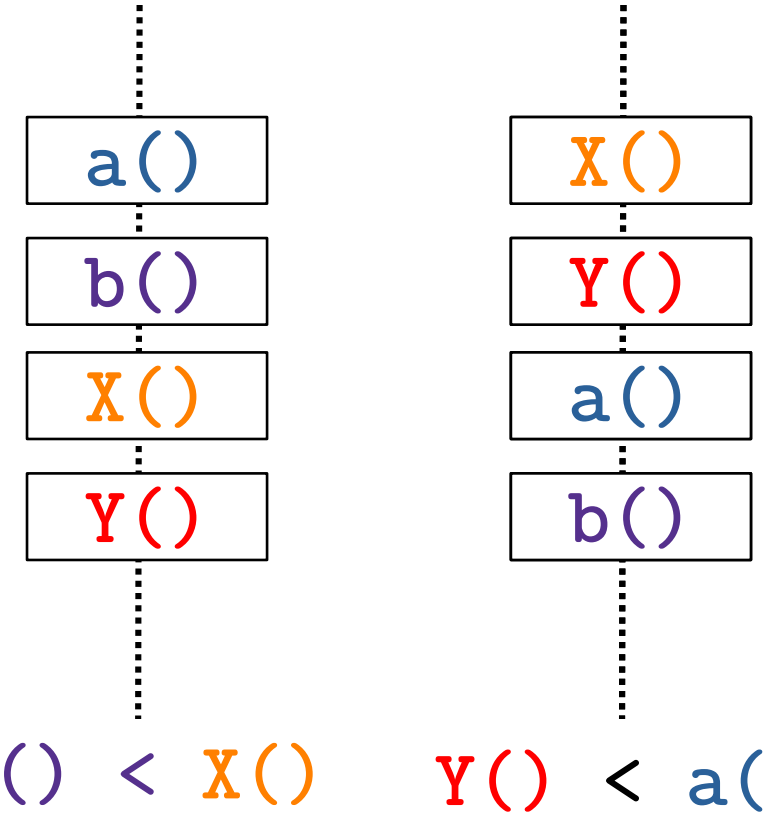

b() < X() Y() < a()

Constraint: We consider only those orders where a() is performed before Y() and X() is performed before b() as **VALID**.

Task/Problem 1b: Barriers: a() < Y(), X() < b()

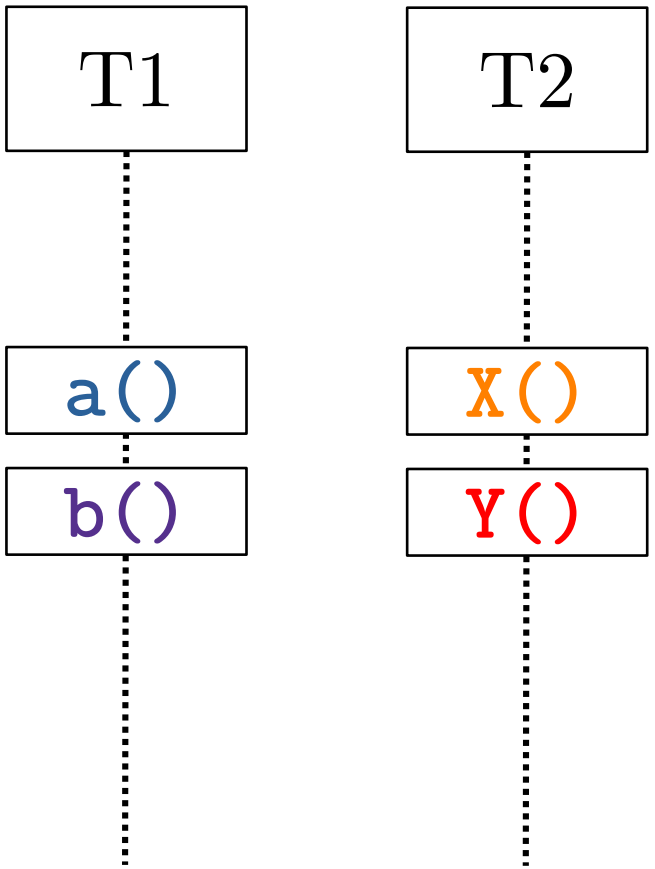

VALID orders

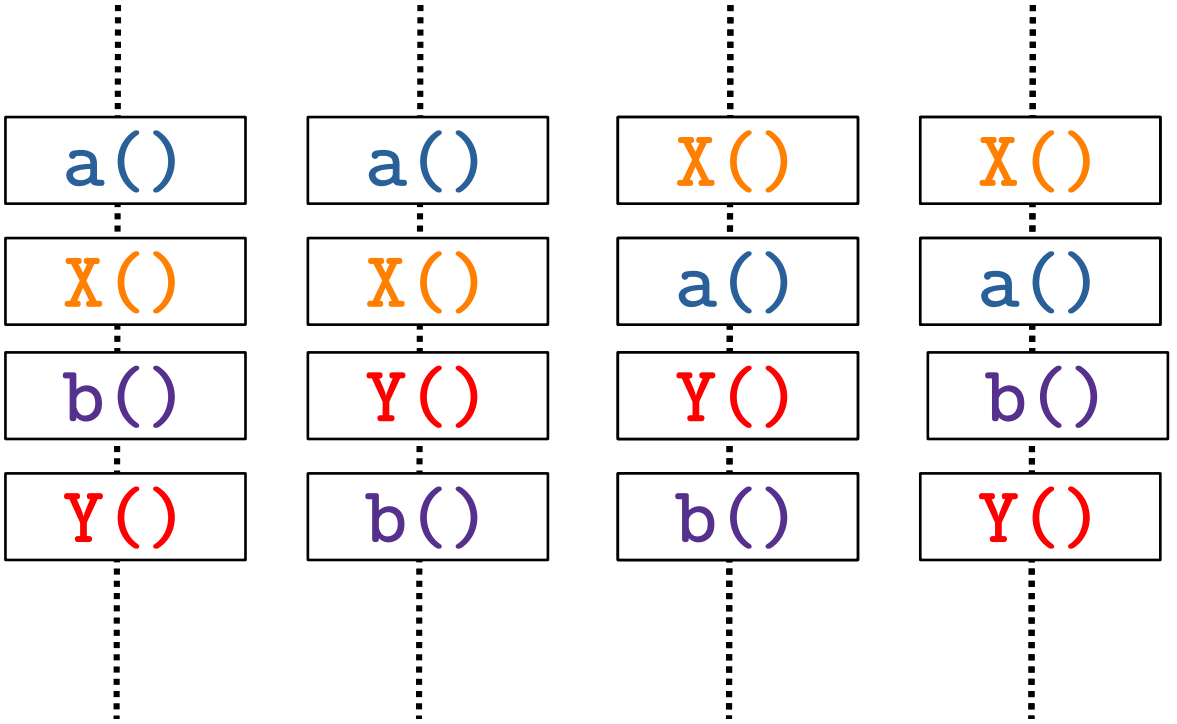

a() < Y(), X() < b()

Constraint: We consider only those orders where a() is performed before Y() and X() is performed before b() as **VALID**.

Task/Problem 1b: Barriers: a() < Y(), X() < b()

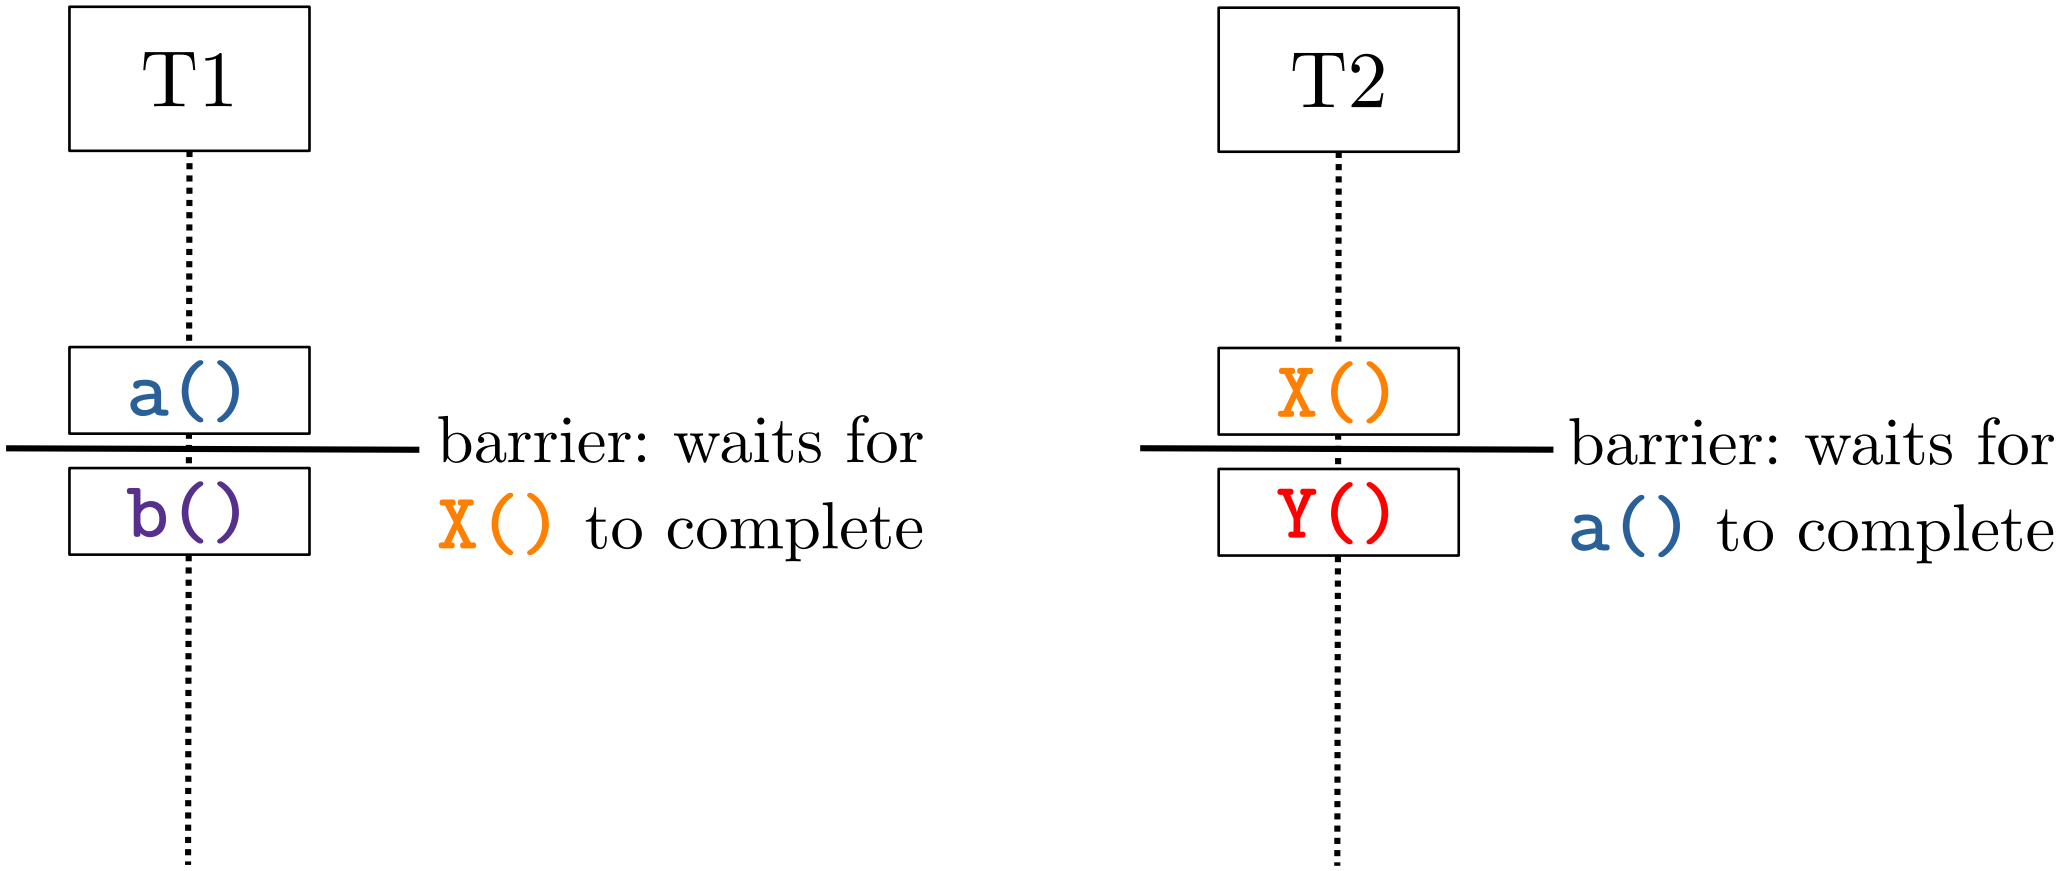

Constraint: We consider only those orders where a() is performed before Y() and X() is performed before b() as **VALID**.

Semaphores

Synchronization Object: Semaphore

Semaphore `S` in an **integer** that signals permission or availability of resources.

`S = n > 0` means there are `n` permits or resources available.

`wait(S)` – a method called by a process/thread in order to request (to **wait** for) a permit or access to a resource. `Decrements S (i.e. S--;)`.

`signal(S)` – a method called by a process/thread in order to return a permit or a resource. `Increments S (i.e. S++;)`.

Synchronization Object: Semaphore

```
wait(S)
{
   while(S <= 0)
   {
      // busy wait
   }
   S--;
}
```

```
signal(S)
{
   S++;
}
```

Semaphores: Other Names for wait and signal

	C	C++
wait(S)	sem_wait(S)	S.acquire()
signal(S)	sem_post(S)	S.release()

Semaphore: Analogy + Some Information

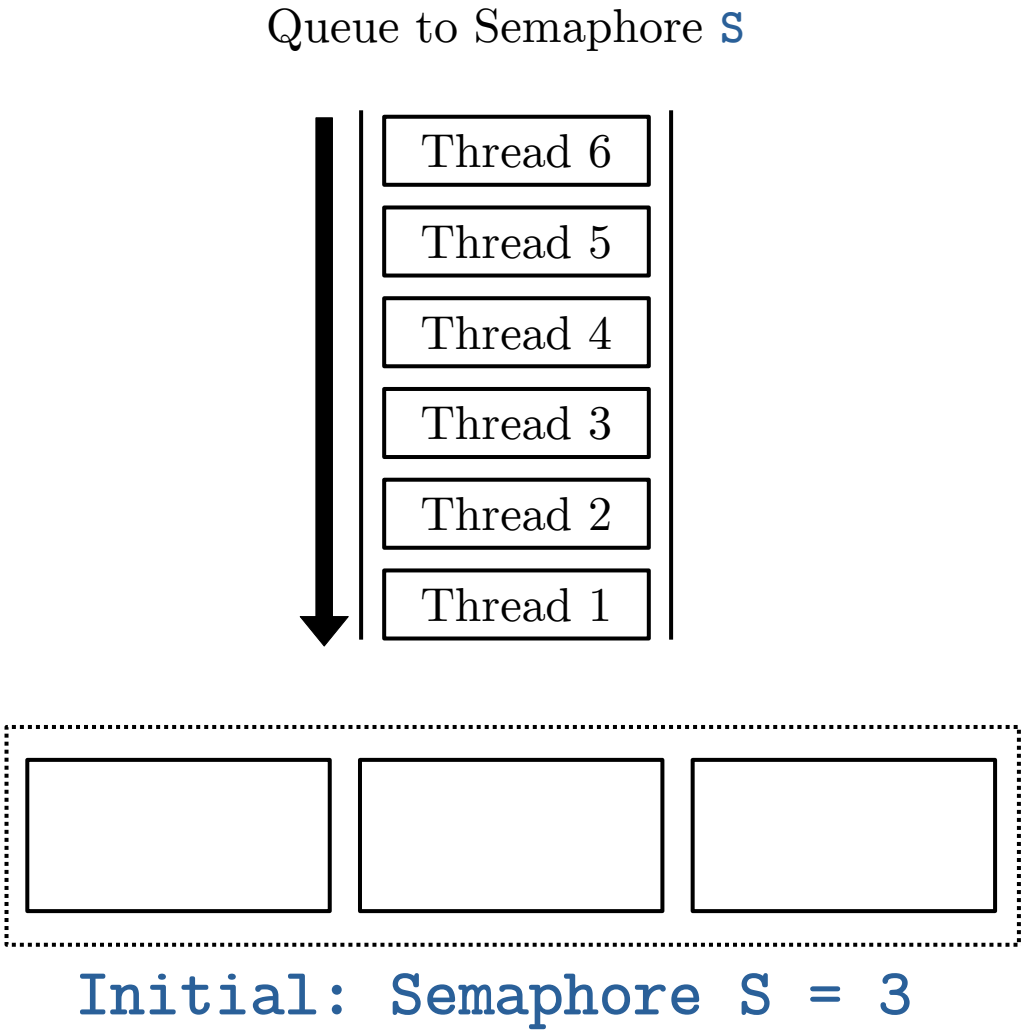

- The boxes represent the permits or available resources.
- The initial number of boxes represent the initial value of the semaphore which is the number of permits or resources available.
- The queue represent all threads that call `wait(S)`. These threads want permits or resources associated with `semaphore S`.
- ***Multiplexing*** is the idea that multiple threads can acquire permits or resources at the same time.

Semaphore: Analogy + Some Information

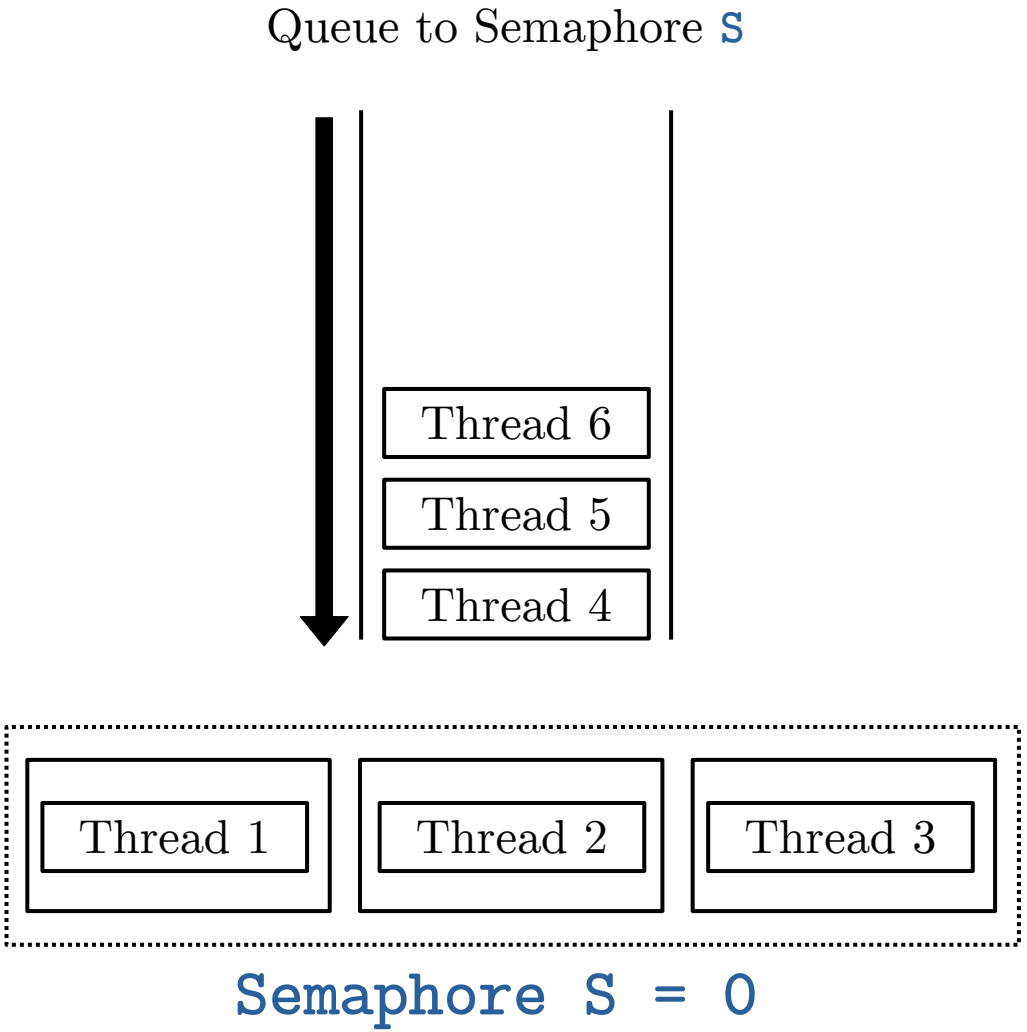

- The boxes represent the permits or available resources.
- The initial number of boxes represent the initial value of the semaphore which is the number of permits or resources available.
- The queue represent all threads that call `wait(S)`. These threads want permits or resources associated with `semaphore S`.
- ***Multiplexing*** is the idea that multiple threads can acquire permits or resources at the same time.

Semaphore: Some Information

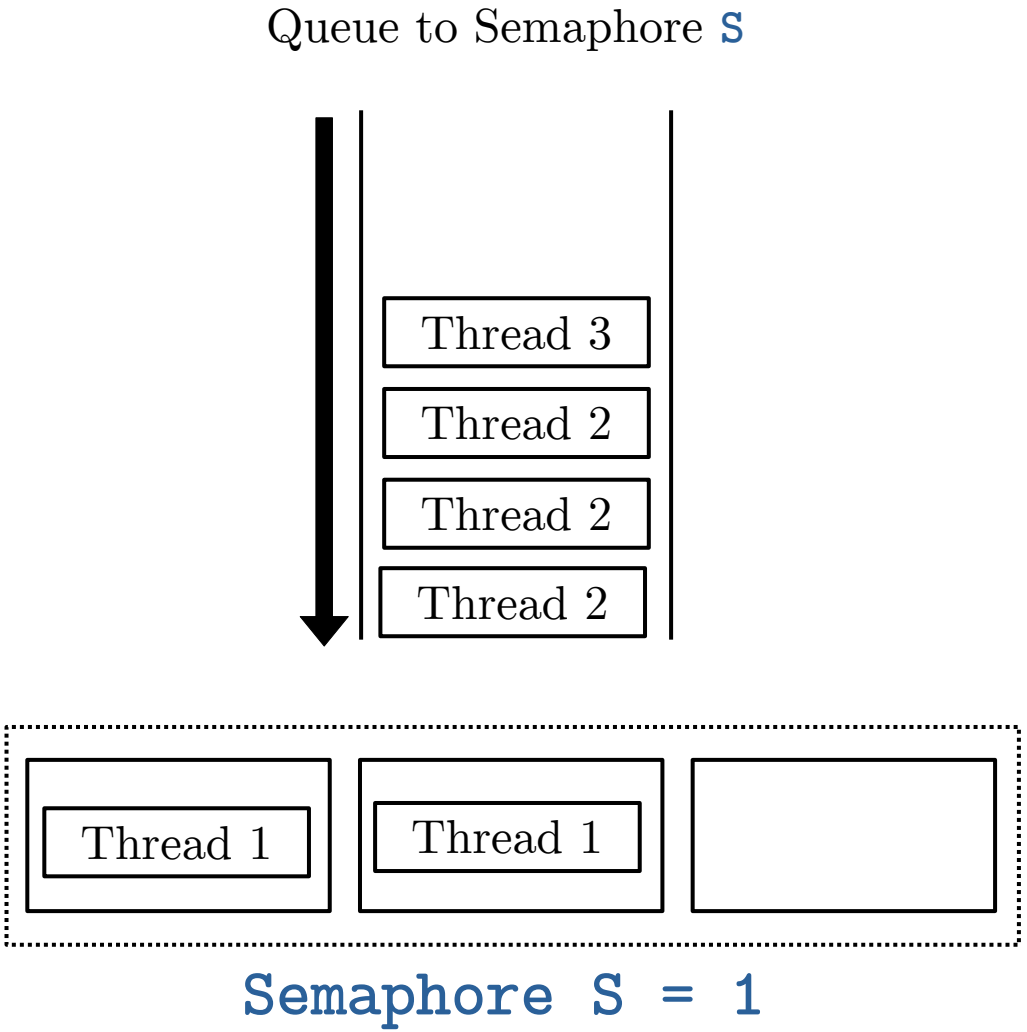

- ***Multiplexing*** is the idea that multiple threads can acquire permits or resources at the same time.
- A semaphore is called ***counting*** if it can give more than 1 permit or resource.
- A semaphore is called ***binary*** if at most it can give 1 permit or 1 resource. S $\leq$ 1.
- A thread calling `wait(S)` multiple times means it wants multiple permits or access to multiple resources associated with semaphore `semaphore S`.

Semaphore: Analogy + Information

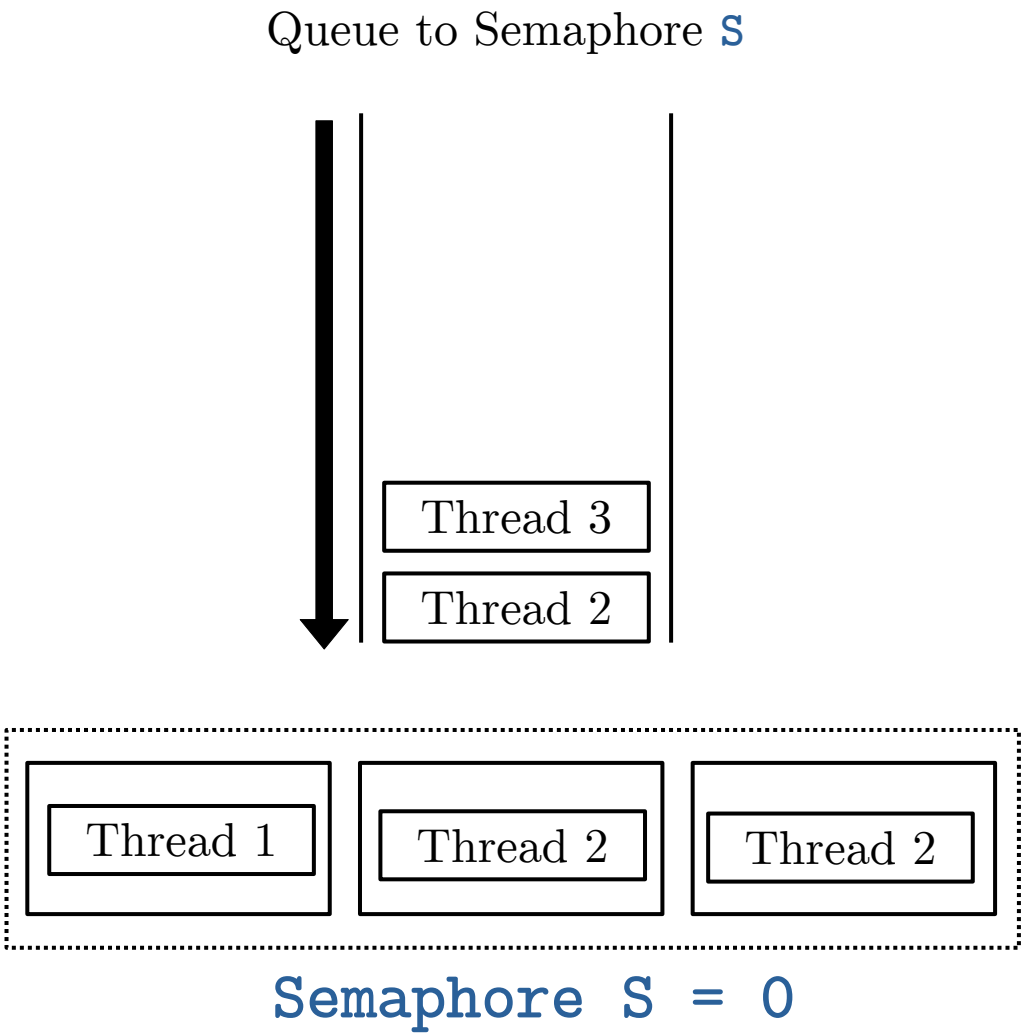

- Threads calling the `signal(S)` increments `semaphore S` and thus increasing the number of permits or resources available that are associate with `semaphore S`.
- In different programming languages, you can specify a semaphore's upper limit value which is the maximum number of permits or resources associated with that semaphore.

Synchronization Task/Problem 1a: **(Thread) Barrier**

[Solution Description]

Solution 1a: Barrier: a() < Y()

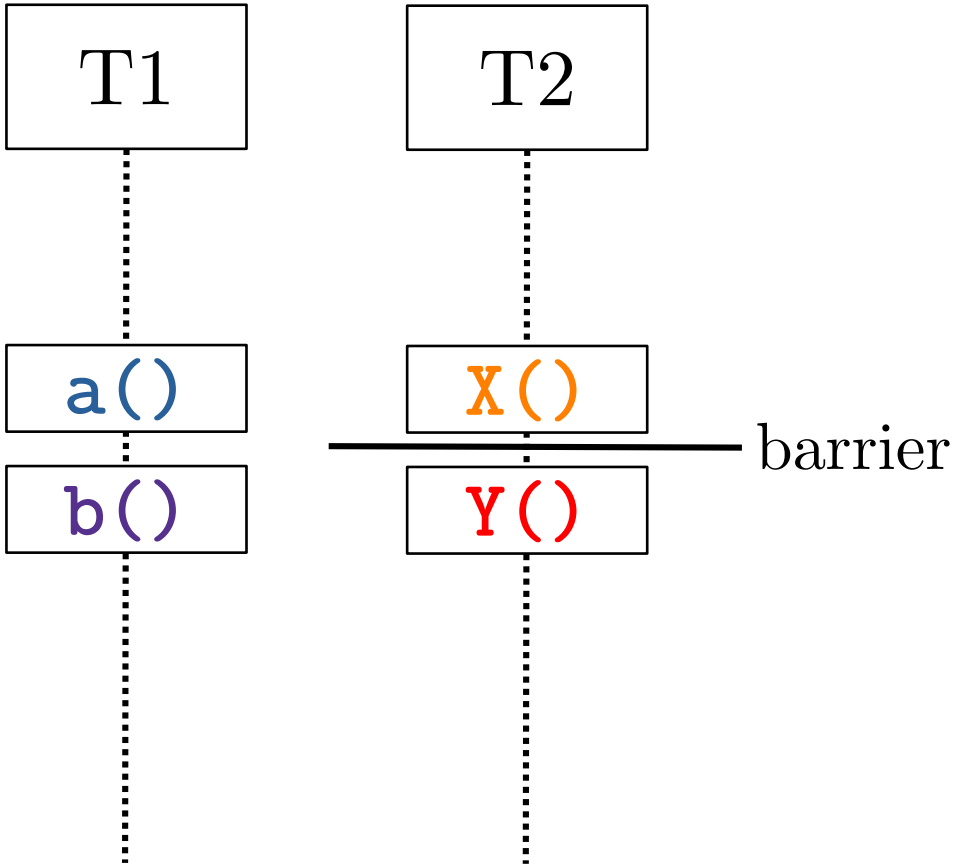

A *barrier* should be placed before the call Y() in T2.

This barrier means T2 has to wait for T1 to complete the a() call before performing Y().

Constraint: We consider only those orders where a() is performed before Y() as VALID.

Solution 1a: Barrier: a() < Y()

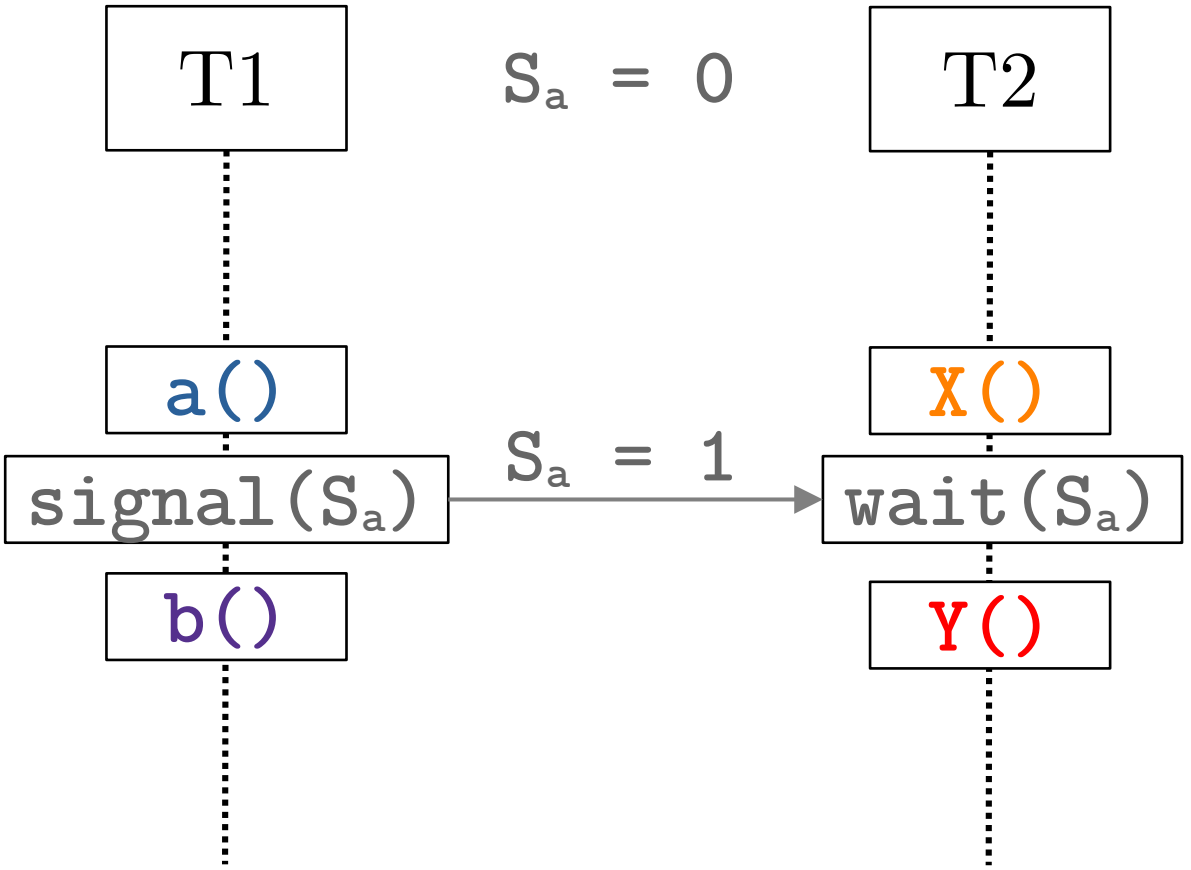

S_a represents the **'Y'** *permit*.
Initially, $S_a = 0$.

T2 can only exit `wait(`S_a`)` and call `Y()` if there is at least *1 permit*.

The *1 permit* is only available when **T1** calls `signal(`S_a`)` after `a()` executes.

Synchronization Task/Problem 1b: **(Thread) Barriers**

[Solution Description]

Solution 1b: Barriers: a() < Y(), X() < b()

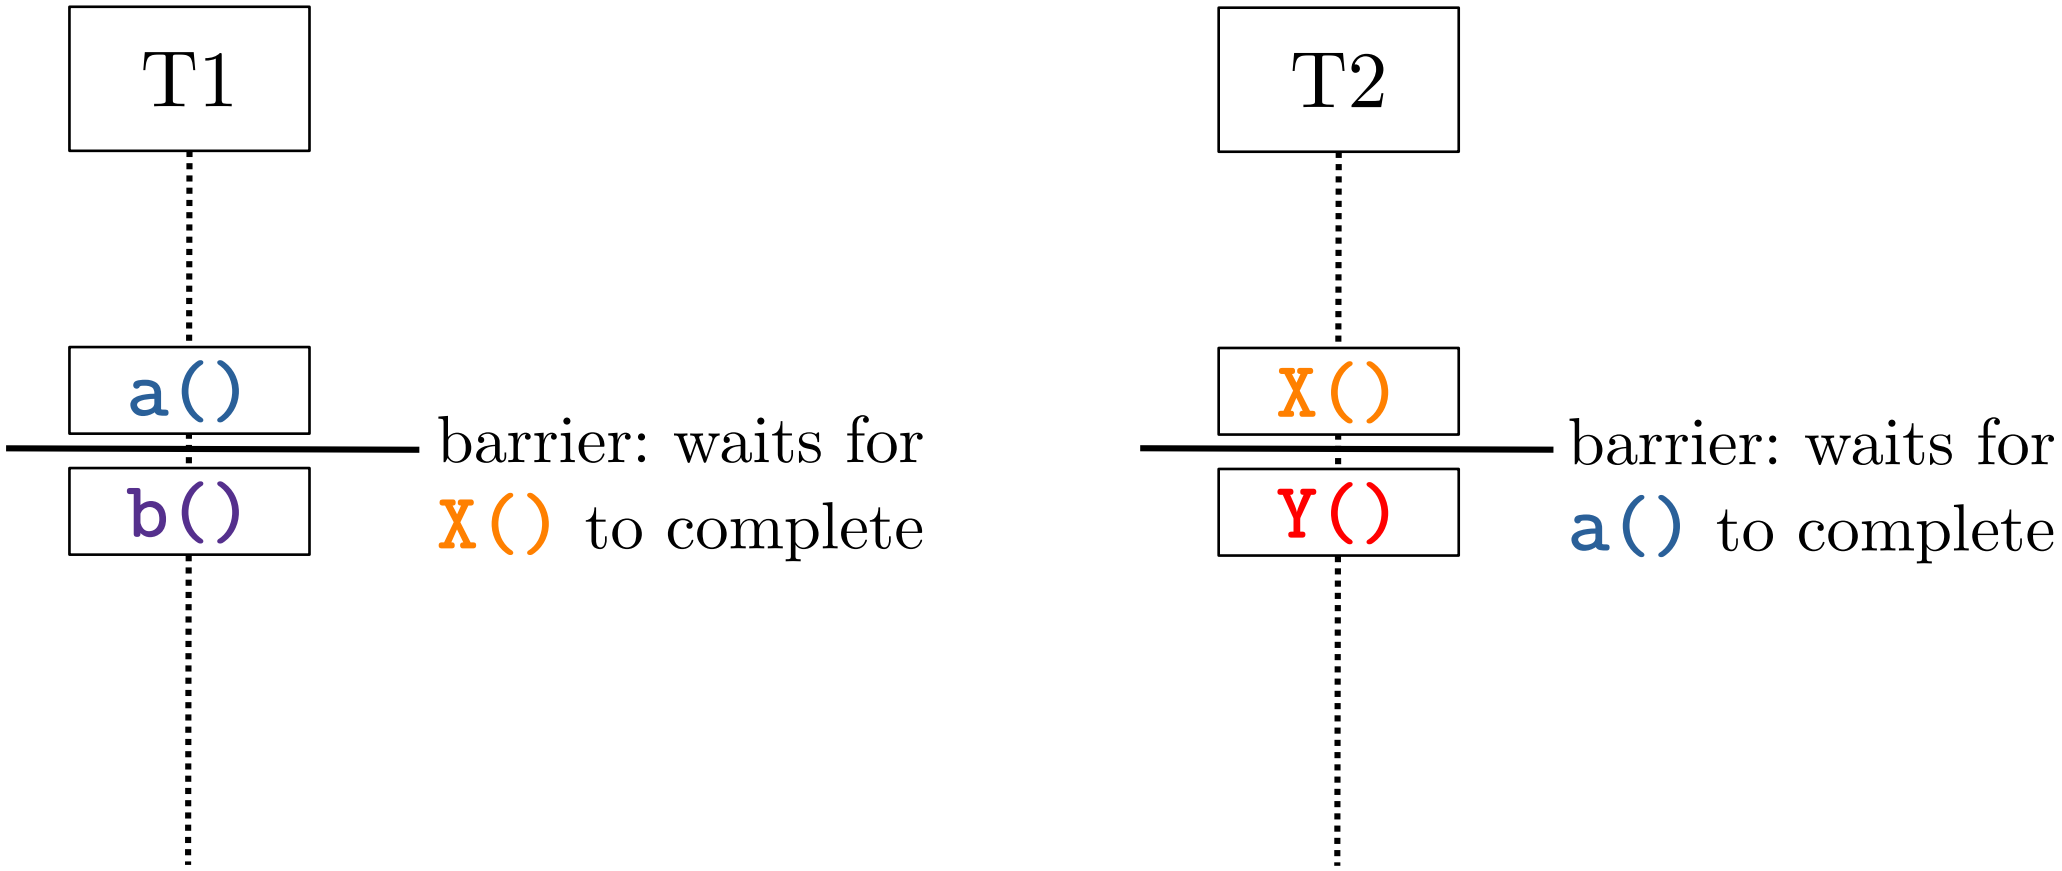

Constraint: We consider only those orders where a() is performed before Y() and X() is performed before b() as **VALID**.

Solution 1b: Barriers: a() < Y(), X() < b()

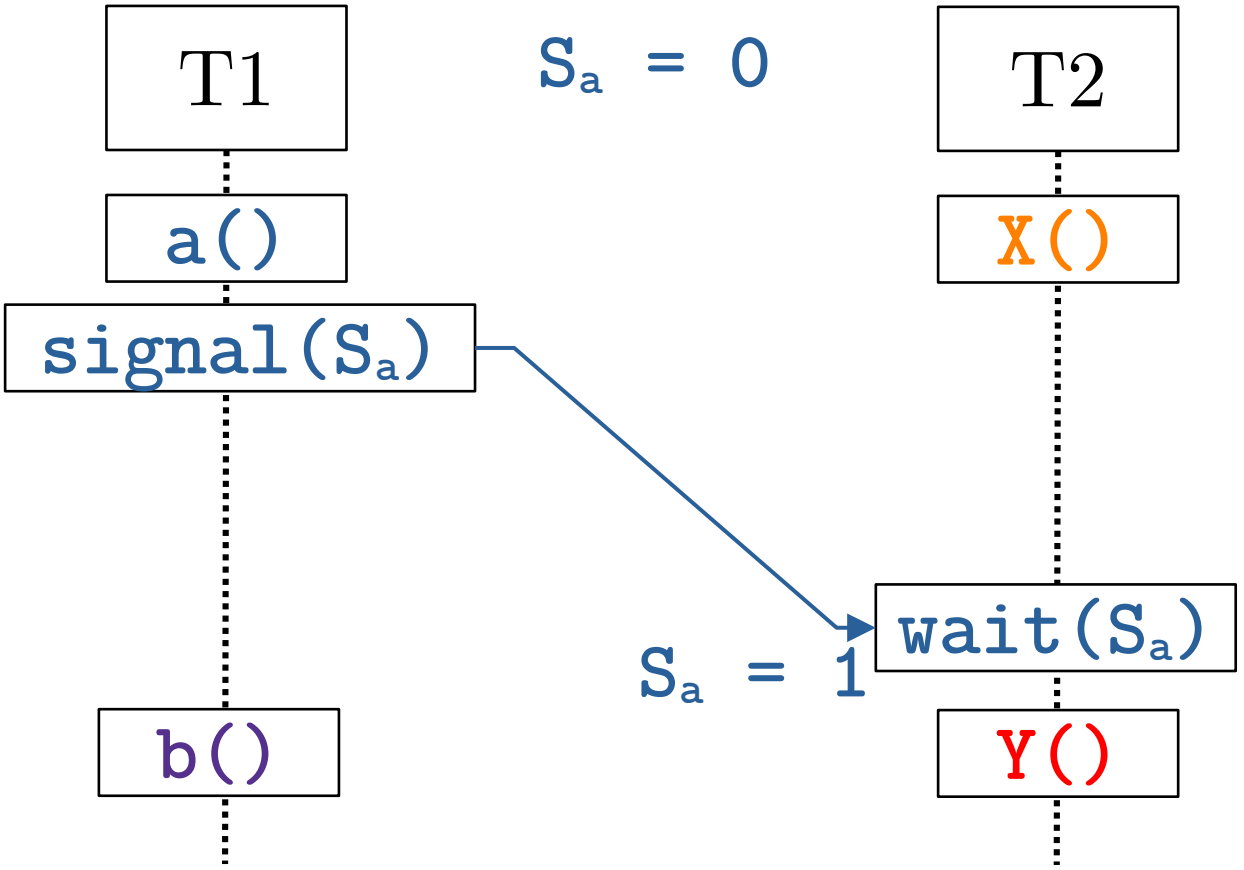

S_a represents the **'Y'** *permit*.
Initially, $S_a = 0$.

T2 can only exit `wait(`S_a`)` and call `Y()` if there is at least *1 permit*.

The *1 permit* is only available when **T1** calls `signal(`S_a`)` after `a()` executes.

Solution 1b: Barriers: a() < Y(), X() < b()

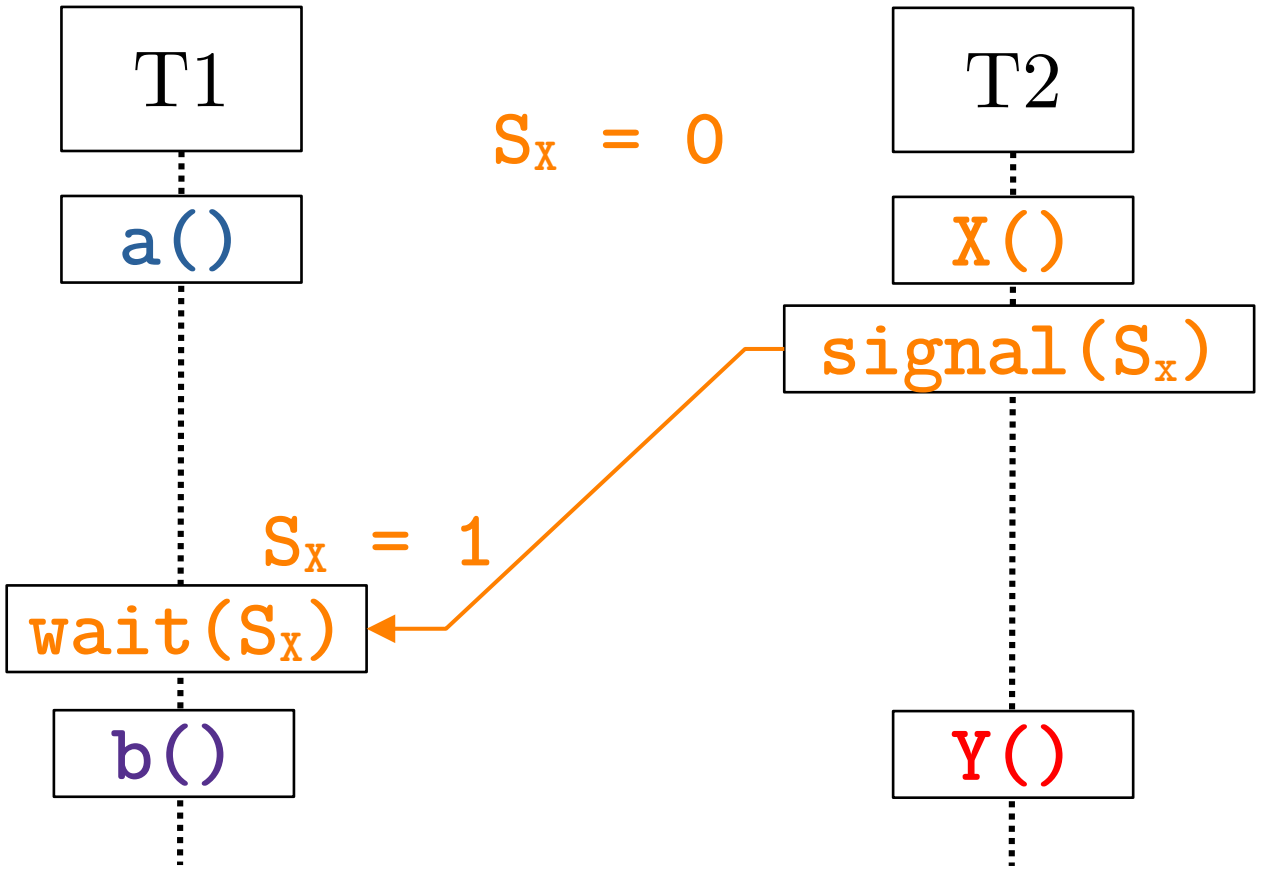

S_X represents the *'b' permit*. Initially, S_X = 0.

T1 can only exit wait(S_X) and call b() if there is at least *1 permit*.

The *1 permit* is only available when **T2** calls signal(S_X) after X() executes.

Solution 1b: Barriers: a() < Y(), X() < b()

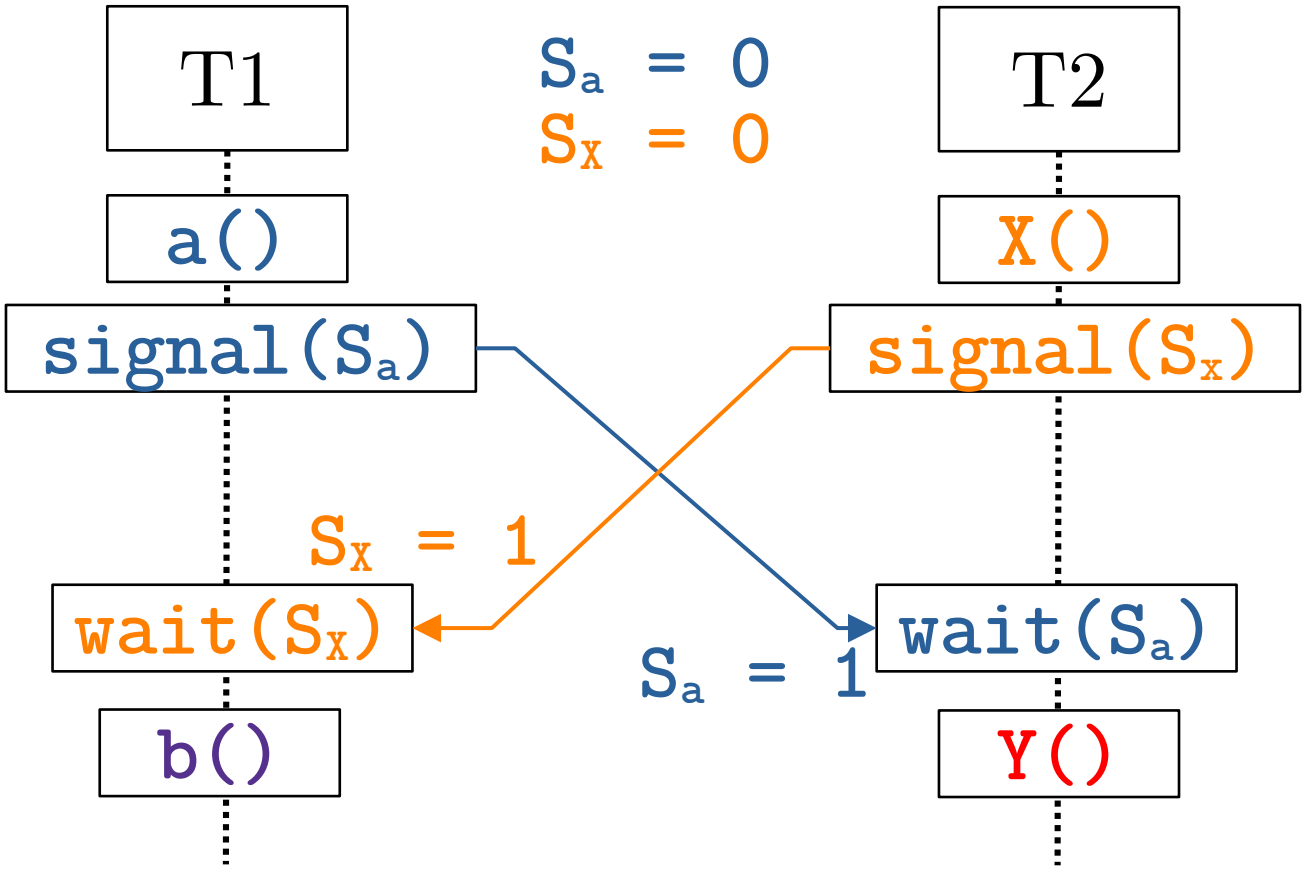

Synchronization Task/Problem 2: **Mutual Exclusion**

[Problem & Solution Descriptions]

Task/Problem 2: Mutual Exclusion

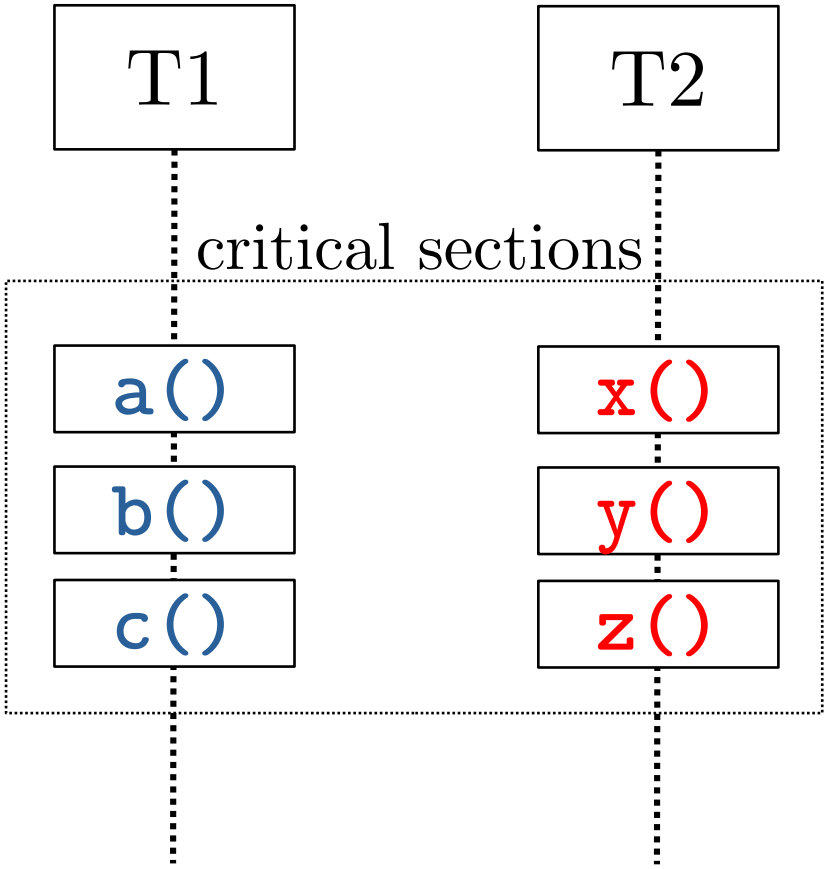

VALID orders

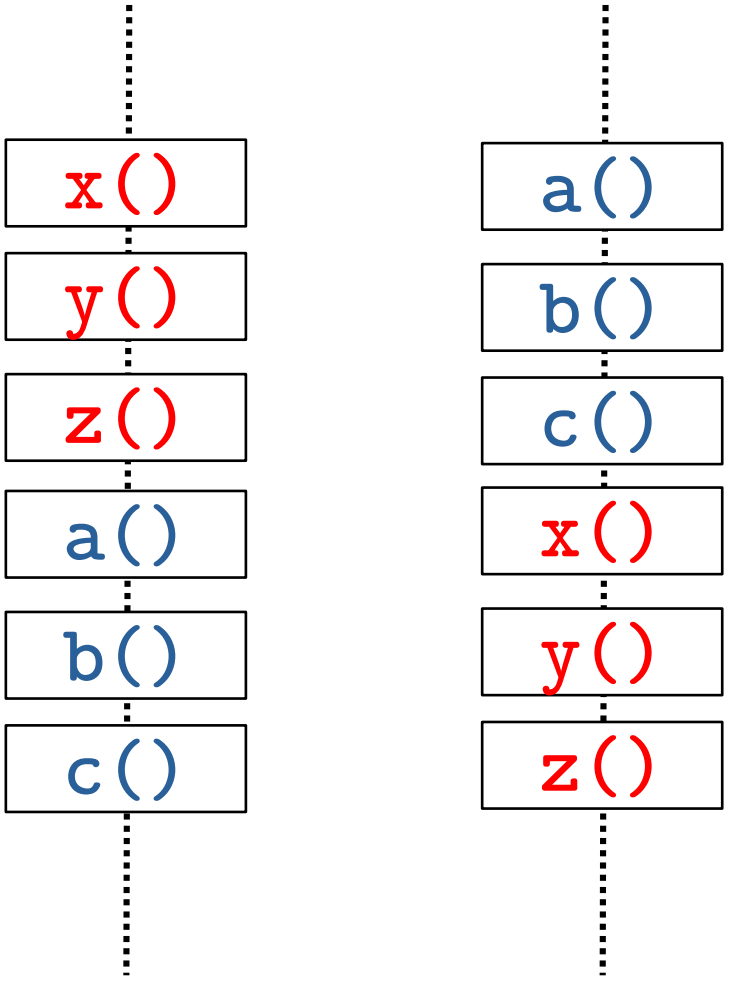

Constraint: Instructions/Operations in T1 and T2's critical sections should NOT interleaved.

Task/Problem 2: Mutual Exclusion

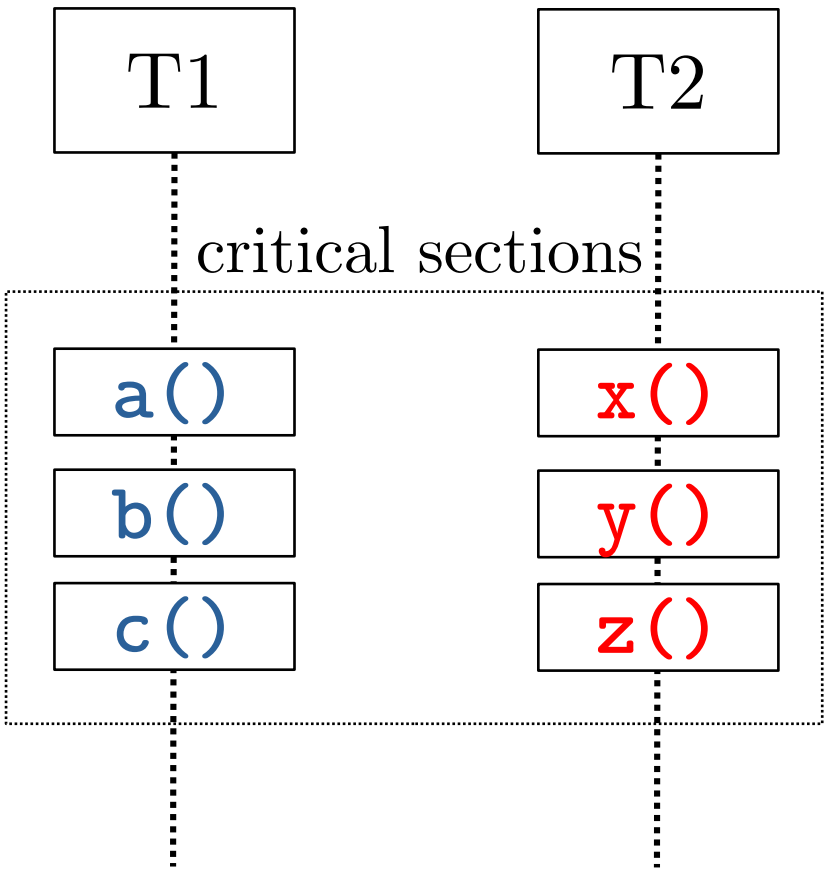

some INVALID orders

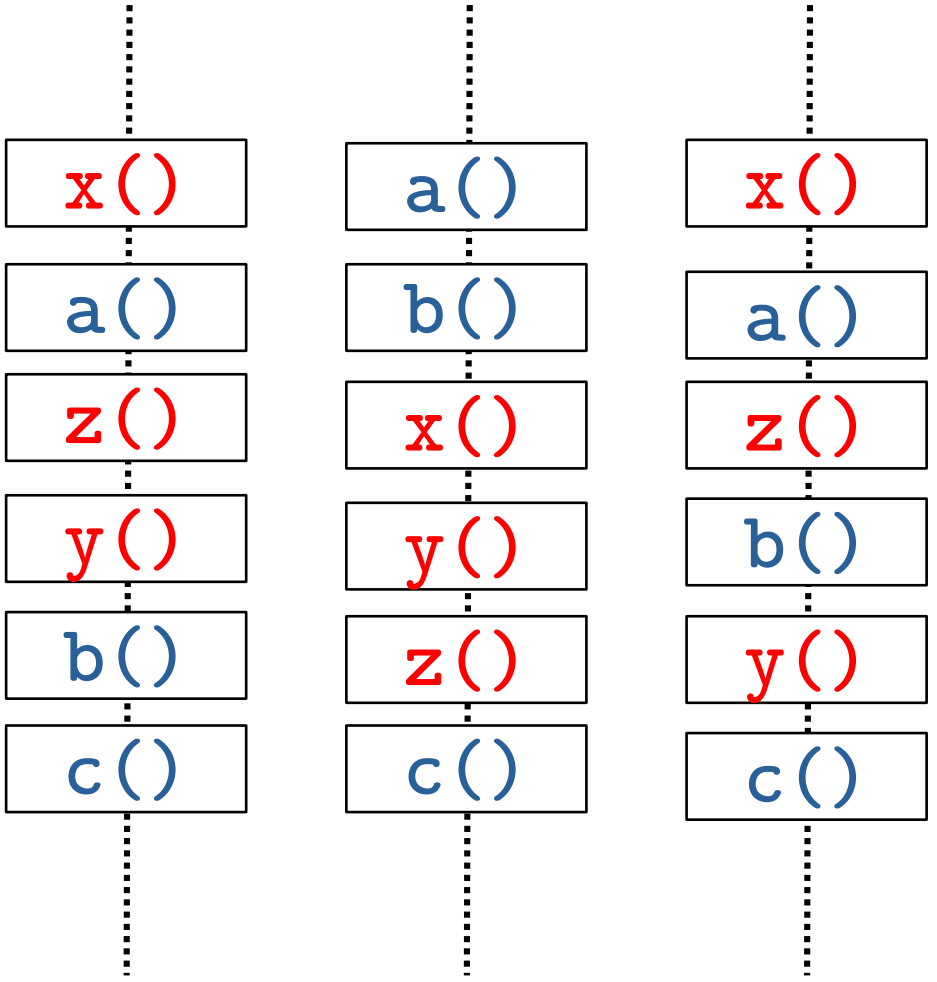

Constraint: Instructions/Operations in T1 and T2's critical sections should NOT interleaved.

Solution 2: Mutual Exclusion

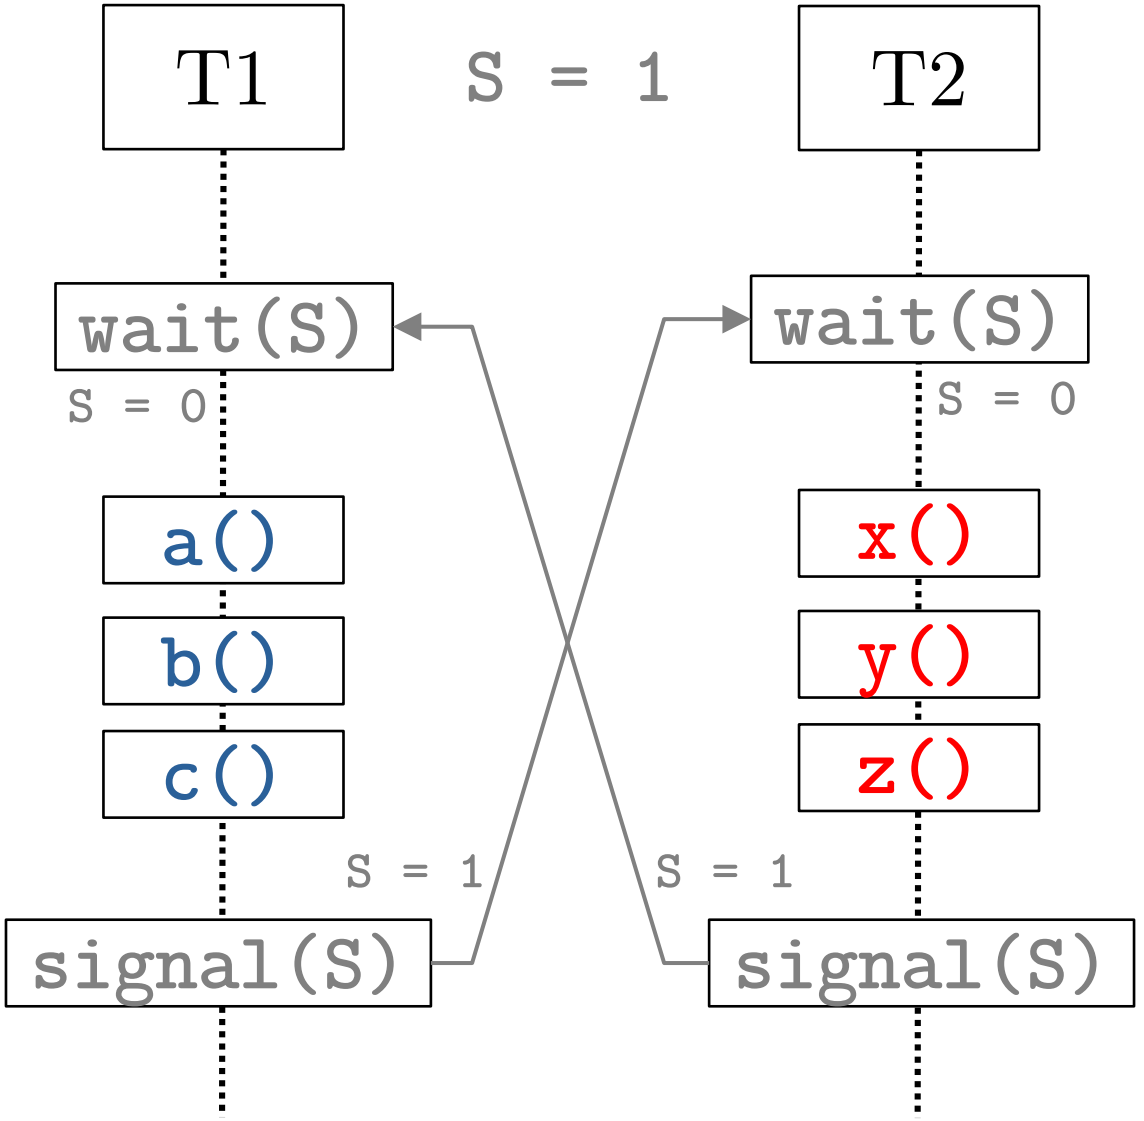

`S = 1` means *1 permit* is available.

Only one thread, **T1** or **T2,** will be given the permit and can exit the `wait(S)` loop, set `S = 0`, and enter its critical section.

After executing its critical section, the thread will now call `signal(S)` and set `S = 1` which gives the other thread a permit to enter its critical section.

Solution 2: Mutual Exclusion

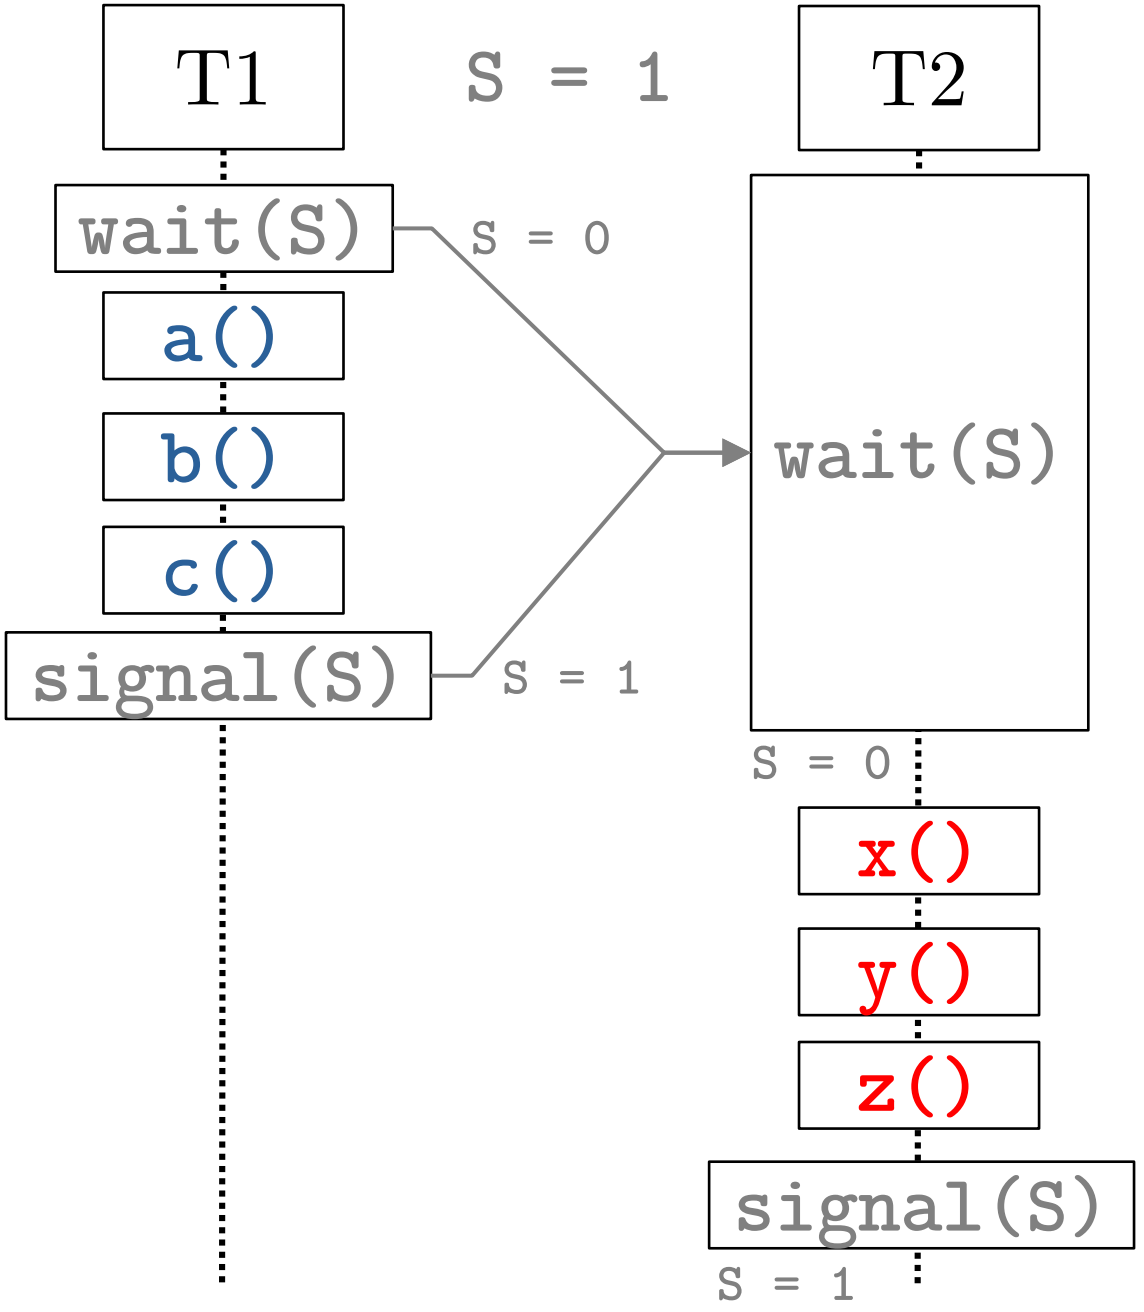

`S = 1` means *1 permit* is available.

Only one thread, **T1** or **T2,** will be given the permit and can exit the `wait(S)` loop, set `S = 0`, and enter its critical section.

After executing its critical section, the thread will now call `signal(S)` and set `S = 1` which gives the other thread a permit to enter its critical section.

Solution 2: Mutual Exclusion

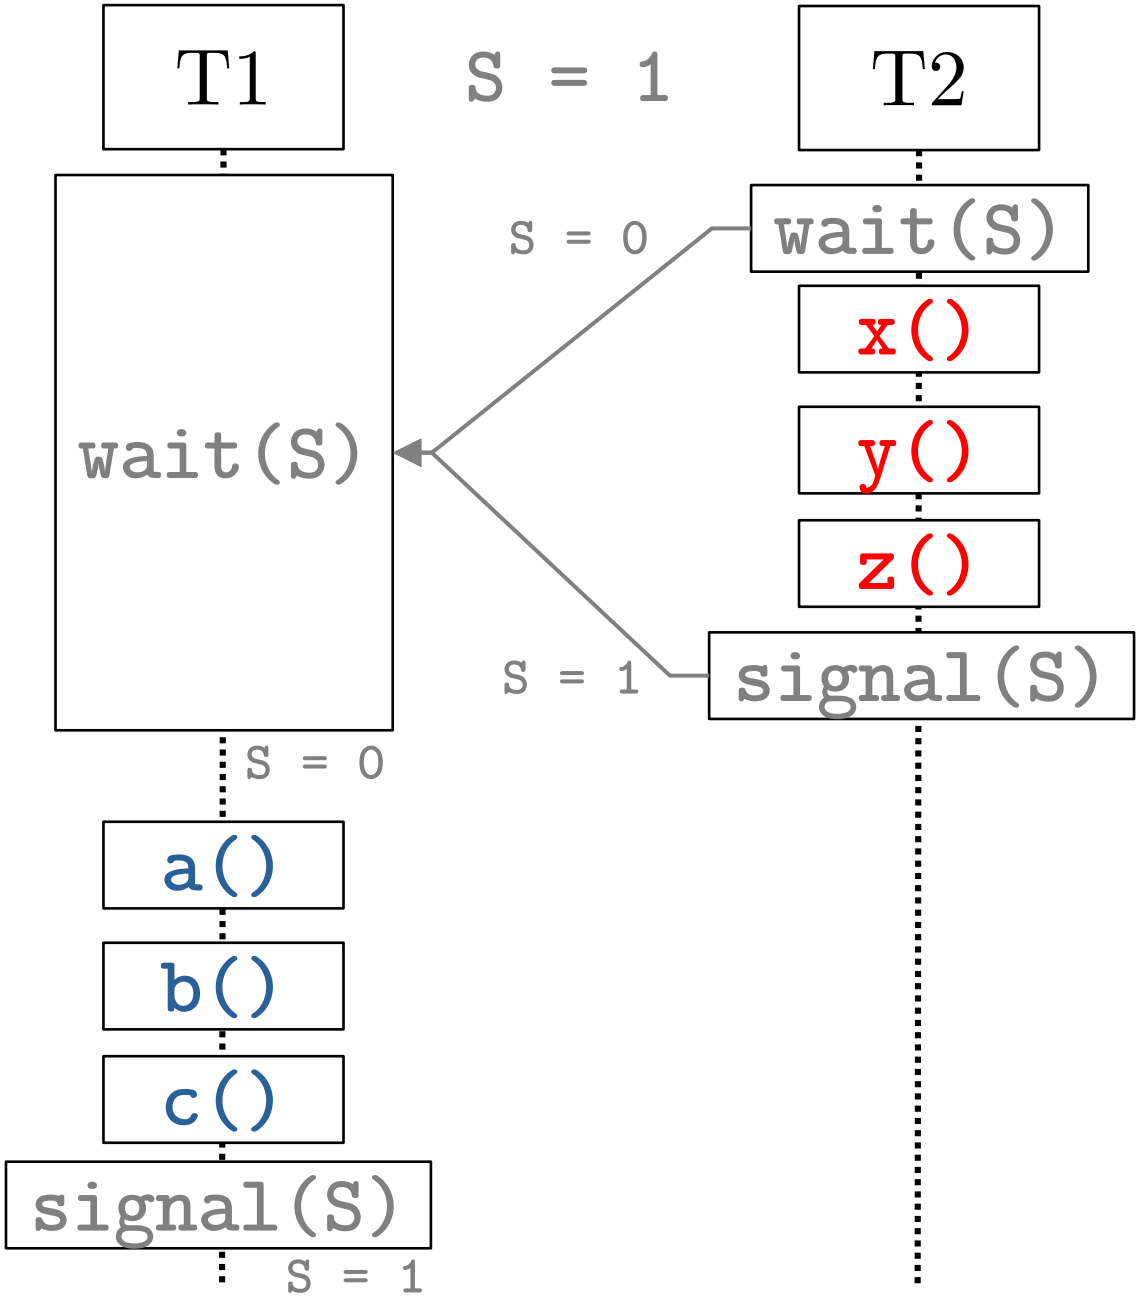

`S = 1` means *1 permit* is available.

Only one thread, **T1** or **T2,** will be given the permit and can exit the `wait(S)` loop, set `S = 0`, and enter its critical section.

After executing its critical section, the thread will now call `signal(S)` and set `S = 1` which gives the other thread a permit to enter its critical section.

Synchronization Task/Problem 3a: **Producer-Consumer**

[Problem & Solution Descriptions]

Task/Problem 3a: Producer(p)-Consumer(c) Problem

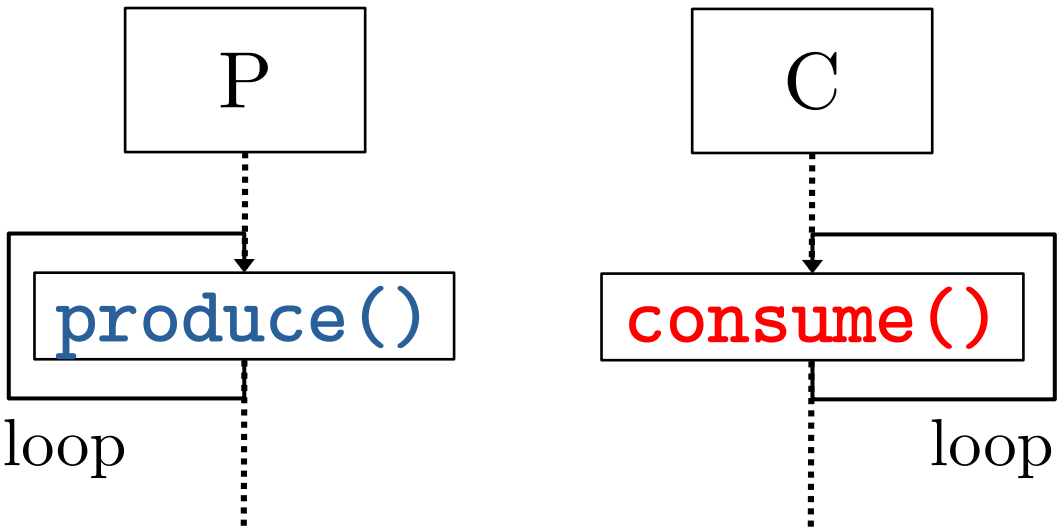

c() = consume()
p() = produce()

some VALID orders

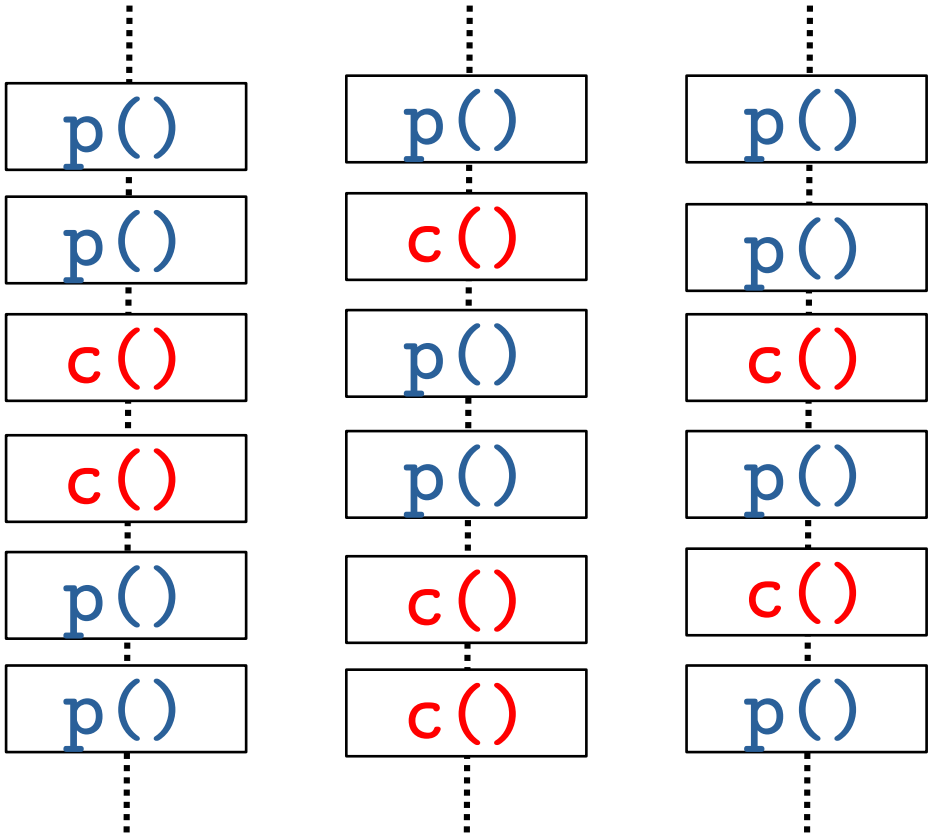

Constraint: Before a `consume()` call can occur a corresponding `produce()` call should have occurred first.

Task/Problem 3a: Producer(p)-Consumer(c) Problem

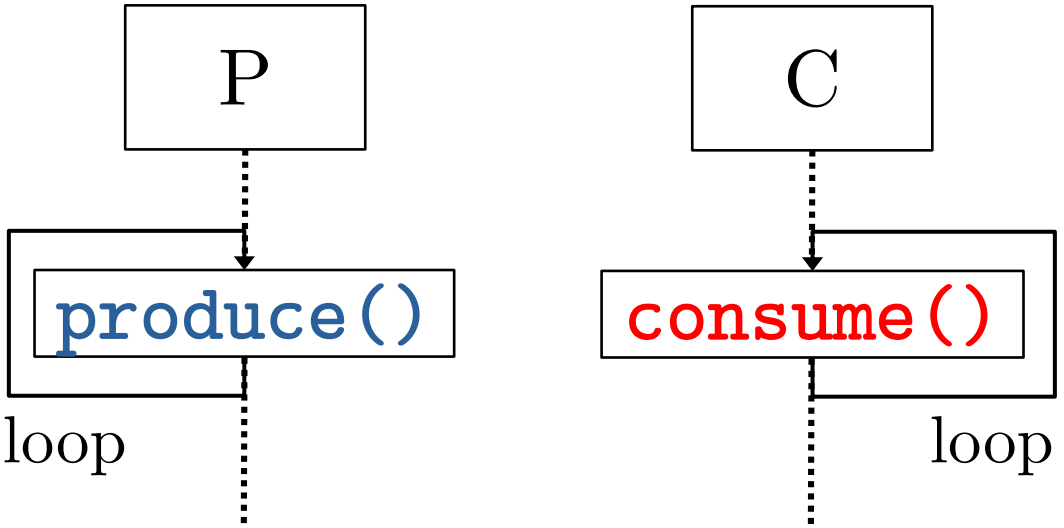

c() = consume()
p() = produce()

some **INVALID** orders

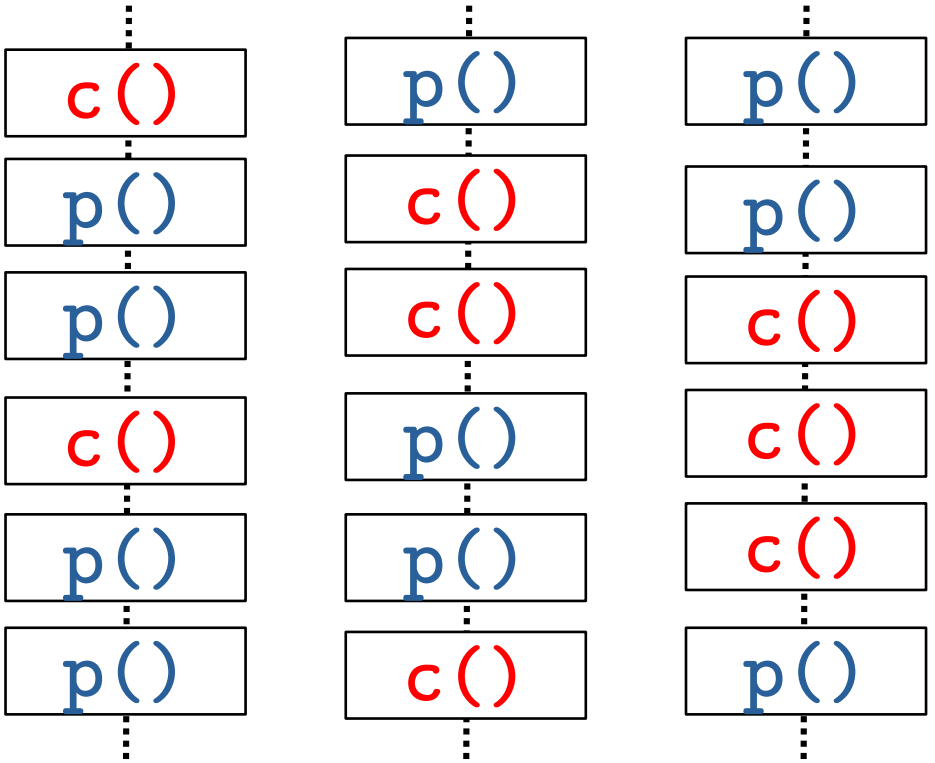

Constraint: Before a `consume()` call can occur a corresponding `produce()` call should have occurred first

Task/Problem 3a: Producer(p)-Consumer(c) Problem

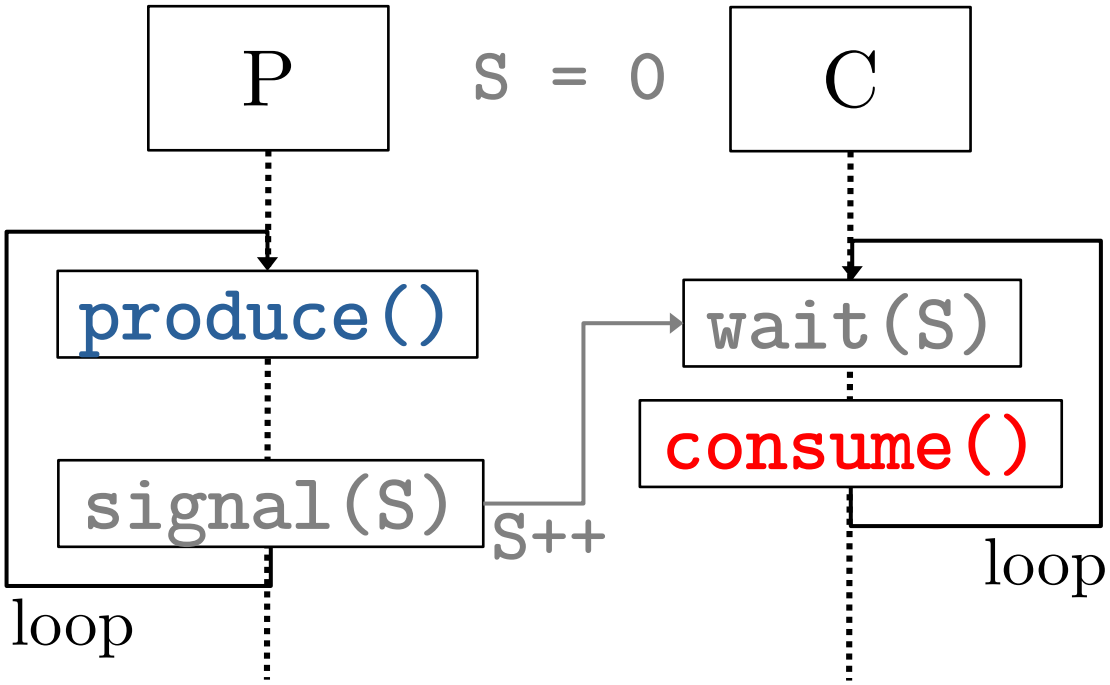

`S = 0` means 0 *permit* is available.

Thread **C** is the one calling `wait(S)` and is waiting for a permit to call `consume()`.

Thread **P** will call `signal(S)` only after calling `produce()`. `signal(S)` will increment `S`.

Thread C can only call `consume()` if there is a corresponding call to `produce()` prior.

Synchronization Task/Problem 3b: **Bounded Buffer**

[Problem & Solution Descriptions]

Task/Problem 3b: Bounded Buffer. i.e. Buffer Size = 3

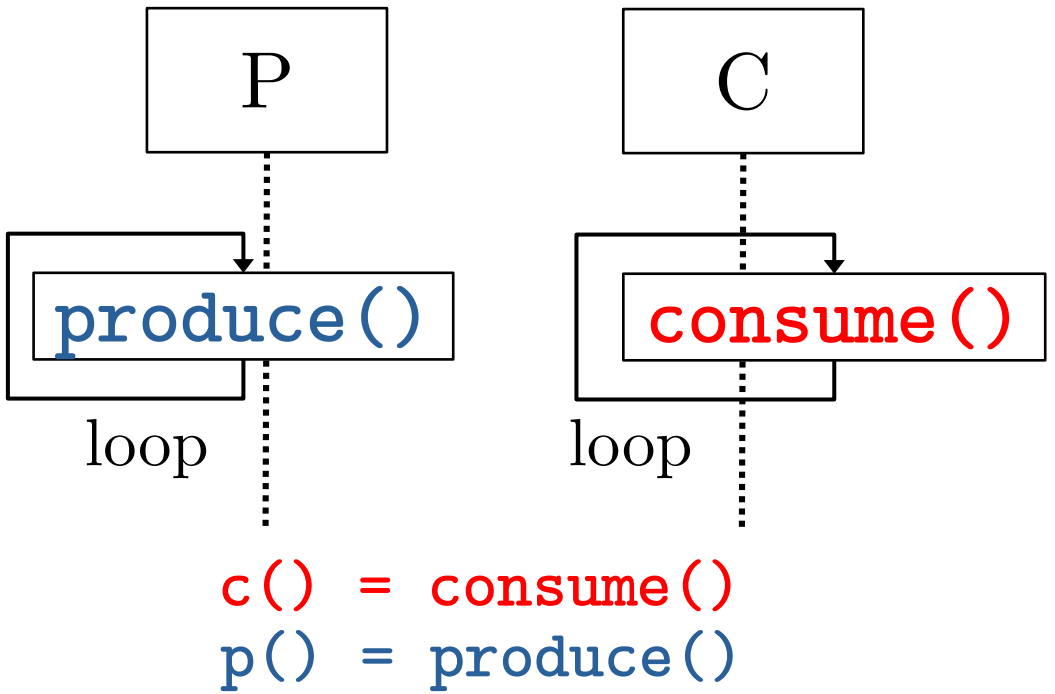

some VALID orders

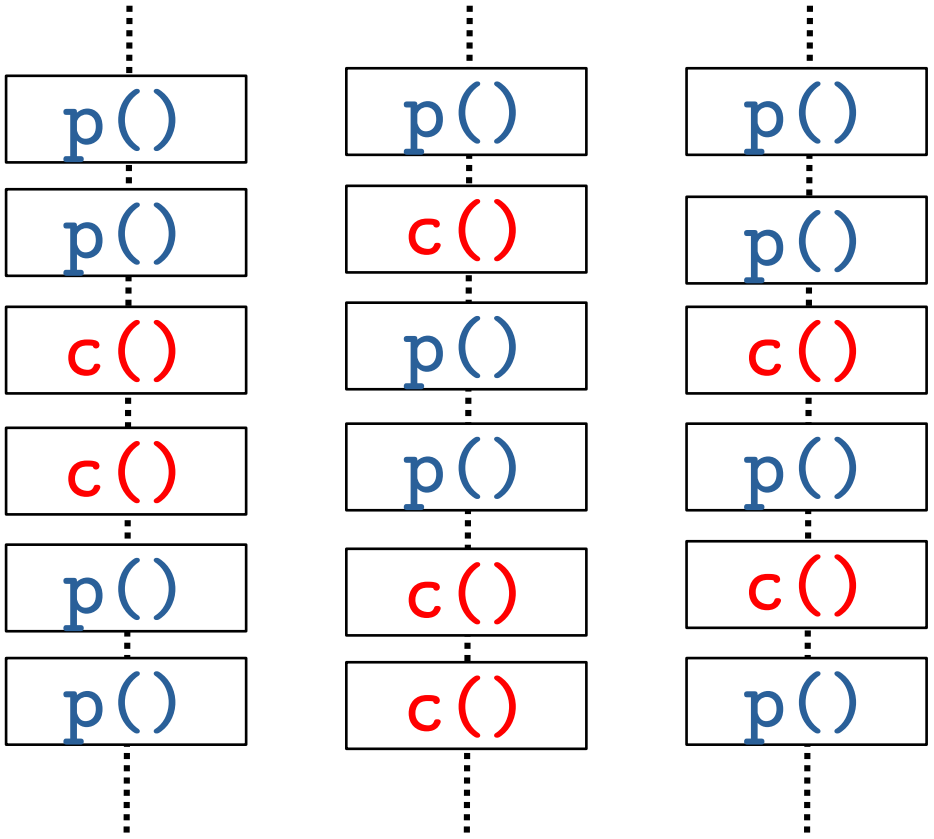

Constraint: Before a `consume()` call can occur a corresponding `produce()` call should have occurred first. At **most 3 calls** to `produce()` can occur that have no corresponding calls to `consume()`. i.e. Buffer is full after 3 calls to `produce()` (and 0 calls to `consume()`).

Task/Problem 3b: Bounded Buffer. i.e. Buffer Size = 3

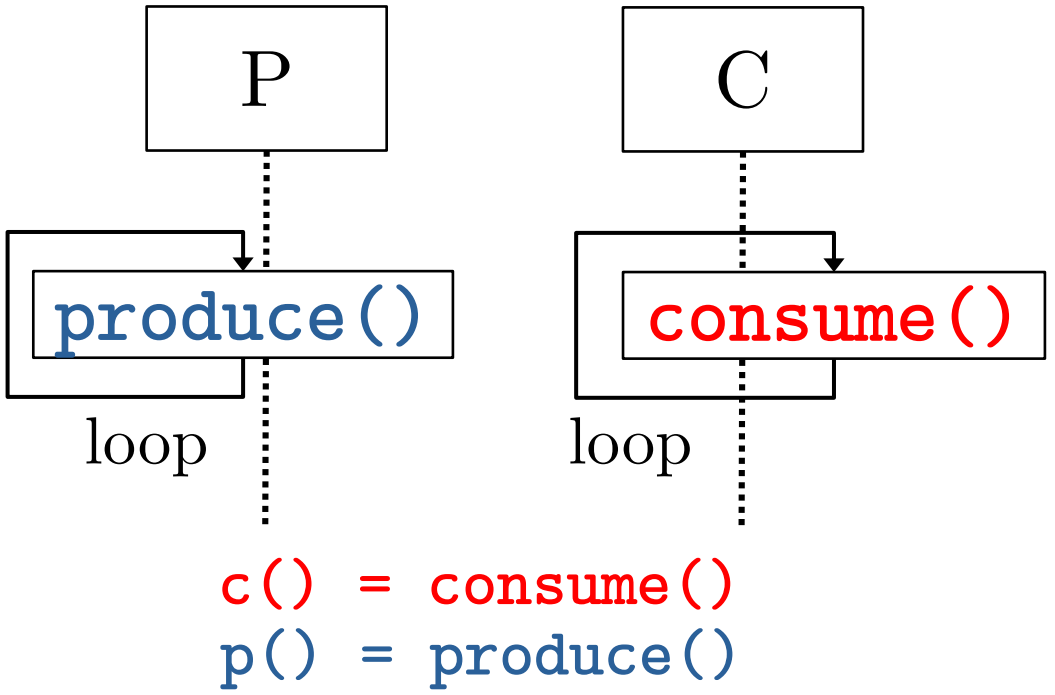

some INVALID orders

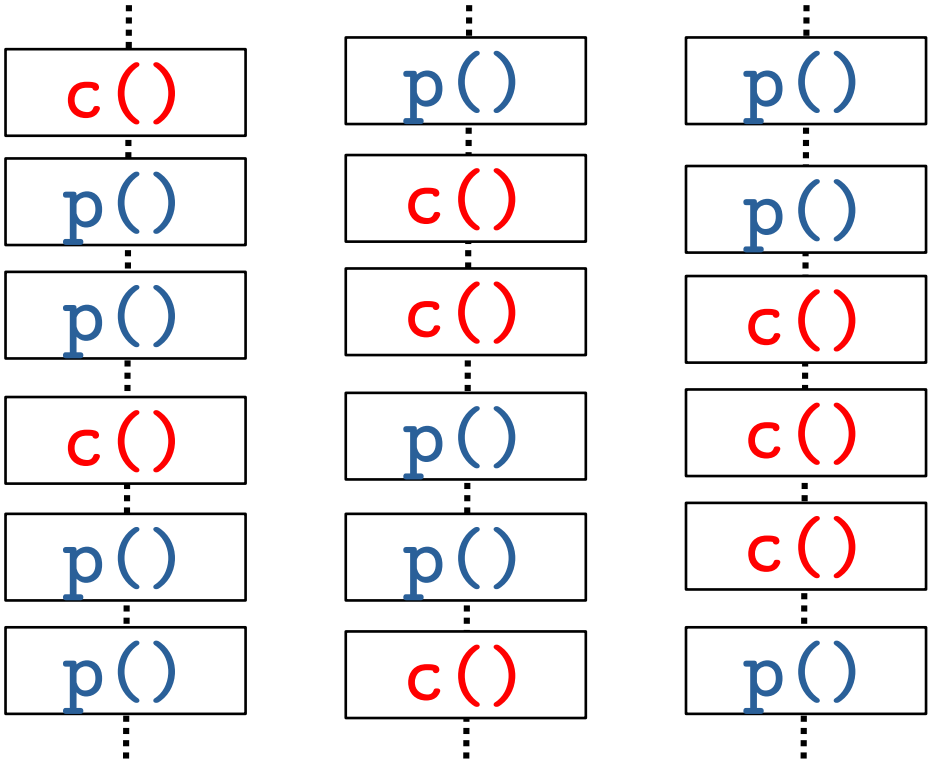

Constraint: Before a `consume()` call can occur a corresponding `produce()` call should have occurred first. At **most 3 calls** to `produce()` can occur that have no corresponding calls to `consume()`. i.e. Buffer is full after 3 calls to `produce()`.

Task/Problem 3b: Bounded Buffer. i.e. Buffer Size = 3

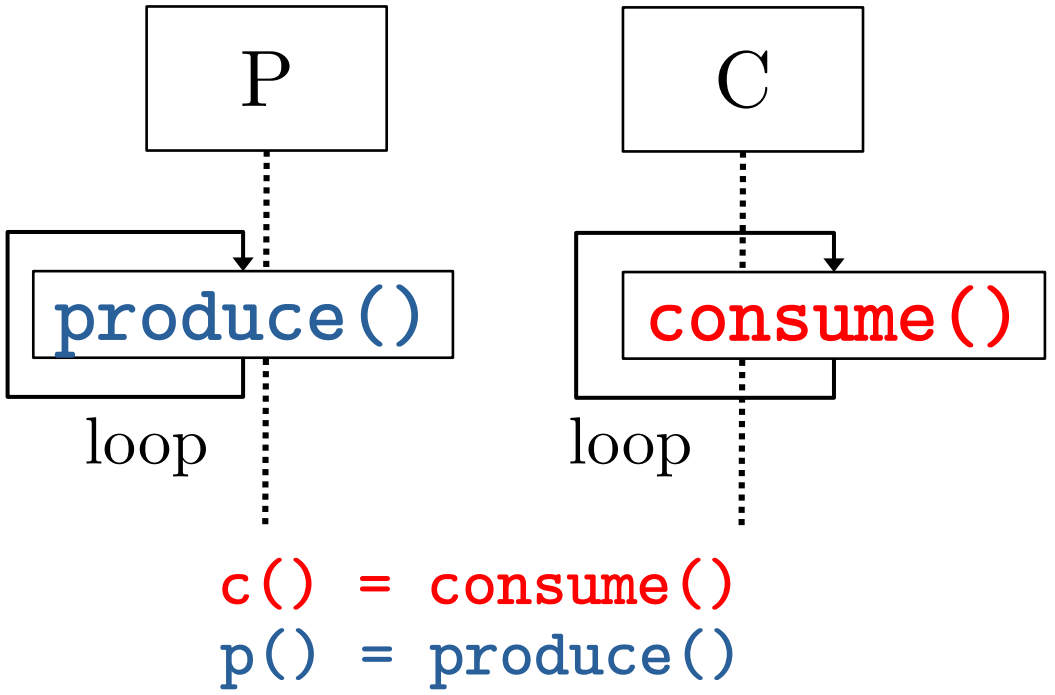

some INVALID orders

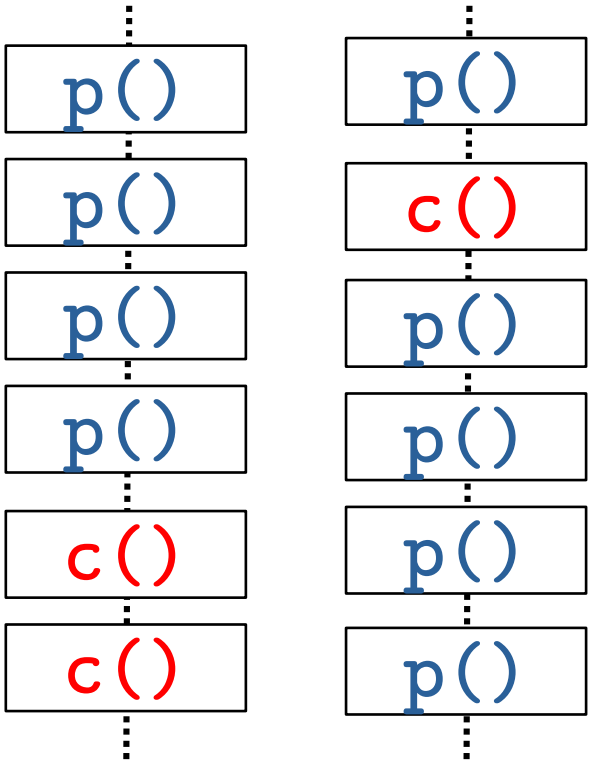

Constraint: Before a `consume()` call can occur a corresponding `produce()` call should have occurred first. At **most 3 calls** to `produce()` can occur that have no corresponding calls to `consume()`. i.e. Buffer is full after 3 calls to `produce()`.

Task/Problem 3b: Bounded Buffer. i.e. Buffer Size = 3

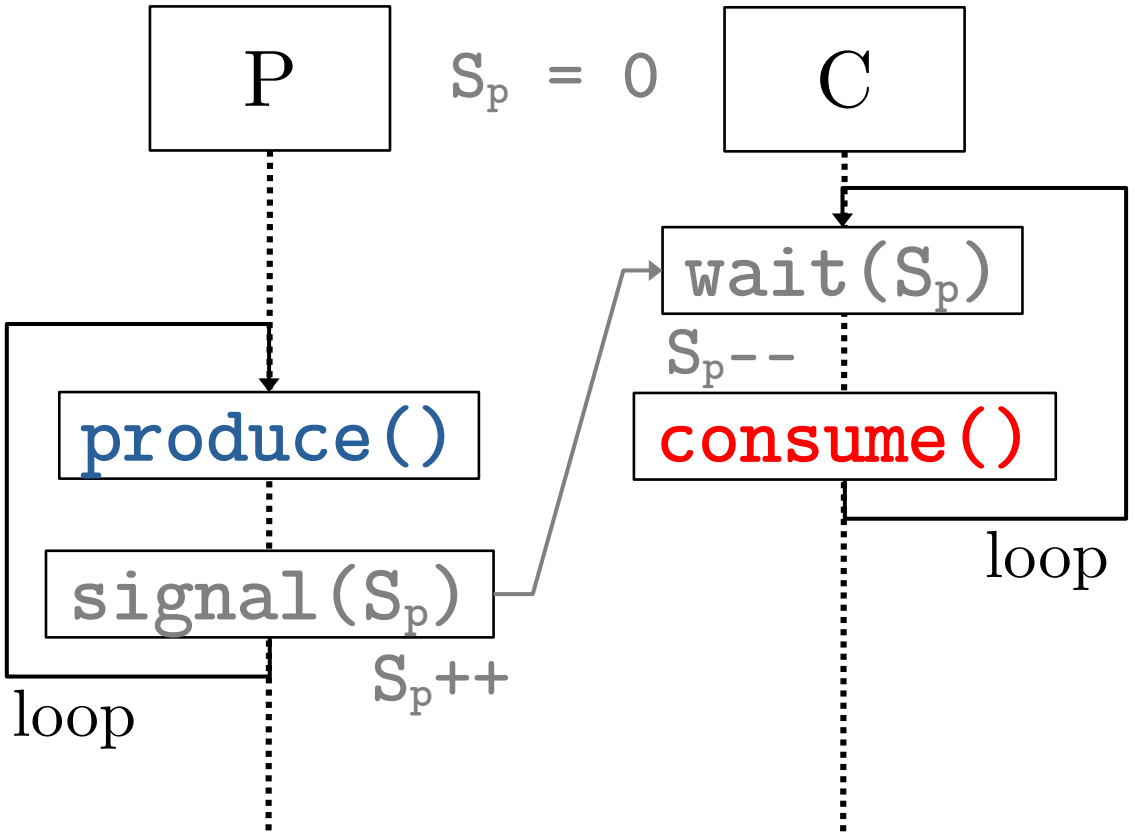

$S_p = 0$ means 0 '**consume**' *permit* is available.

Thread **C** is the one calling `wait(`S_p`)` and is waiting for a permit before calling `consume()`.

Thread **P** will call `signal(`S_p`)` only after calling `produce()`. `signal(`S_p`)` will increment S_p.

Thread C can only call `consume()` if there is a corresponding call to `produce()` prior.

Task/Problem 3b: Bounded Buffer. i.e. Buffer Size = 3

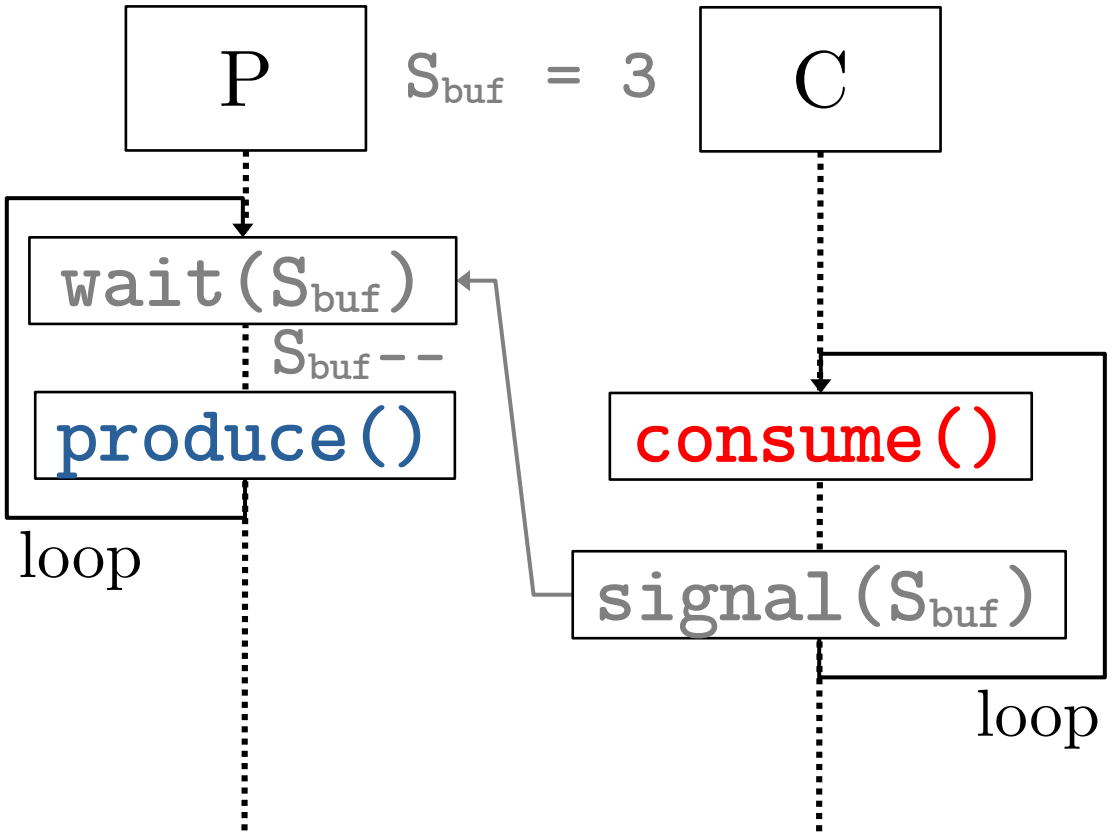

S_{buf} = 3 means 3 'produce' *permit* is available.

Thread **P** is the one calling wait(S_{buf}) and is waiting for a permit before calling produce().

Thread **C** will call signal(S_p) only after calling consume() . signal(S_p) will increment S_p.

Thread C can only call consume() if there is a corresponding call to produce() prior.

Task/Problem 3b: Bounded Buffer. i.e. Buffer Size = 3

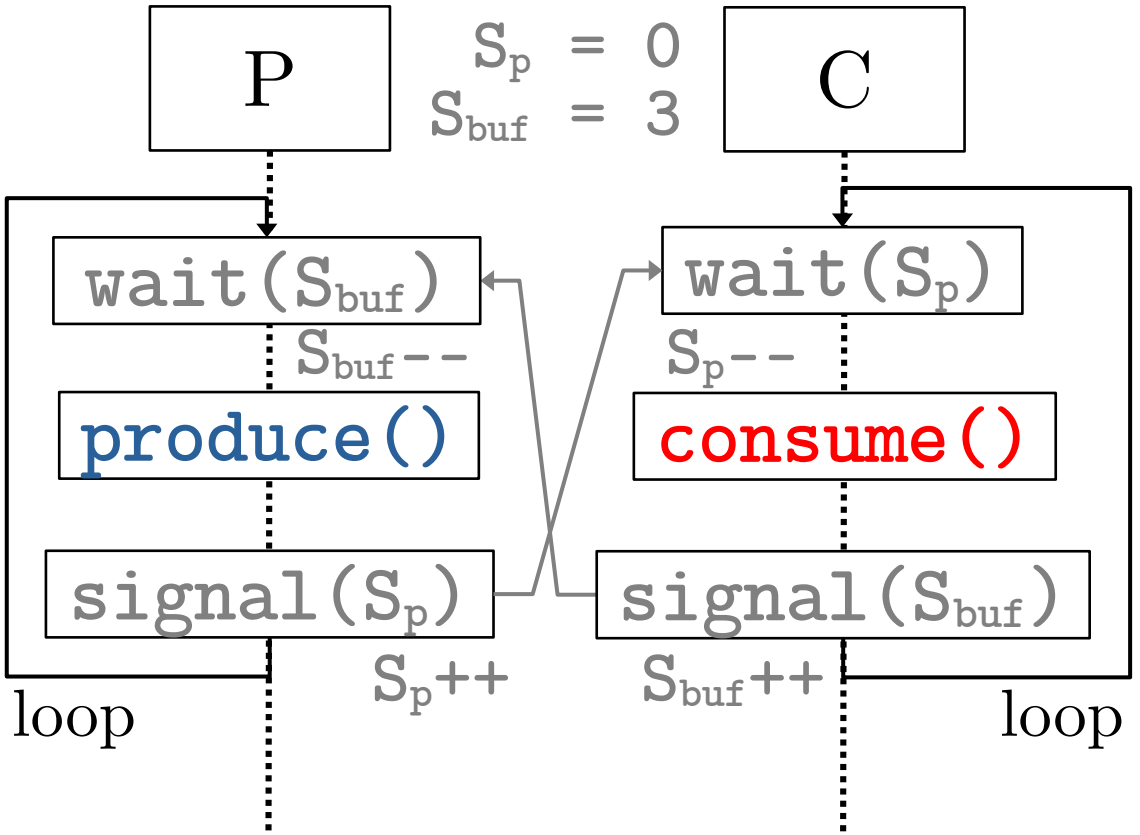

S_p represents the number of '**consume**' *permits*.
S_{buf} represents the number of '**produce**' *permits*.

Thread **P** calls `wait(`S_{buf}`)` to wait for a '**produce**' *permit* before calling **produce()**. After calling **produce()**, thread **P** calls `signal(`S_p`)` to increment the **consume** permit.

Thread **C** calls `wait(`S_p`)` to wait for a '**consume**' *permit* before calling **consume()**. After calling **consume()**, thread **C** calls `signal(`S_{buf}`)` to increment the **produce** permit.

Synchronization Task/Problem 4: **Readers-Writers**

[Problem & Solution Descriptions]

Task/Problem 4: Reads-Writers

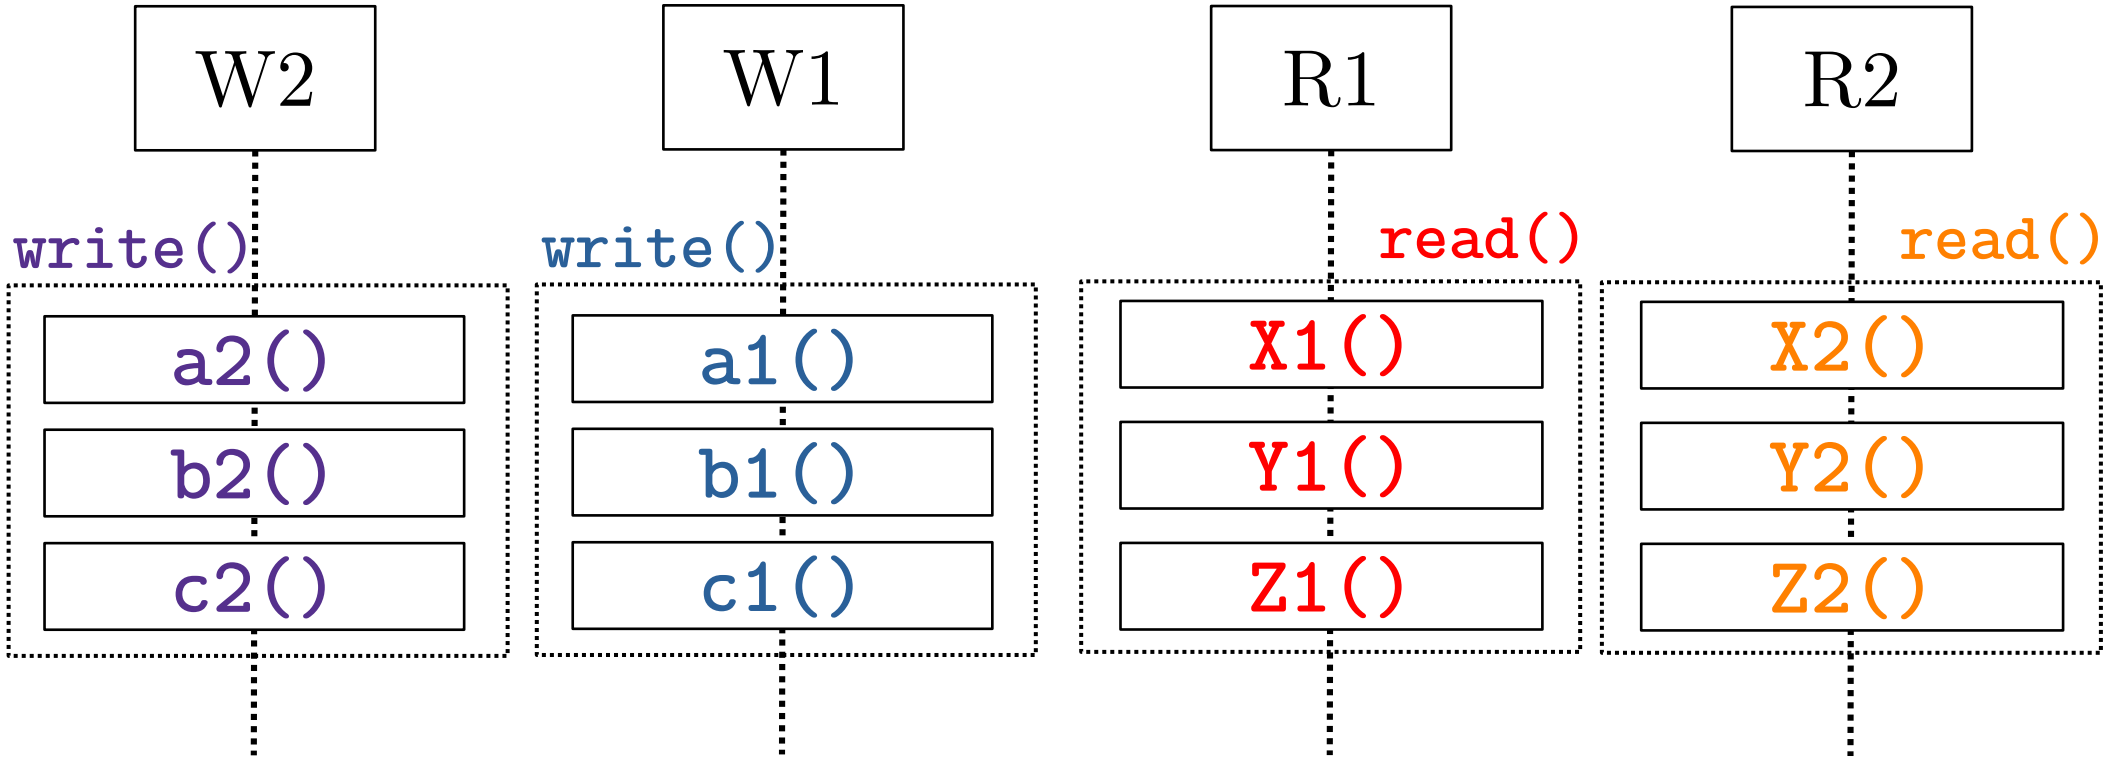

Constraint: A **writer thread** has exclusive access. No other threads (writer or reader) can access their `write()` or `read()` function if there is a writer thread calling `write()`. **Reader threads**, as a group, can access call their `read()` function concurrently and no writer threads can call their `write()` function while at least one reader thread is calling `read()`.

Task/Problem 4: Reads-Writers

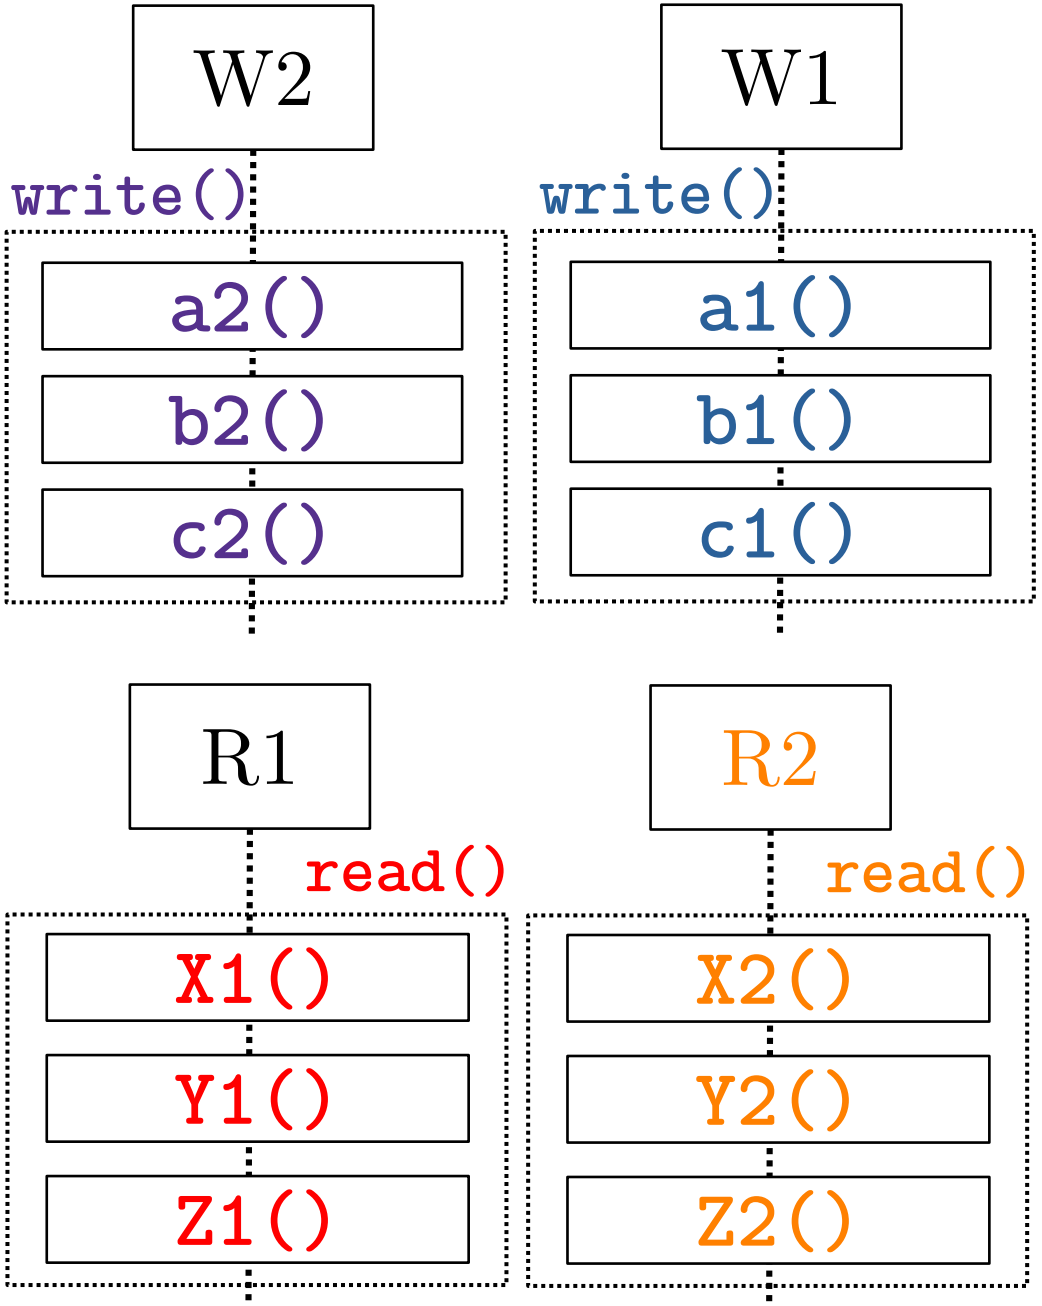

some INVALID orders

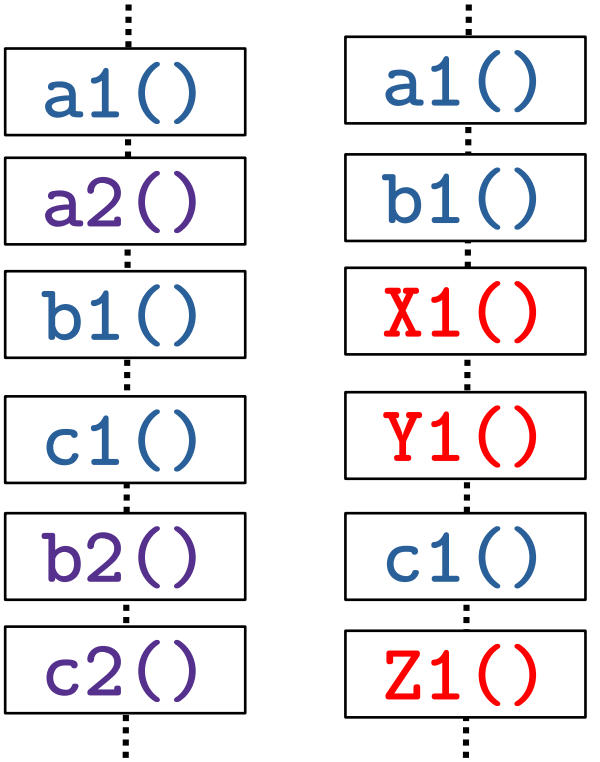

Task/Problem 4: Reads-Writers

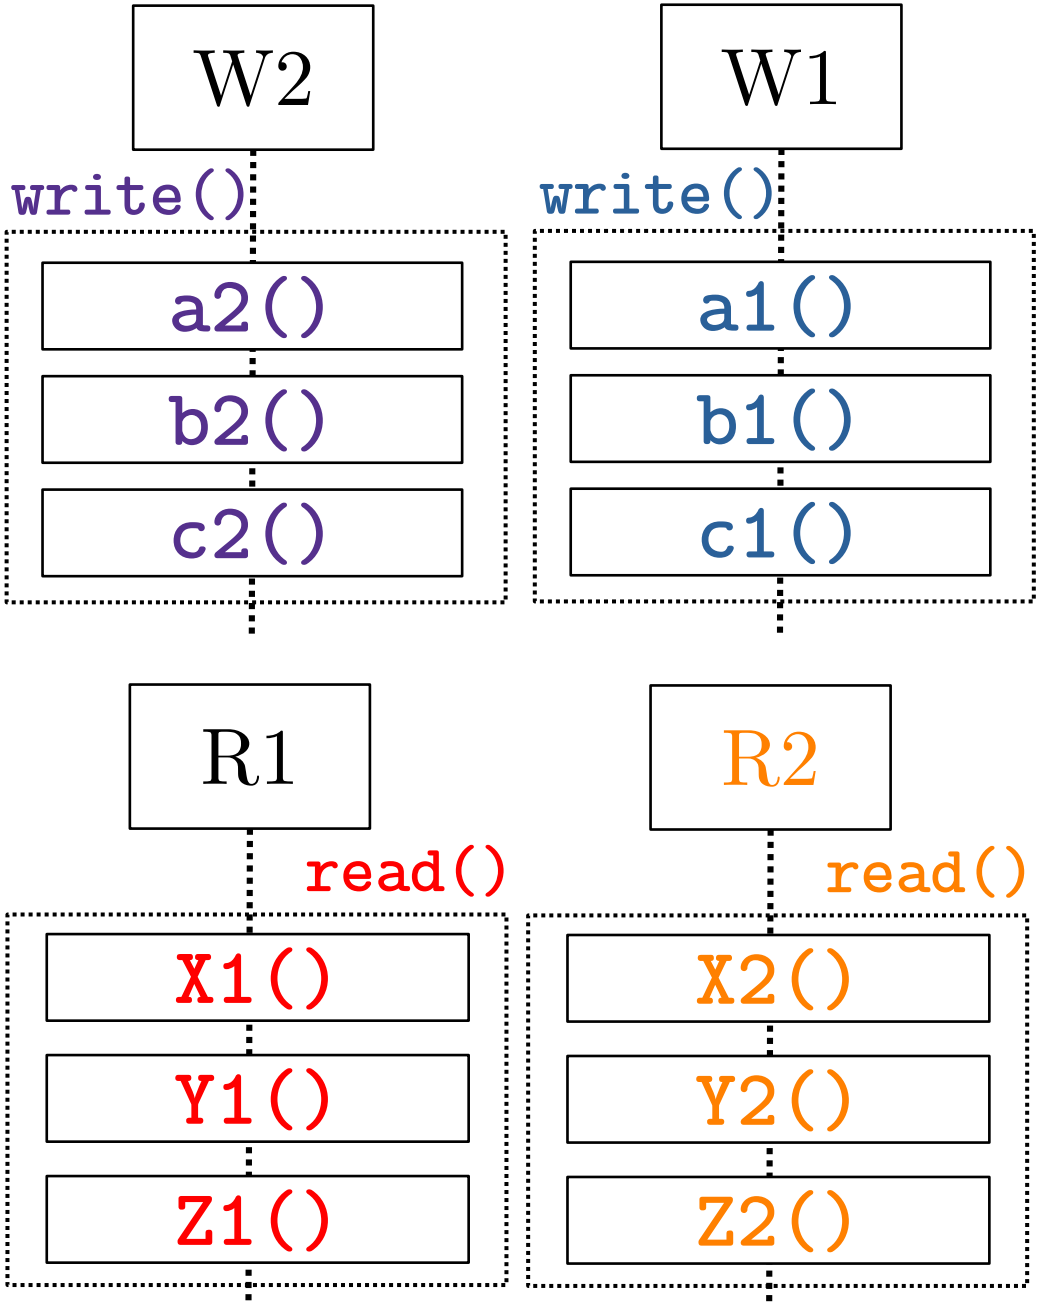

some VALID orders

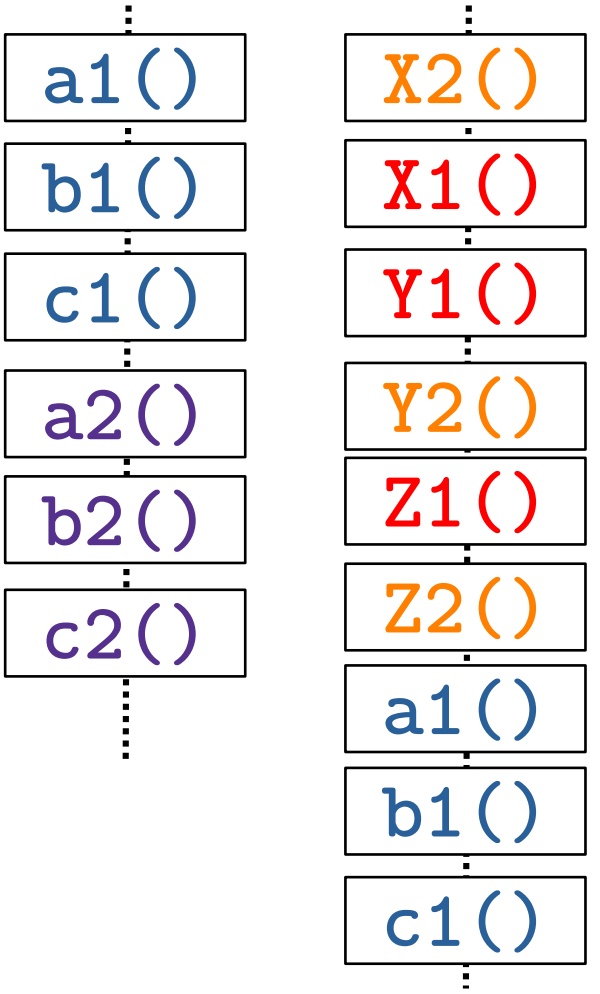

Solution 4: Reads-Writers

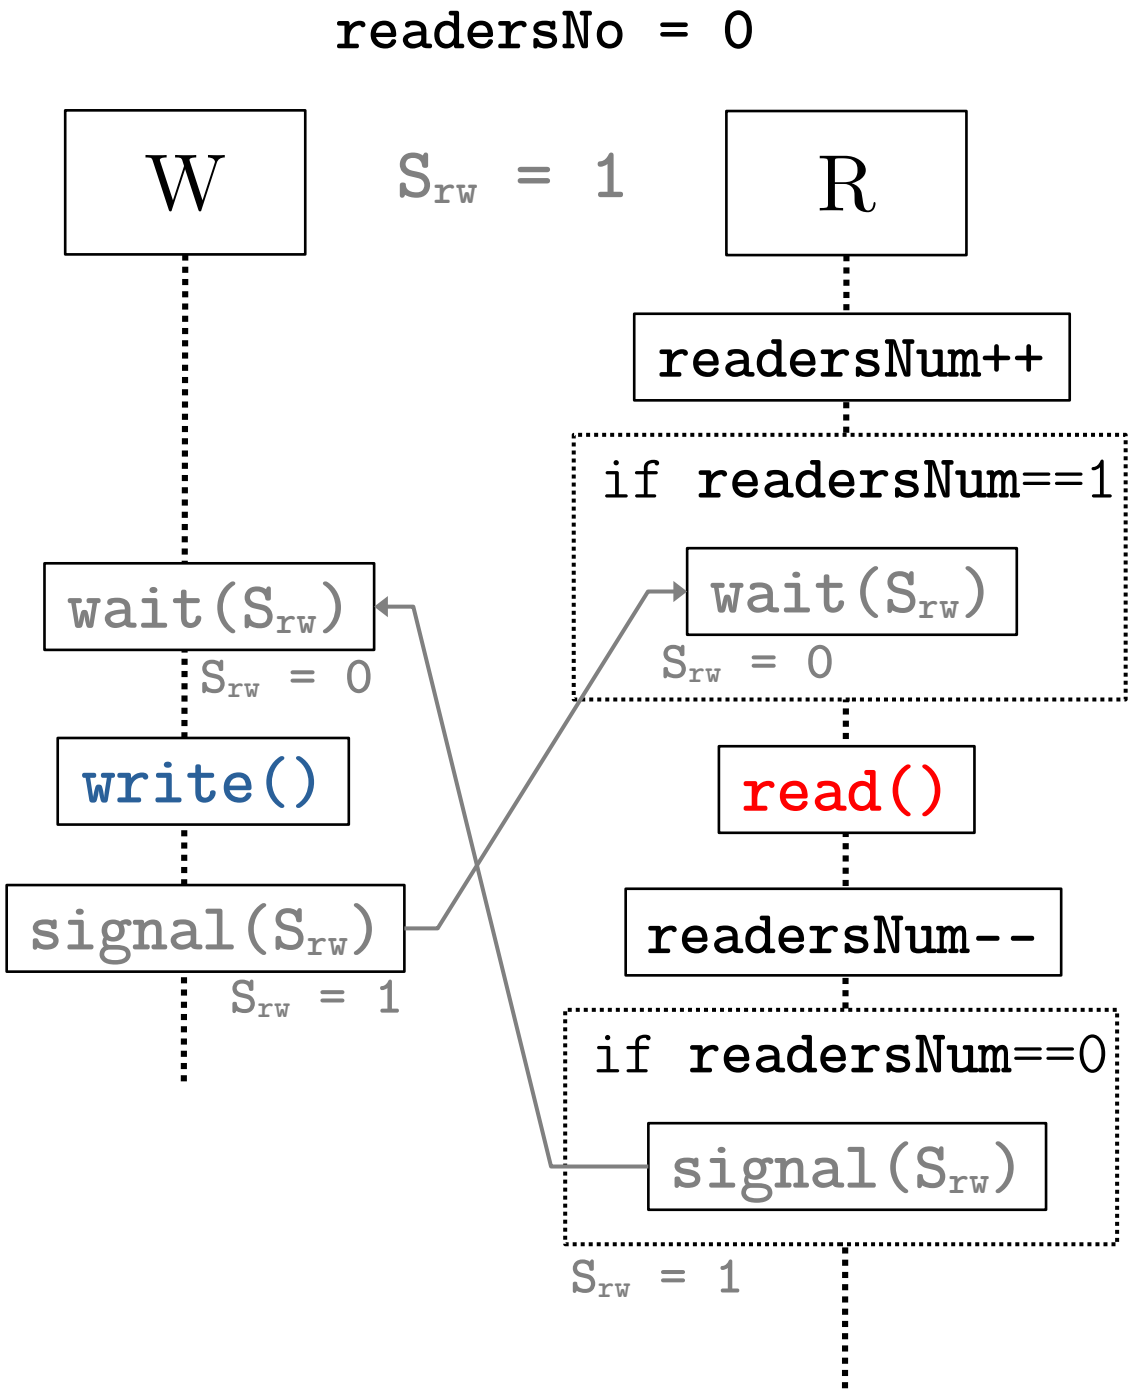

S_{rw} represents the **read-write** *permit*. Initially, S_{rw} = 1.

Writer **W** calls wait(S_{rw}) to wait for a **read-write** *permit* before calling **write()**. After calling **write()**, writer **W** calls signal(S_{rw}) to release the **read-write** *permit*. S_{rw} = 1

Reader **R** increments the variable **readersNo**. If **readersNo==1** after being incremented, then Reader R is the first reader thread trying to perform the **read()** operation.

Solution 4: Reads-Writers

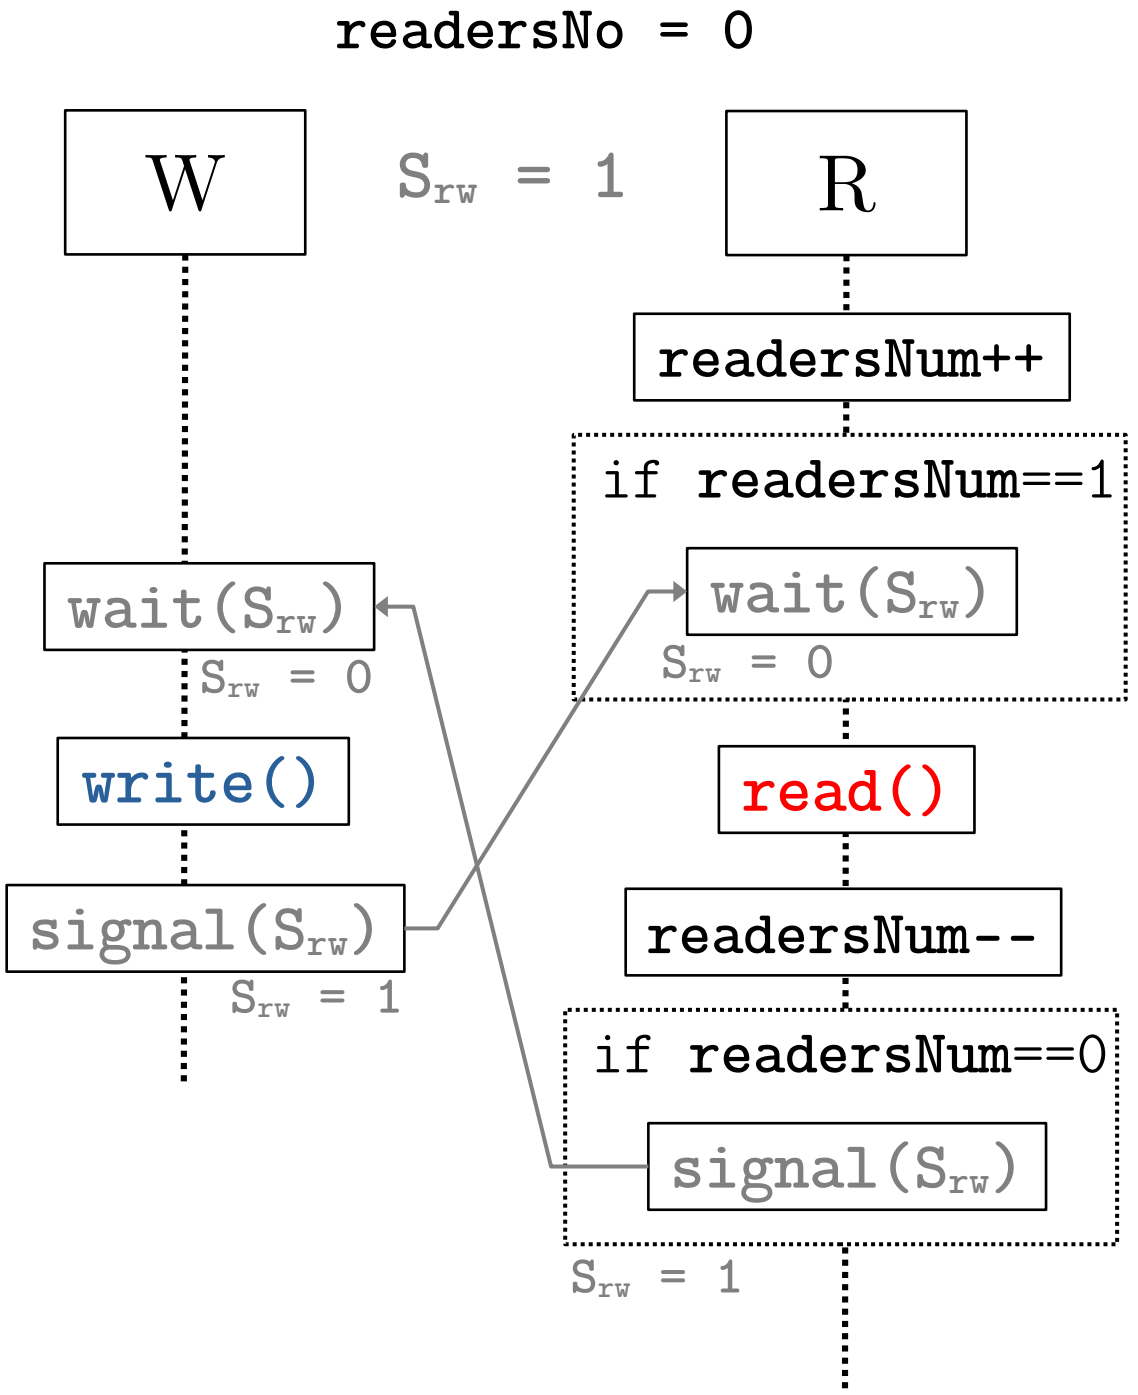

Reader **R** increments the variable **readersNum**. If **readersNum==1** after being incremented, then reader **R** is the first reader thread trying to perform the **read()** operation. If it is the first reader thread, then it will call wait(S_{rw}) to wait for the **read-write** *permit*.

If reader **R** is not the first reader (i.e. **readersNum>1**), then it does not need to wait for the **read-write** *permit* because some reader thread is currently calling **read()** has acquired it previously.

Solution 4: Reads-Writers

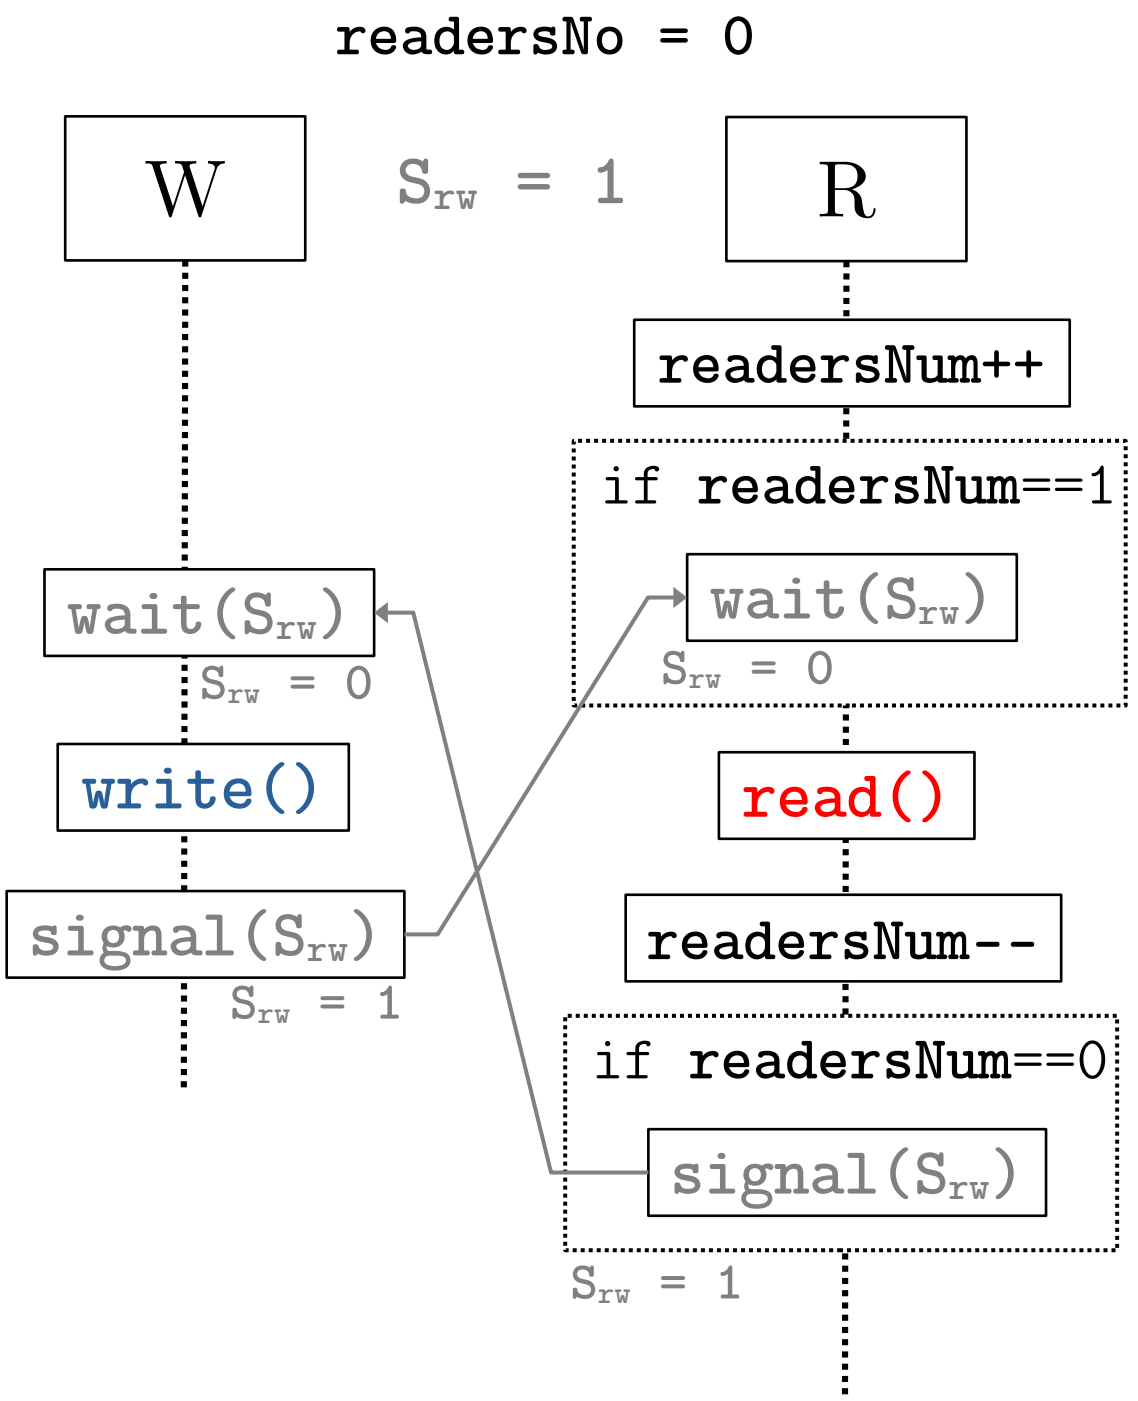

If reader **R** is not the first reader (i.e. **readersNum>1**), then it does not need to wait for the **read-write** *permit* because some reader thread is currently calling **read()** has acquired it previously.

After the call to **read()**, reader **R** will decrement **readersNum**, and it will release the **read-write** *permit* if it is the last reader that performed **read()**. i.e. If **readersNum**==0.

Solution 4: Reads-Writers: Lightswitch Design Pattern

Threads waiting for room access

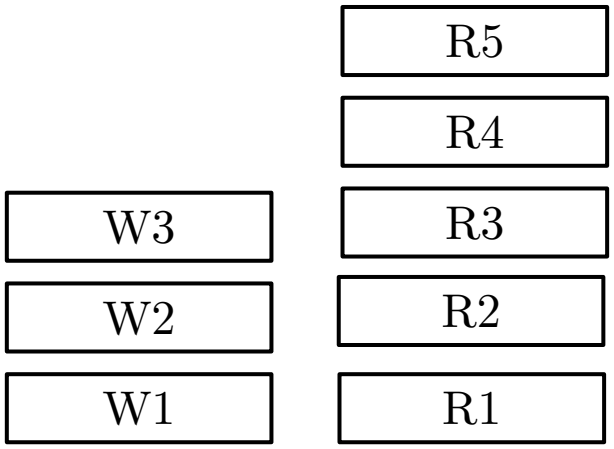

Room for reading/writing

Lights OFF = Free Room

Semaphore S_{rw}=1

Analogy:

- The room is for reading/writing.
- Free Room – Lights OFF
- Room in Use – Lights ON
- Semaphore S_{rw} = 1 means a free room (lights OFF)
- Semaphore S_{rw} = 0 means a room in use (lights ON)
- Threads want to enter the room in order to `read()` or `write()`.

Solution 4: Reads-Writers: Lightswitch Design Pattern

Threads waiting for room access

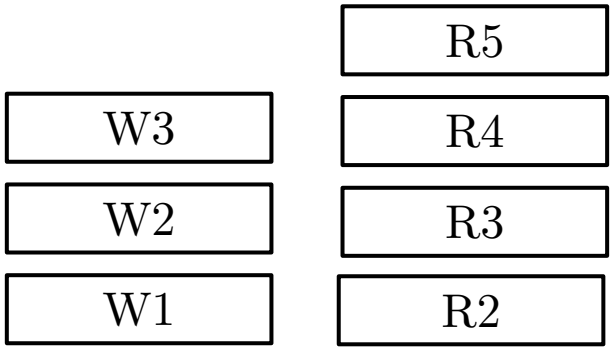

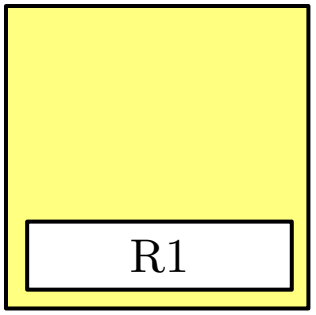

Room for reading/writing

Lights ON = Room in Use

Semaphore S_{rw}=0

- R1 enters the room and turns ON the light.
- S_{rw}=0 → **lights ON → Room in Use**
- When a reader is already in the room, other readers can enter the rule and perform `read()`.

Solution 4: Reads-Writers: Lightswitch Design Pattern

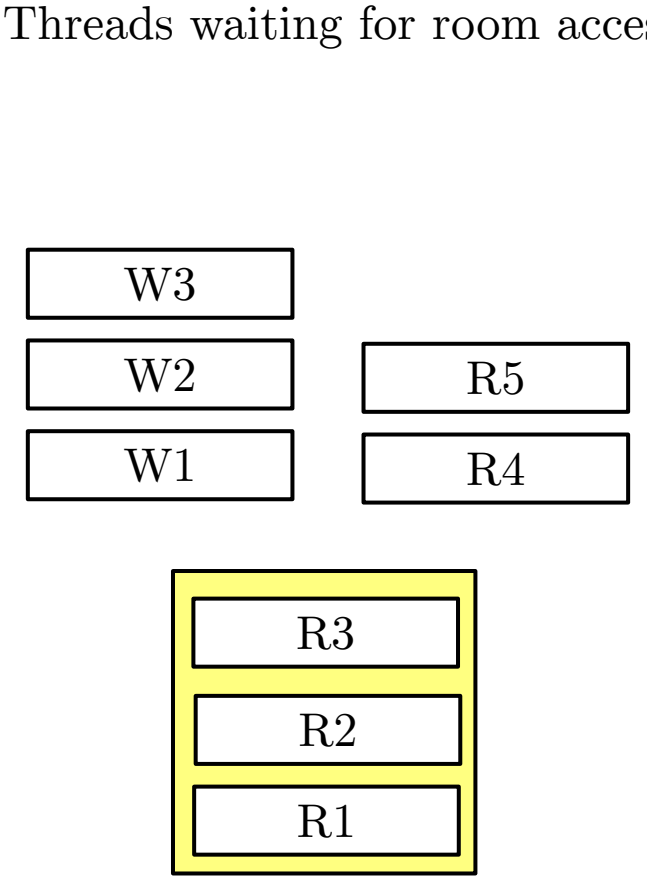

- R1 enters the room and turns ON the light.
- S_{rw}=0 → **lights ON → Room in Use**
- When a reader is already in the room, other readers can enter the rule and perform `read()`.

Solution 4: Reads-Writers: Lightswitch Design Pattern

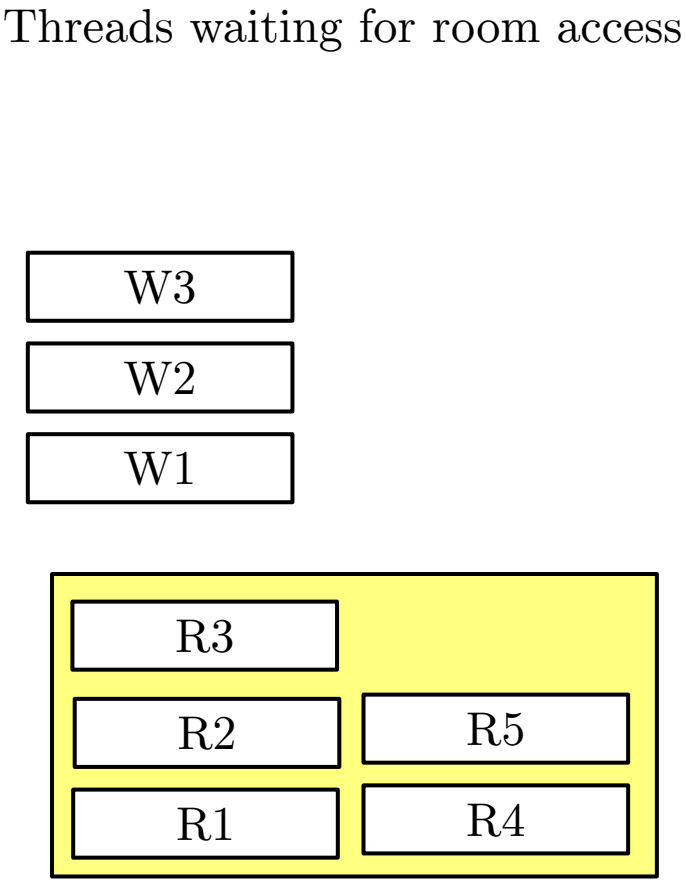

- R1 enters the room and turns ON the light.
- S_{rw}=0 → **lights ON → Room in Use**
- When a reader is already in the room, other readers can enter the rule and perform `read()`.

Solution 4: Reads-Writers: Lightswitch Design Pattern

Threads waiting for room access

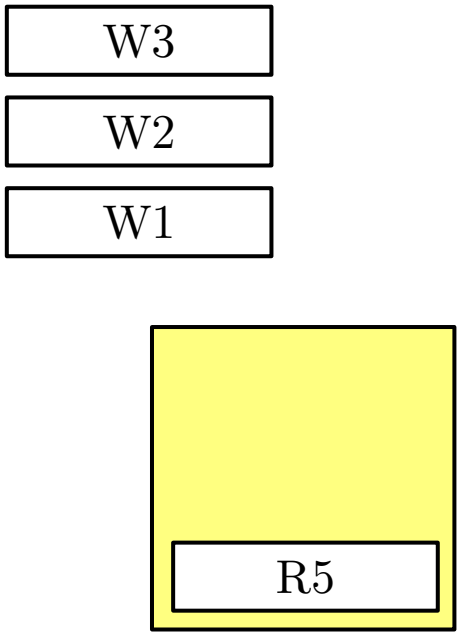

Room for reading/writing

Lights ON = Room in Use

Semaphore S_{rw}=0

- The last reader in the room will be responsible for turning the light OFF.
- **lights OFF** $\rightarrow$ S_{rw}=0 $\rightarrow$ **Free Room**

Solution 4: Reads-Writers: Lightswitch Design Pattern

Threads waiting for room access

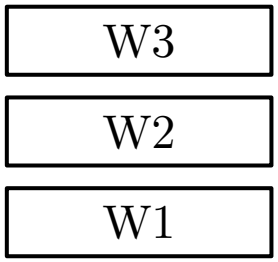

Room for reading/writing

Lights OFF = Free Room

Semaphore $S_{rw}=1$

- The last reader in the room will be responsible for turning the light OFF.
- **lights OFF** $\rightarrow$ $S_{rw}=0$ $\rightarrow$ **Free Room**

Solution 4: Reads-Writers: Lightswitch Design Pattern

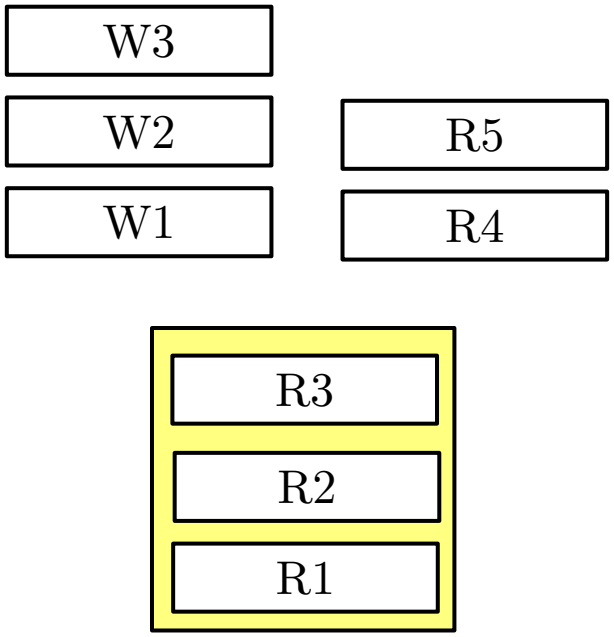

- **ISSUE:** If there is a continuous stream of readers while the room is occupied by at least one reader, the writers might '*starve*'.

Solution 4: Reads-Writers: Turnstile Design Pattern

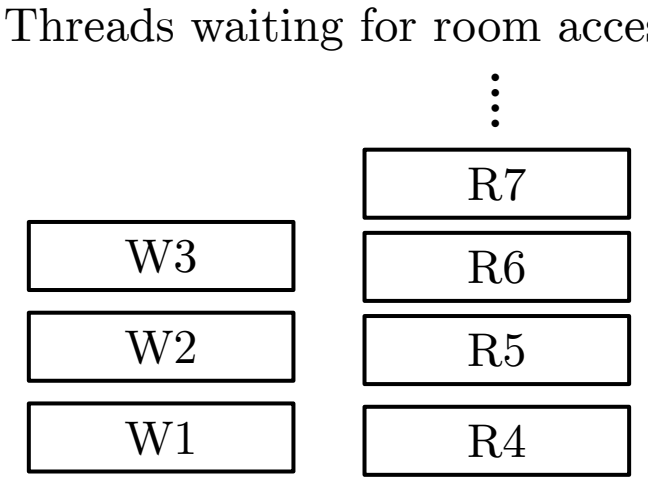

- **ISSUE:** If there is a continuous stream of readers while the room is occupied by at least one reader, the writers might '*starve*'.

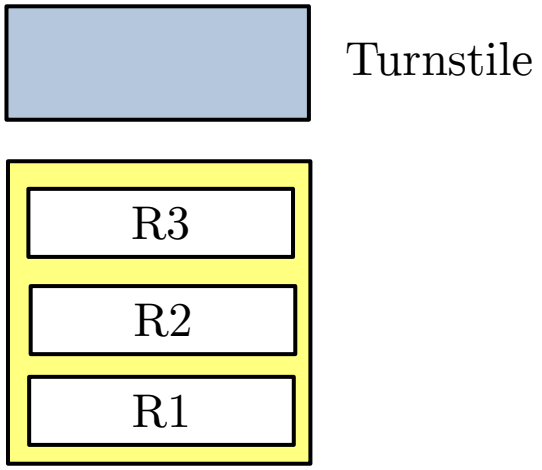

Room for reading/writing

Lights ON = Room in Use

Semaphore S_{rw}=0

Turnstiles

Solution 4: Reads-Writers: Turnstile Design Pattern

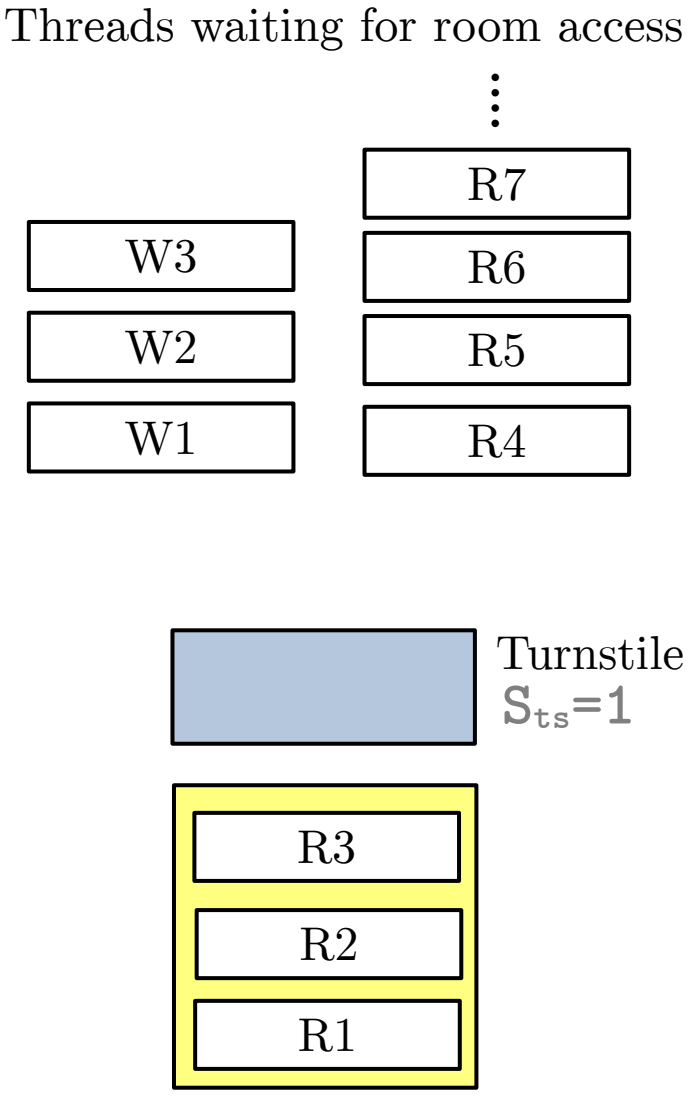

- **ISSUE:** If there is a continuous stream of readers while the room is occupied by at least one reader, the writers might '*starve*'.
- All threads should try to access the turnstile first before trying to enter the room.
- Only one thread at a time can access the turnstile.
- S_{ts} is the semaphore that hold that one *permit* to access the turnstile. Initially, S_{ts}=1.

Solution 4: Reads-Writers: Turnstile Design Pattern

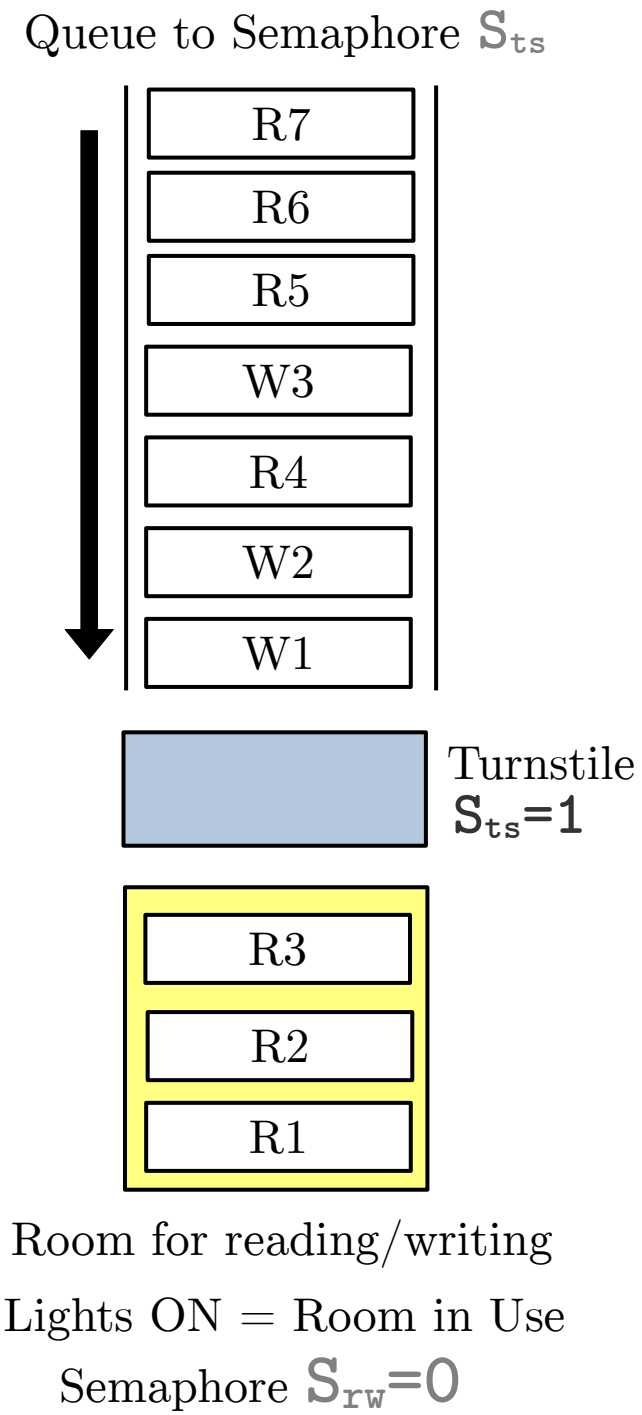

- The turnstile mechanism forces all threads to form a line because all of the thread need to acquire the permit to access the turnstile which means all threads will call `wait(S_ts)`.
- Without the turnstile, if the room is in use by a reader (lights ON), other readers do form a queue in order to access the room and call `read()`. A reader can immediate use to room and call `read()` concurrently with other readers.

Solution 4: Reads-Writers: Turnstile Design Pattern

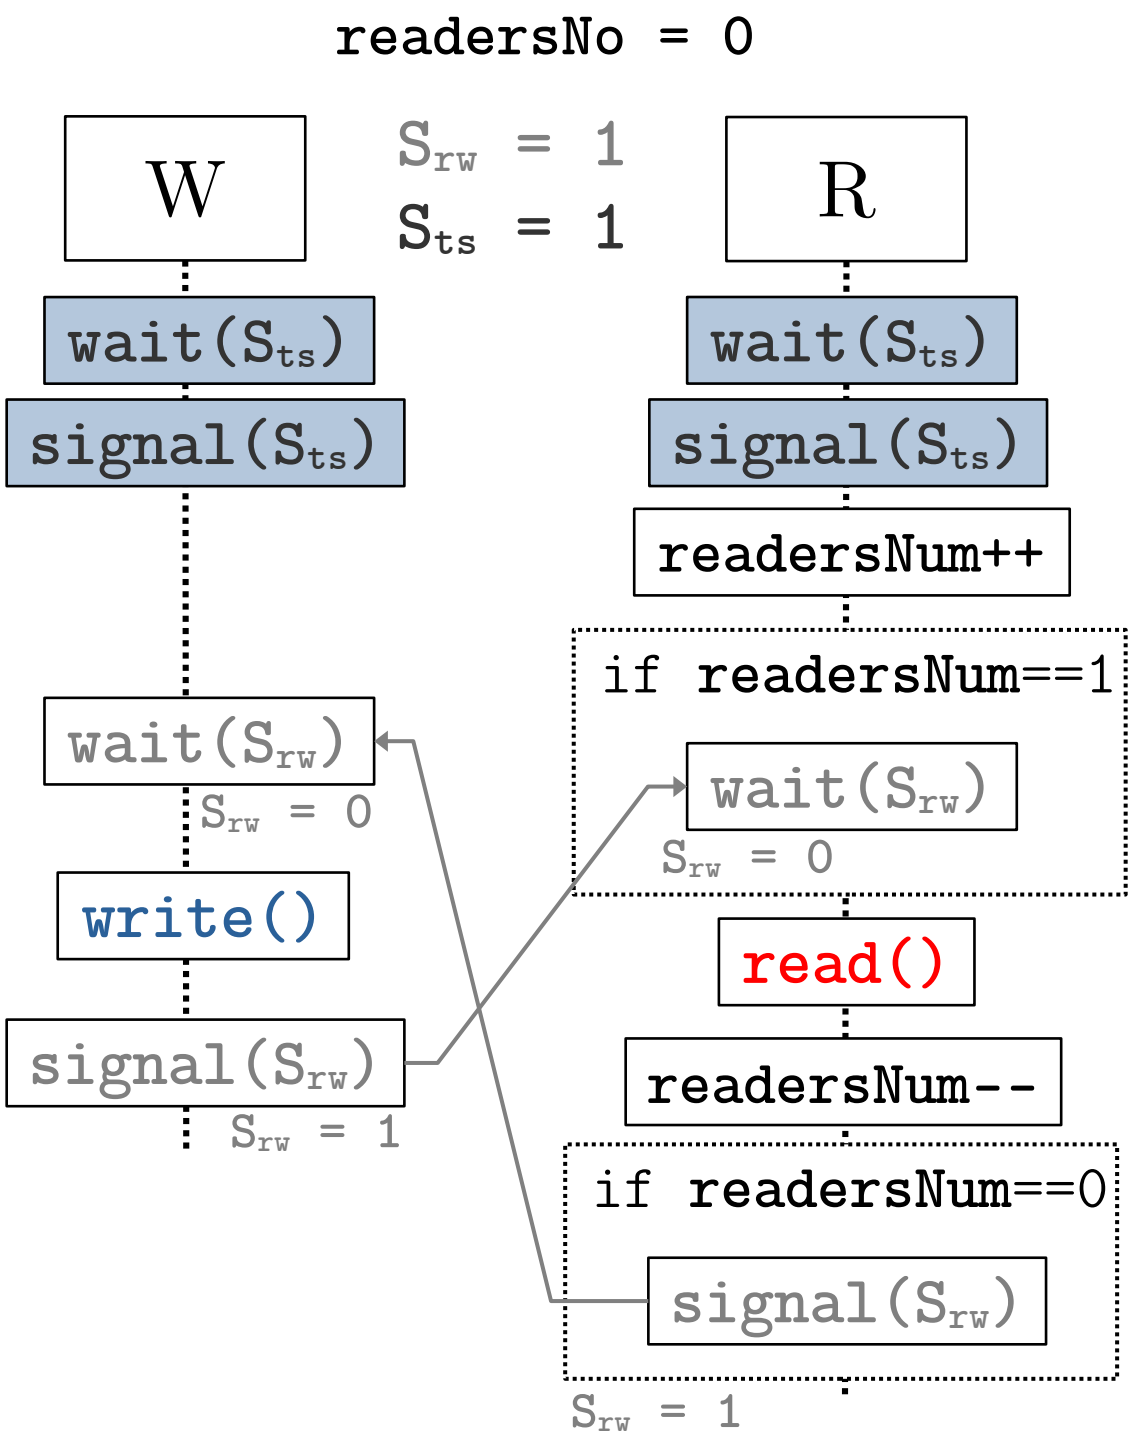

All threads now needs to go through the turnstile by calling `wait(`S_{ts}`)` in order to ask for turnstile access permit then `signal(`S_{ts}`)` is called afterwards.

Solution 4: Reads-Writers: Turnstile Design Pattern

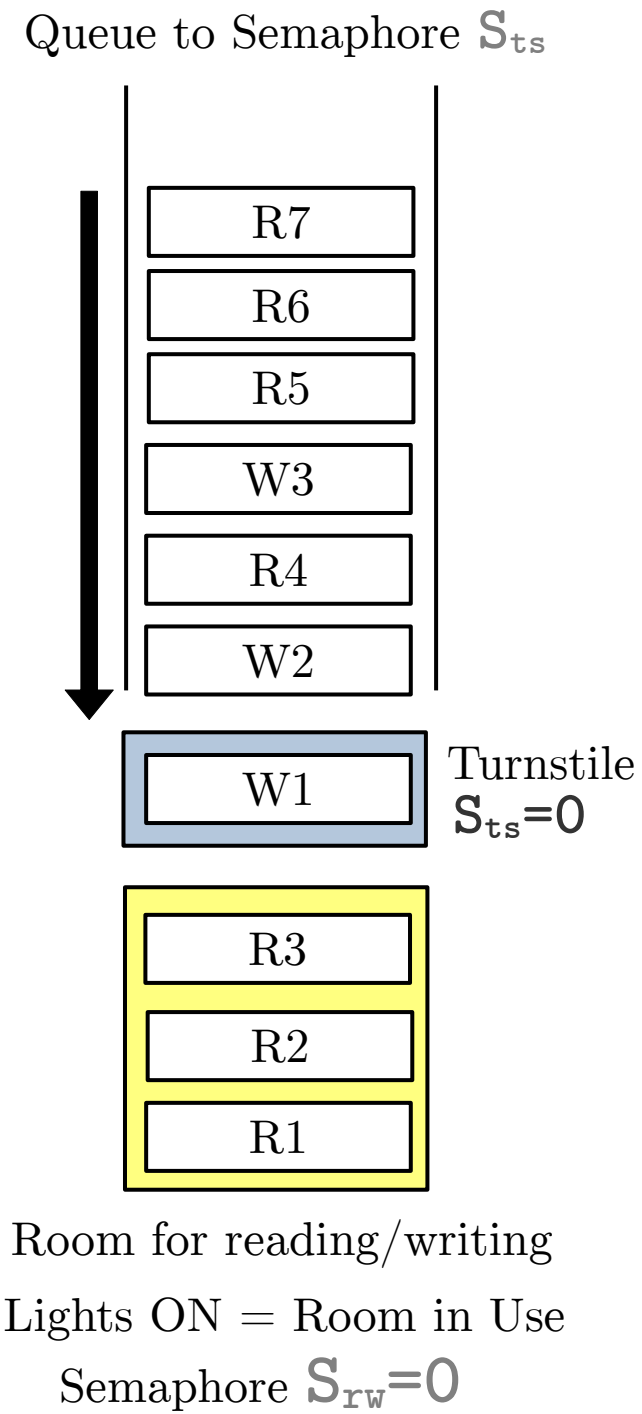

- **ISSUE:** If there is a continuous stream of readers while the room is occupied by at least one reader, the writers might '*starve*'.
- All threads should try to access the turnstile first before trying to enter the room.
- Only one thread at a time can access the turnstile.
- S_{ts} is the semaphore that hold that one *permit* to access the turnstile. Initially, S_{ts}=1.

Solution 4: Reads-Writers: Turnstile Design Pattern

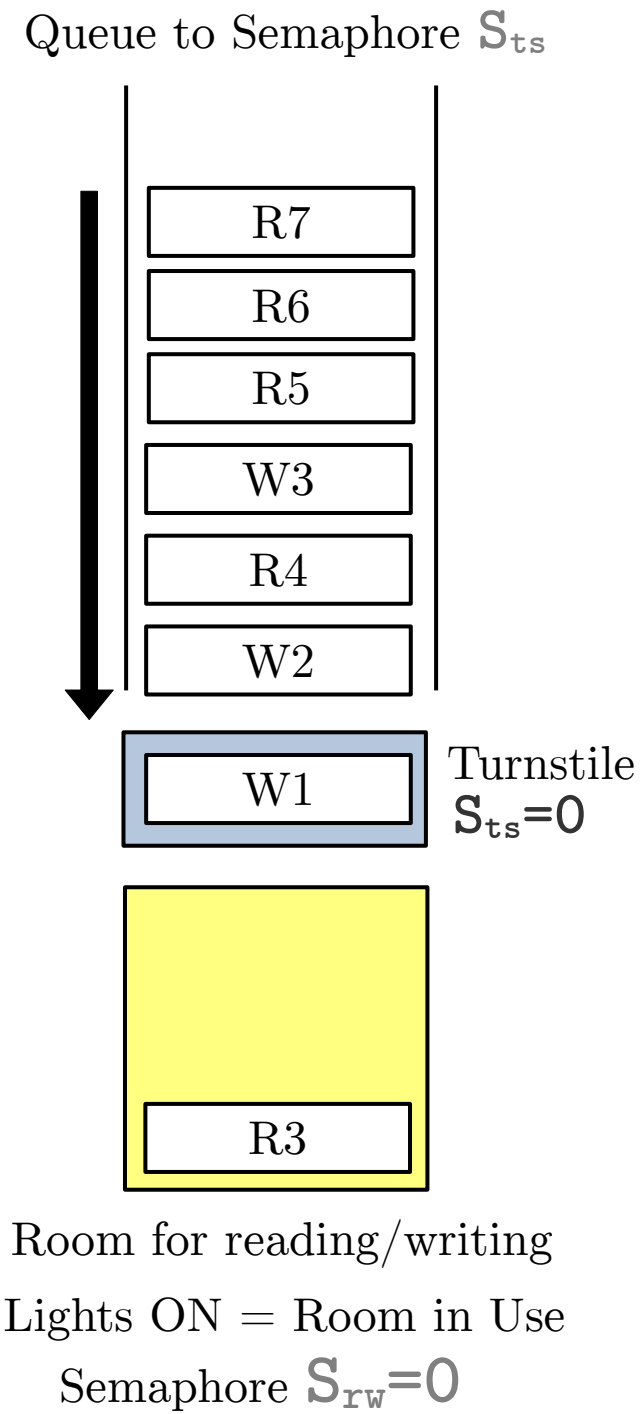

- **ISSUE:** If there is a continuous stream of readers while the room is occupied by at least one reader, the writers might '*starve*'.
- All threads should try to access the turnstile first before trying to enter the room.
- Only one thread at a time can access the turnstile.
- S_{ts} is the semaphore that hold that one *permit* to access the turnstile. Initially, S_{ts}=1.

Solution 4: Reads-Writers: Turnstile Design Pattern

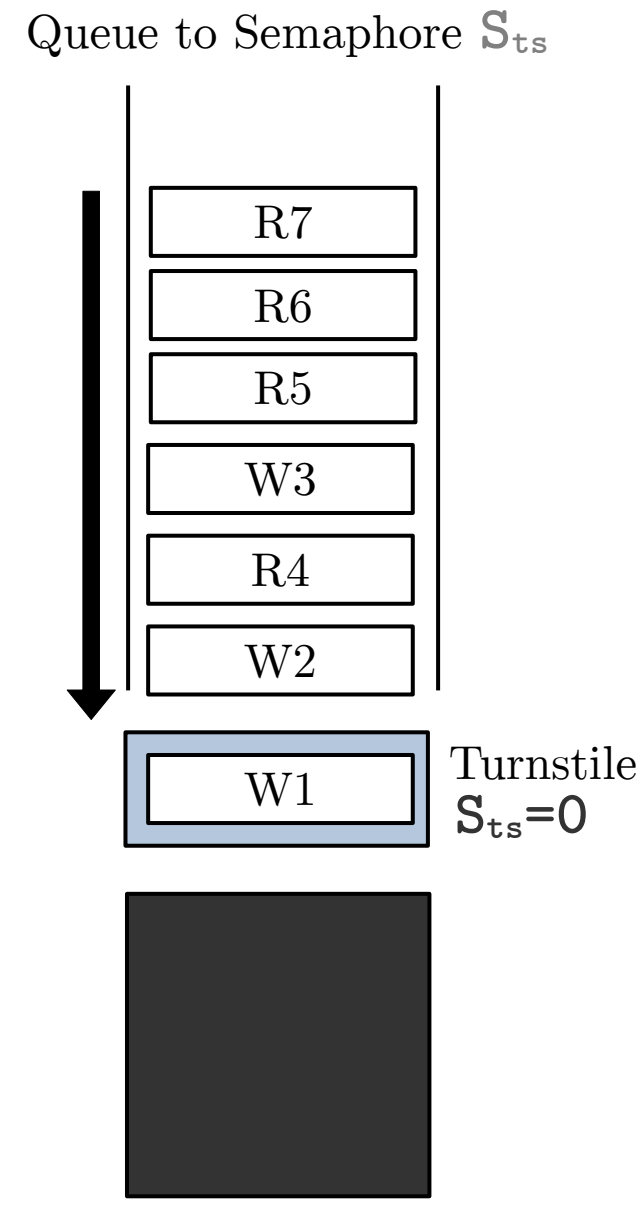

- **ISSUE:** If there is a continuous stream of readers while the room is occupied by at least one reader, the writers might '*starve*'.
- All threads should try to access the turnstile first before trying to enter the room.
- Only one thread at a time can access the turnstile.
- S_{ts} is the semaphore that hold that one *permit* to access the turnstile. Initially, S_{ts}=1.

Synchronization Task/Problem 5: **Dining Philosophers**

[Problem & Solution Descriptions]

Task/Problem 5: Dining Philosophers

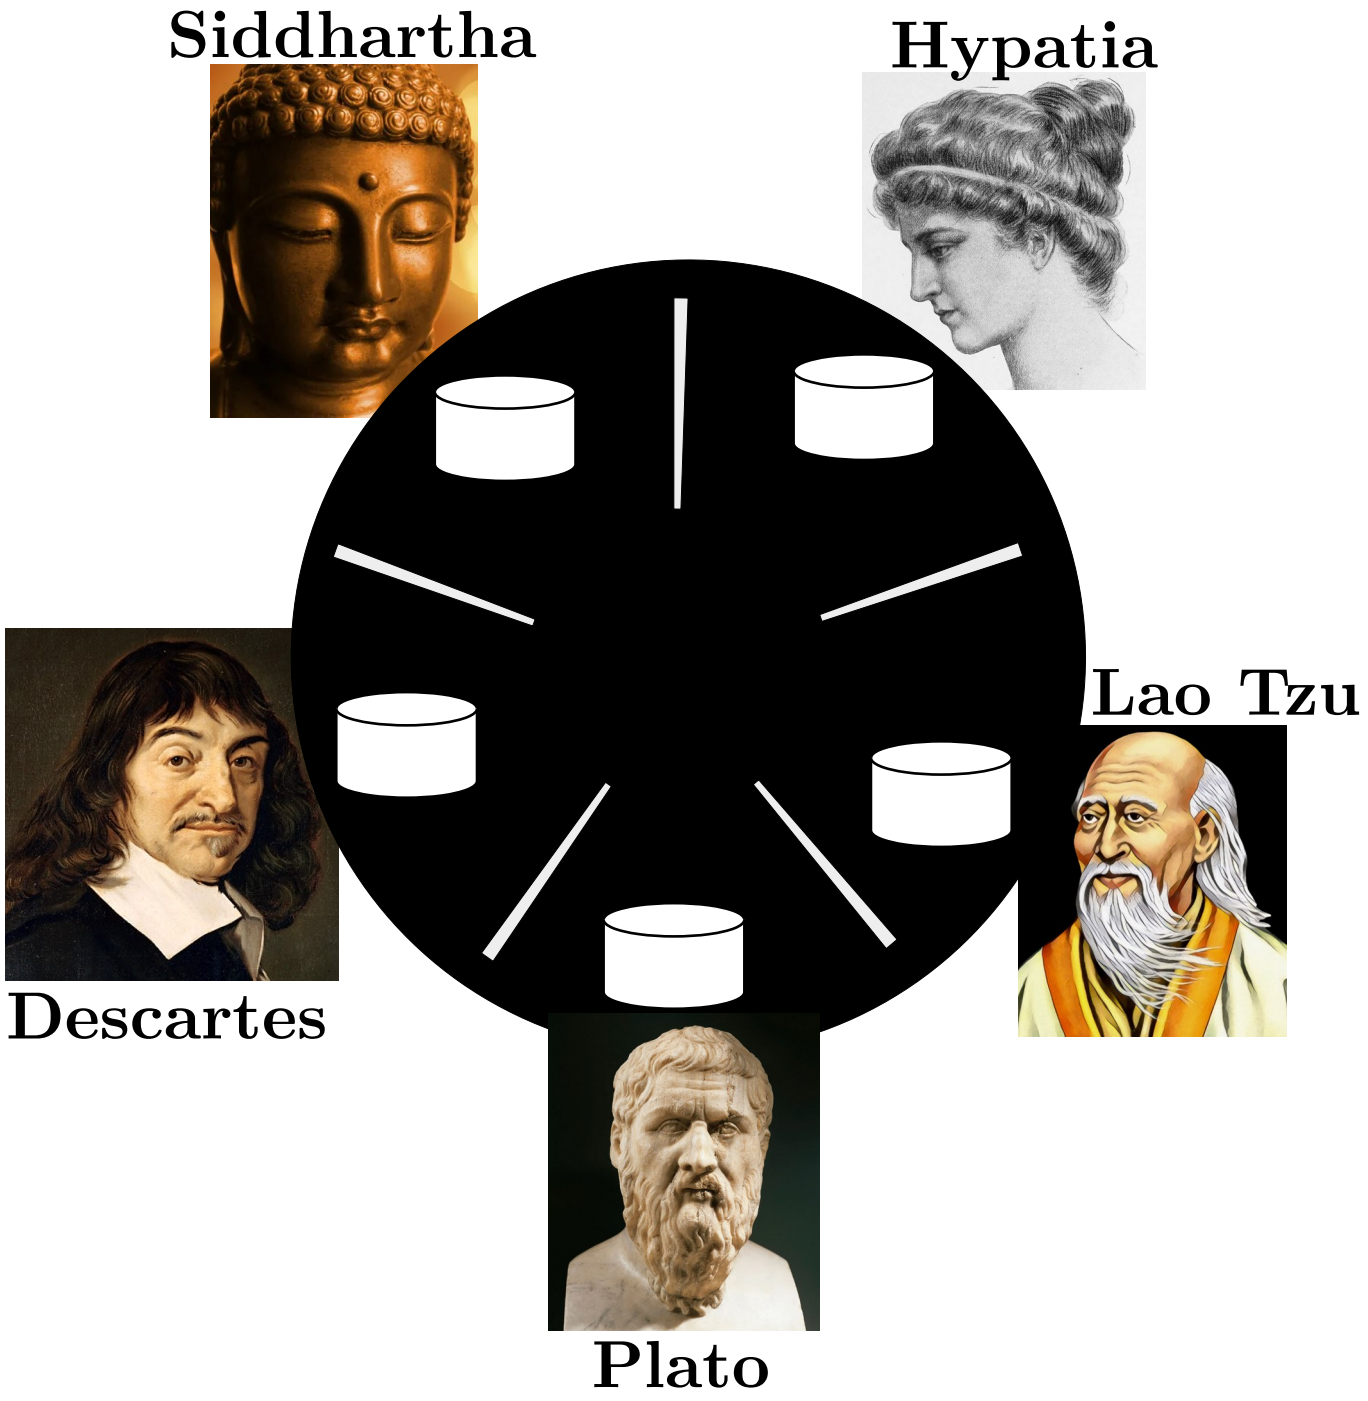

Task/Problem 5: Dining Philosophers

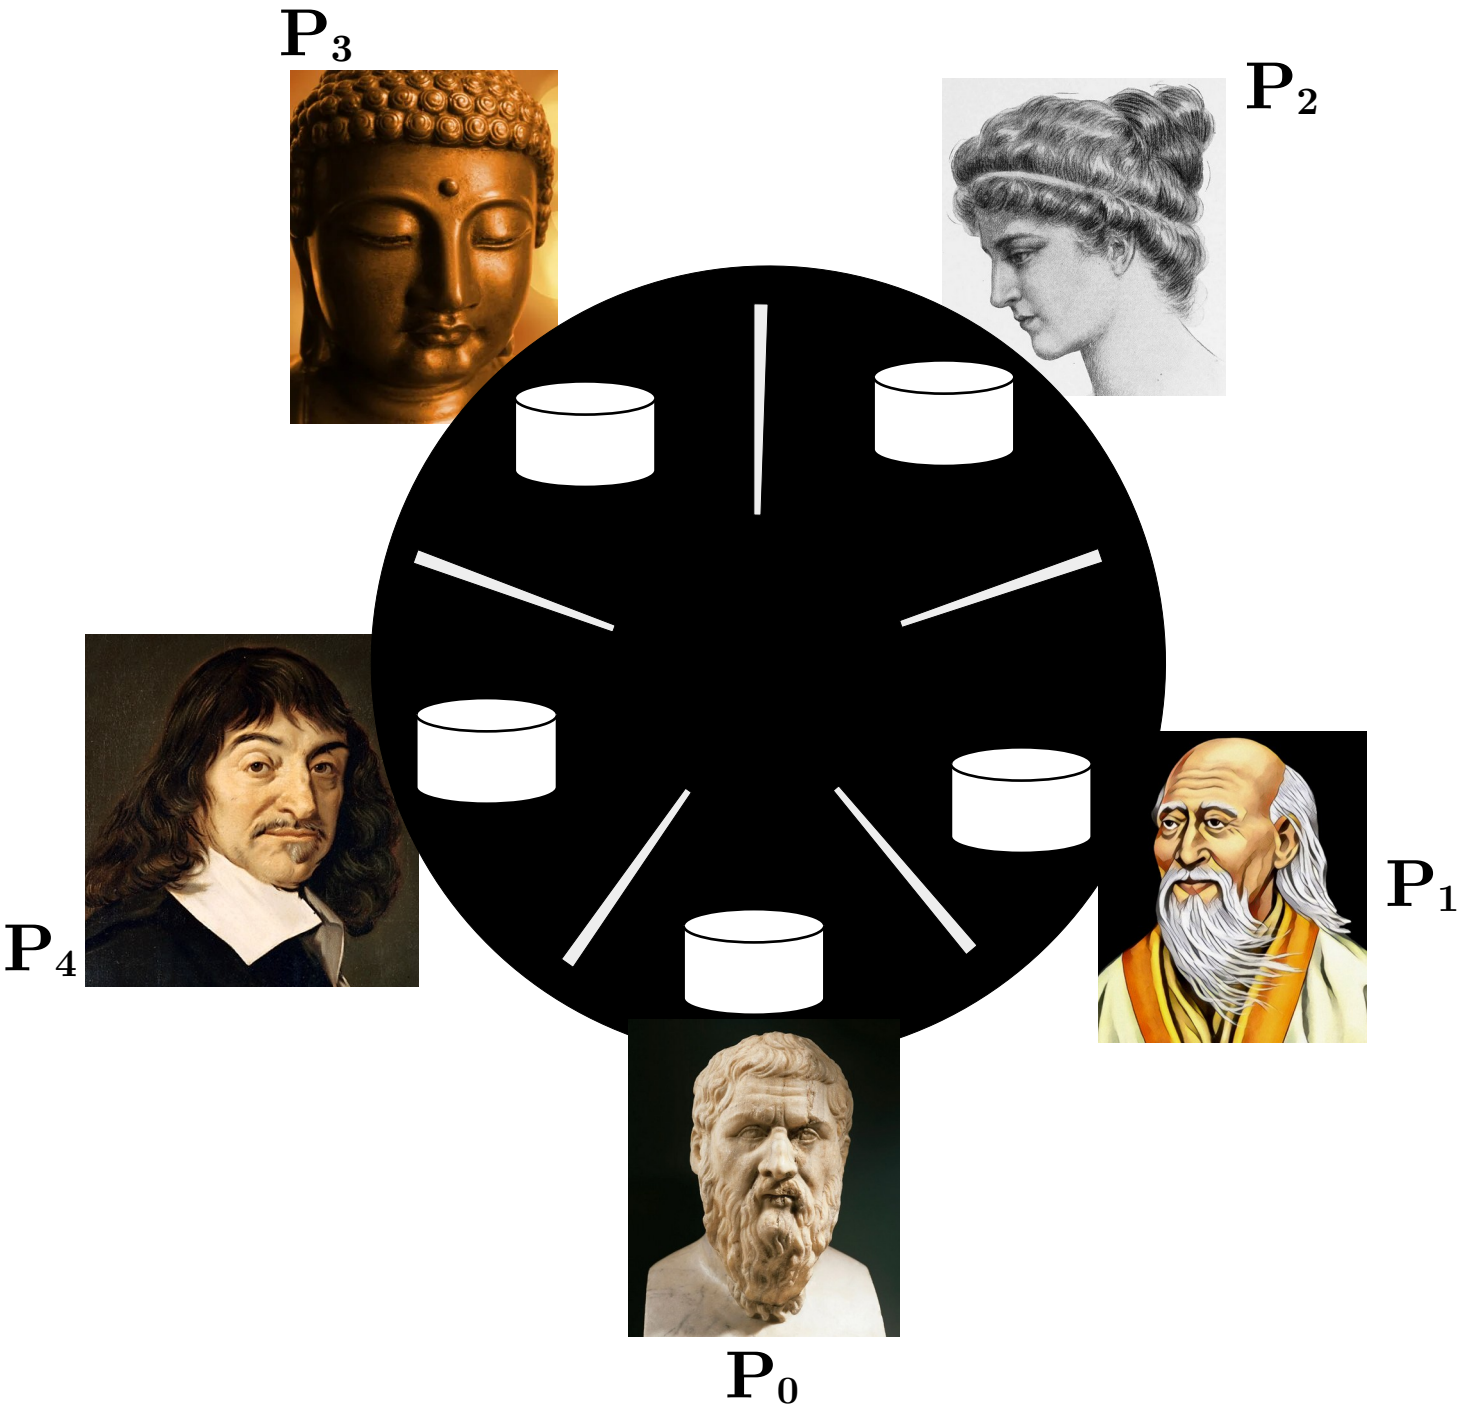

Task/Problem 5: Dining Philosophers

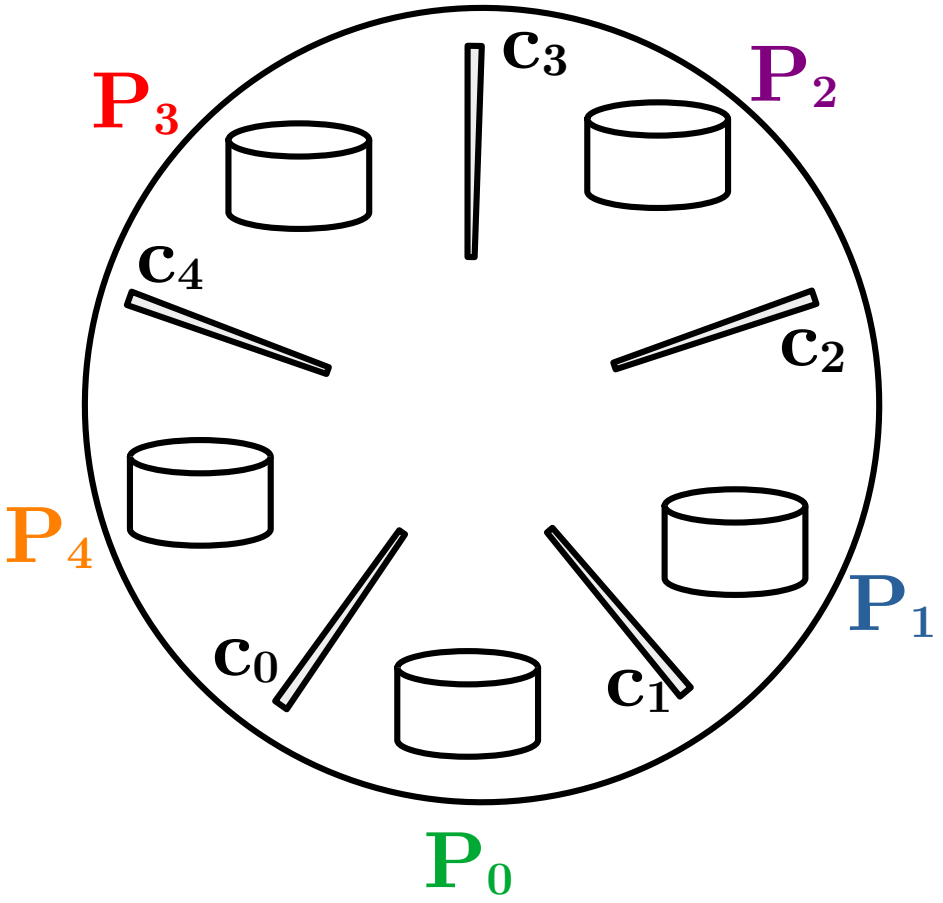

Task/Problem 5: Dining Philosophers

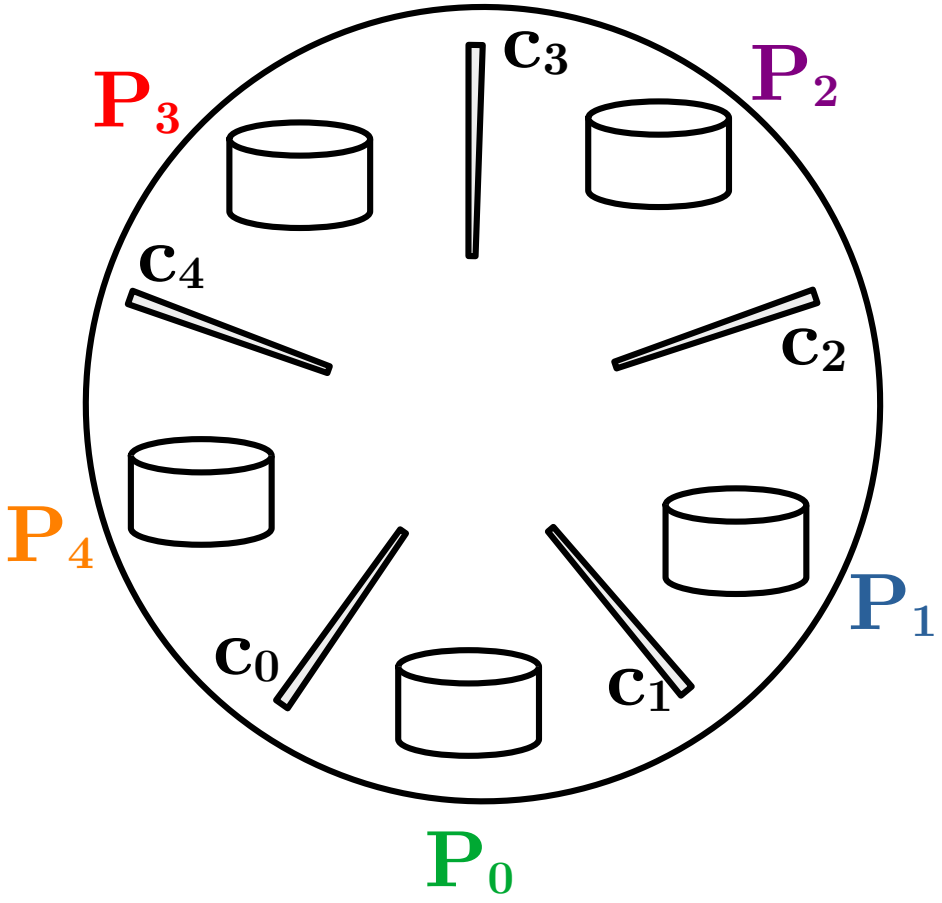

P_i needs c_i and c_{i+1} or more accurately $c_{(i+1)\%5}$

Constraint: Each philosopher needs two chopsticks (left & right) to eat.

P_0 needs c_0 and c_1,

P_1 needs c_1 and c_2,

P_2 needs c_2 and c_3,

P_3 needs c_3 and c_4,

P_4 needs c_4 and c_0,

Task/Problem 5: Dining Philosophers

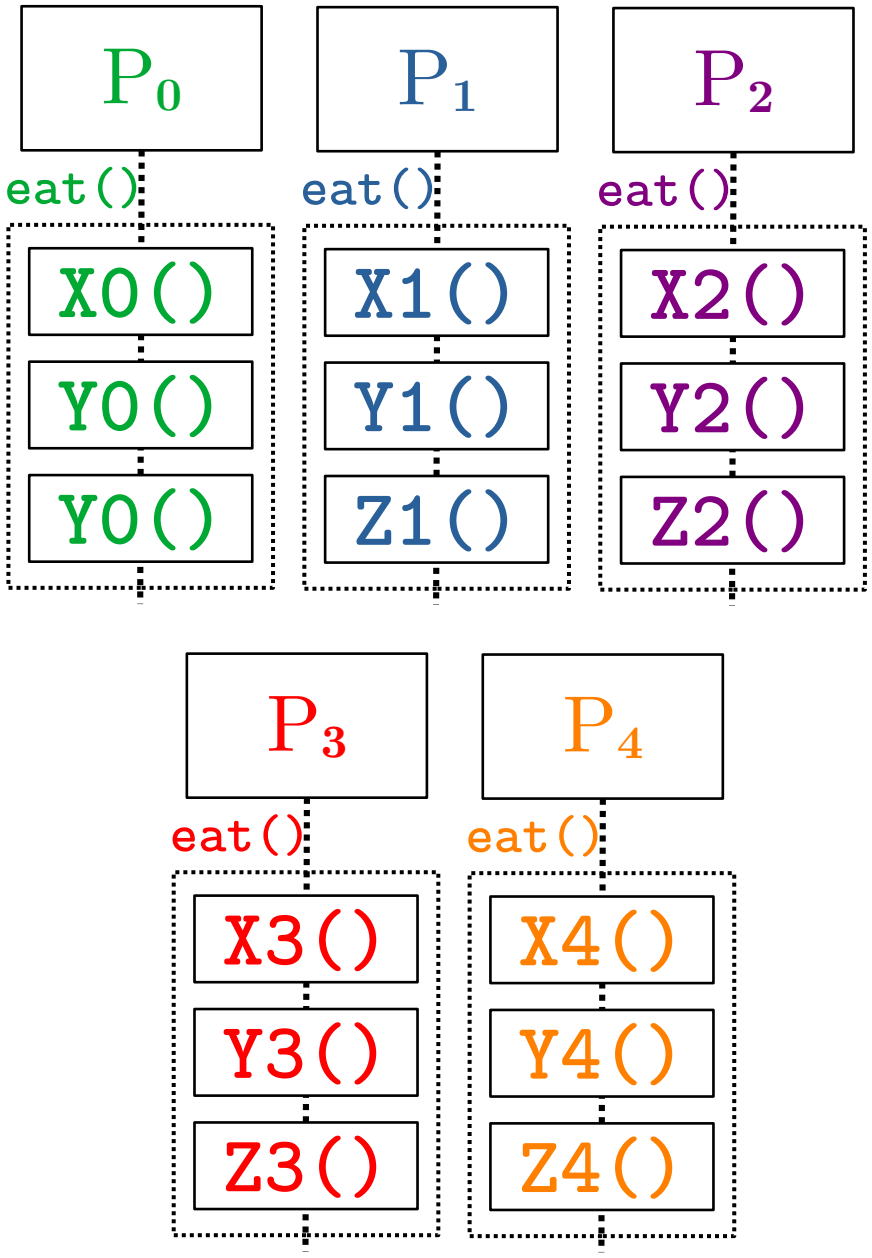

some VALID orders

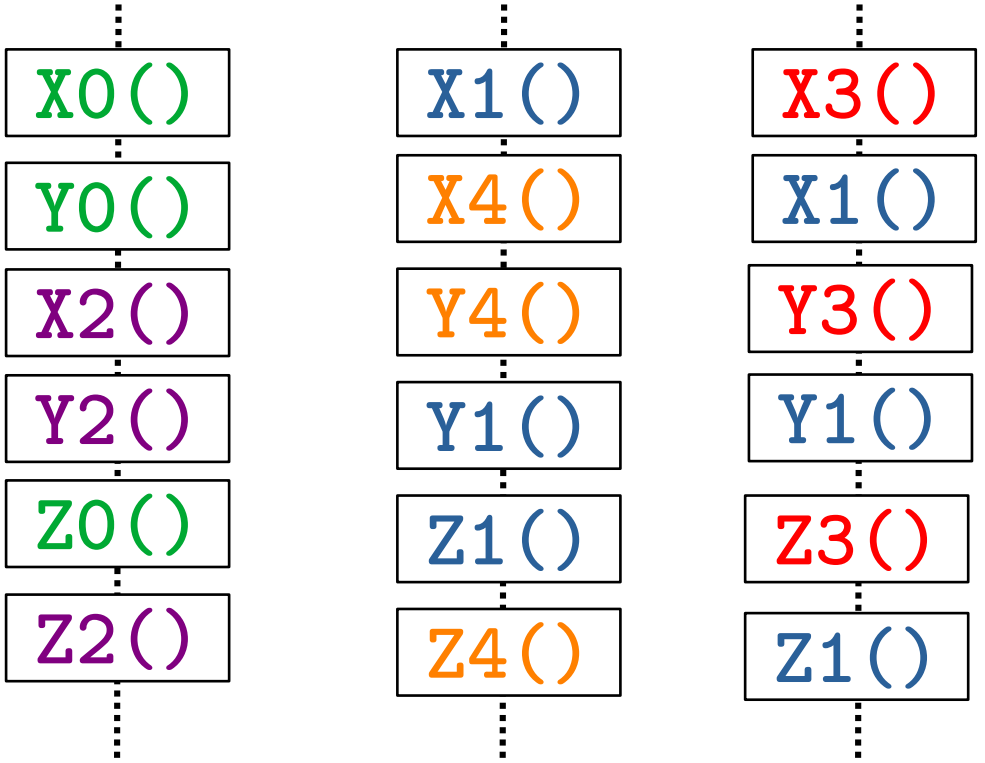

P0 & P2 eating in parallel

P1 & P4 eating in parallel

P1 & P3 eating in parallel

Task/Problem 5: Dining Philosophers

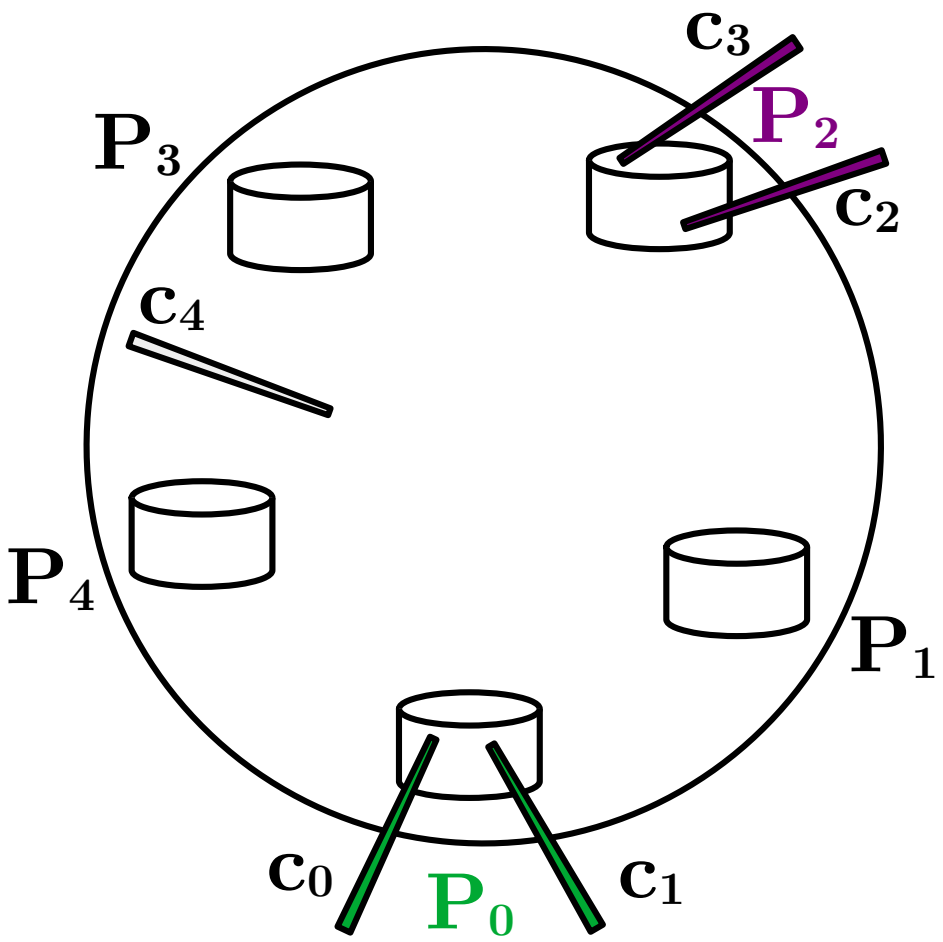

P0 & P2
eating in parallel

some VALID orders

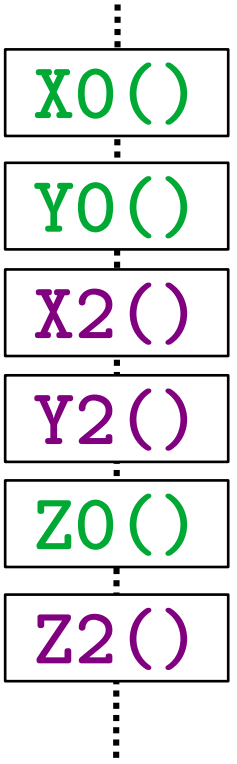

P0 & P2
eating in parallel

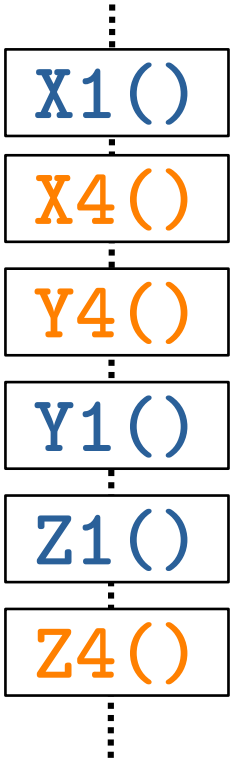

P1 & P4
eating in parallel

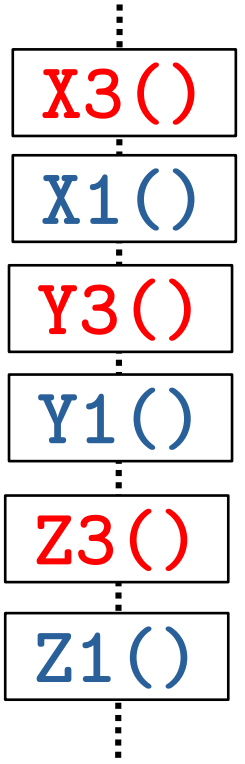

P1 & P3
eating in parallel

Task/Problem 5: Dining Philosophers

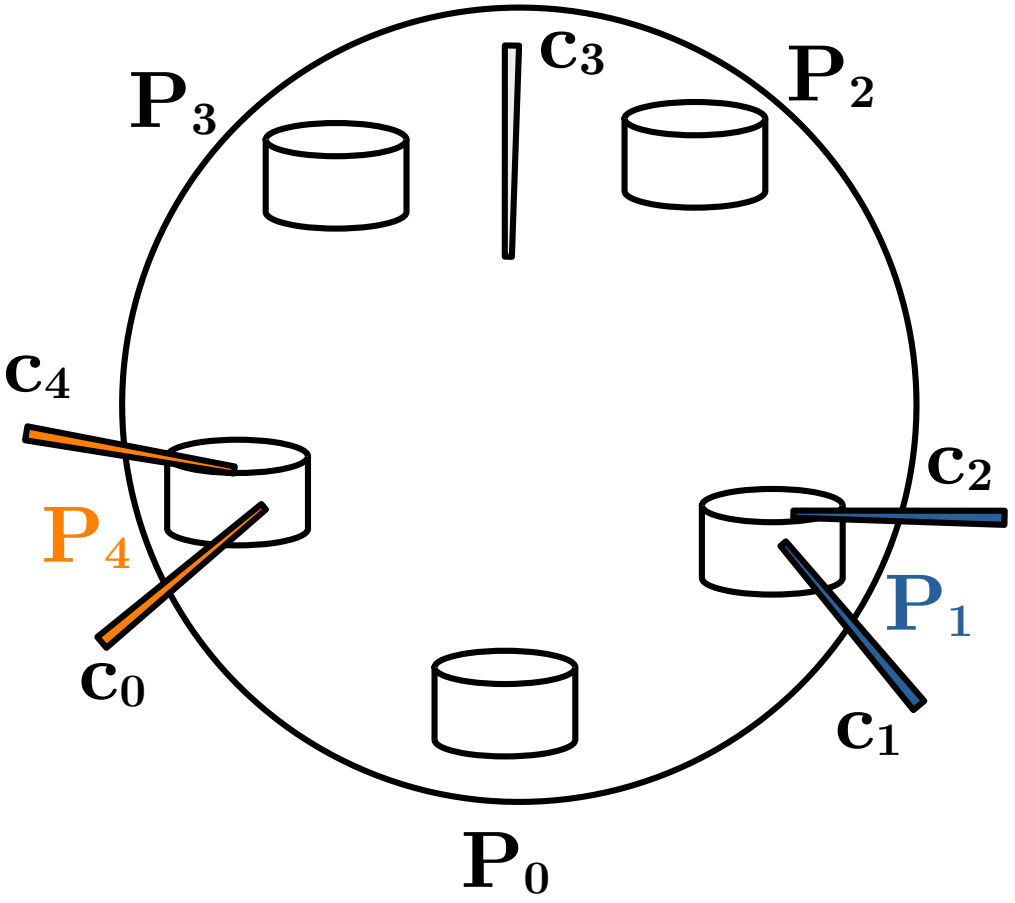

P1 & P4
eating in parallel

some VALID orders

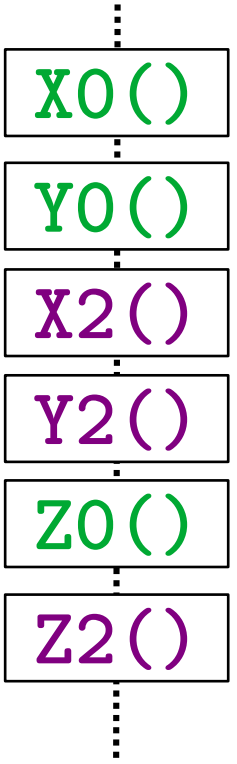

P0 & P2
eating in parallel

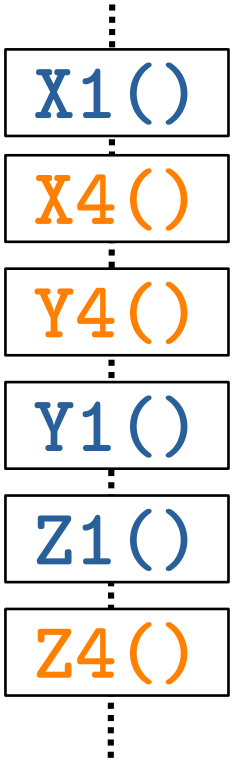

P1 & P4
eating in parallel

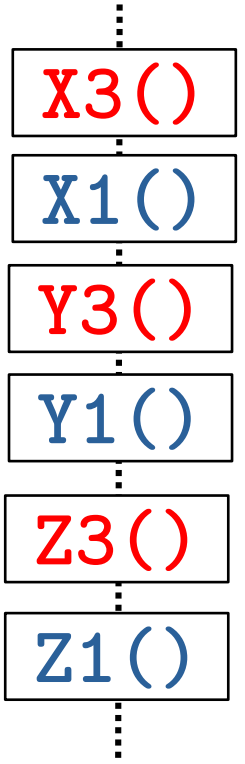

P1 & P3
eating in parallel

Task/Problem 5: Dining Philosophers

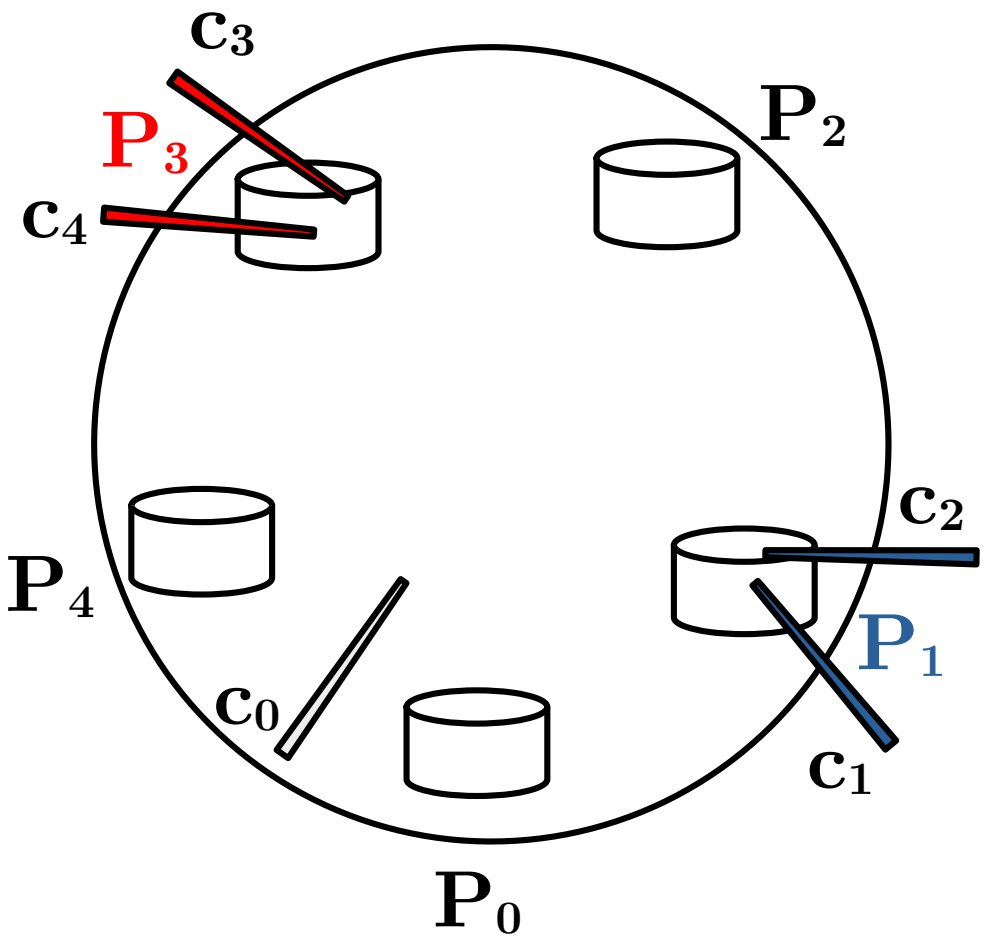

P1 & P3
eating in parallel

some VALID orders

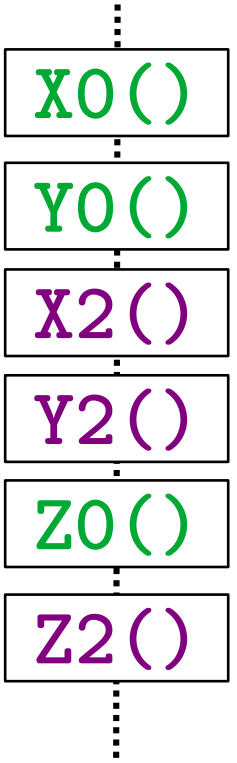

P0 & P2
eating in parallel

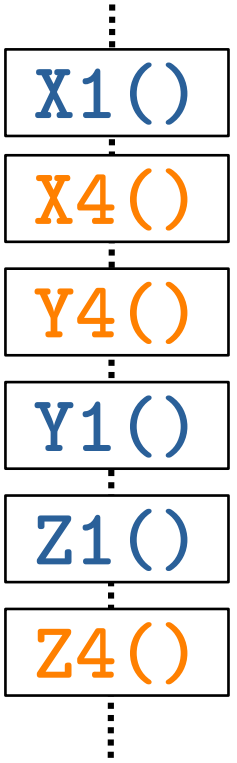

P1 & P4
eating in parallel

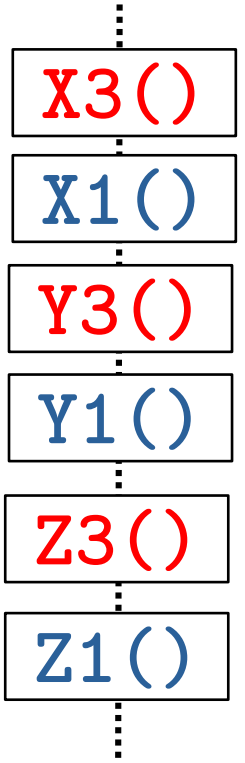

P1 & P3
eating in parallel

Task/Problem 5: Dining Philosophers

some INVALID orders

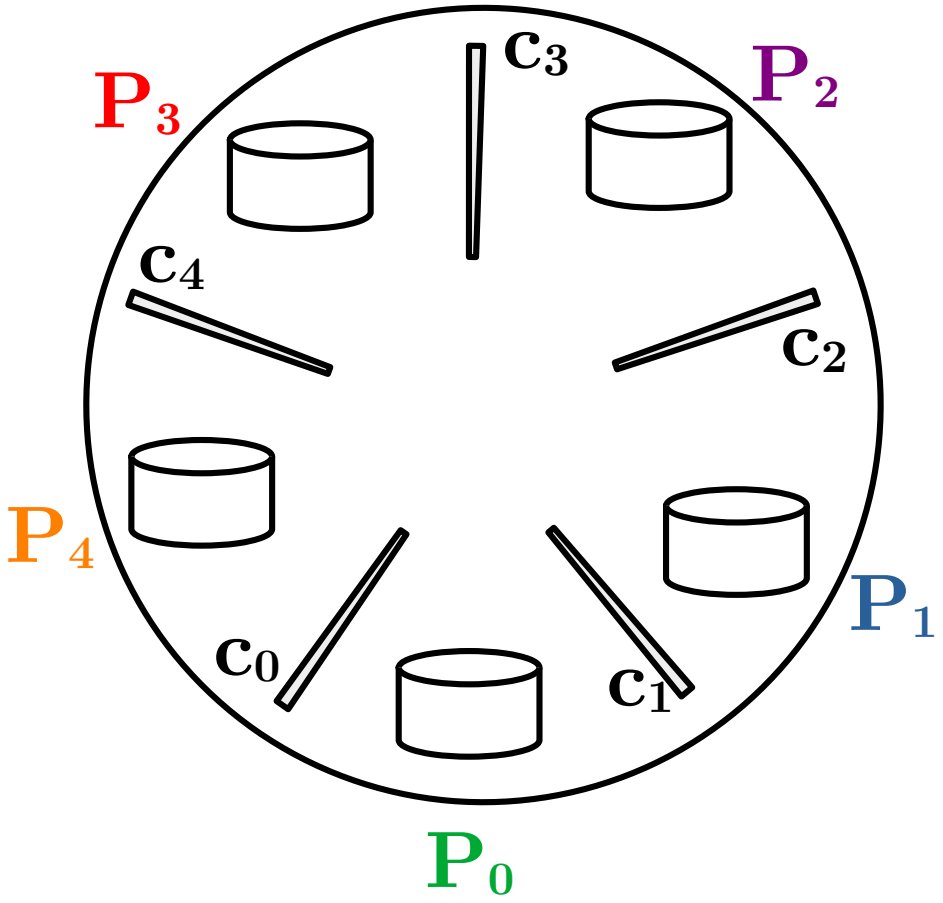

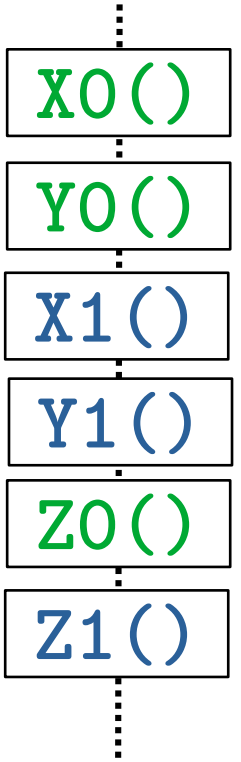

P0 & P1 eating in parallel

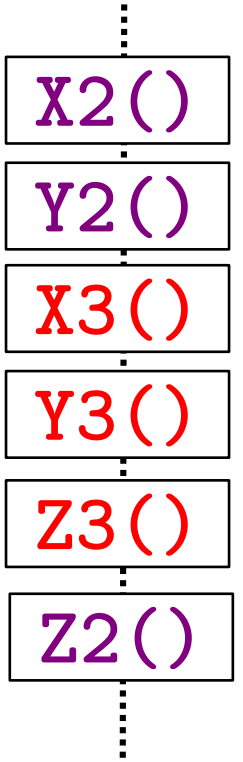

P2 & P3 eating in parallel

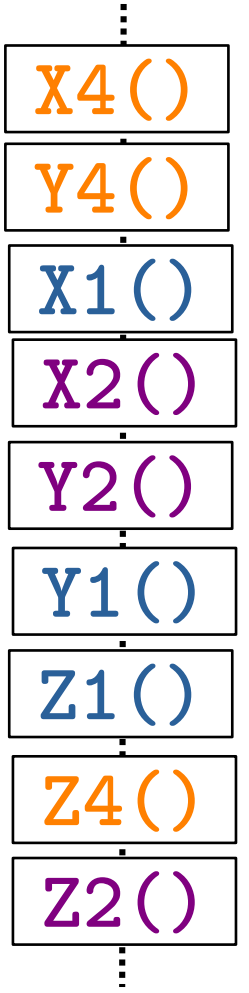

P1, P2, P4 eating in parallel

Task/Problem 5: Dining Philosophers

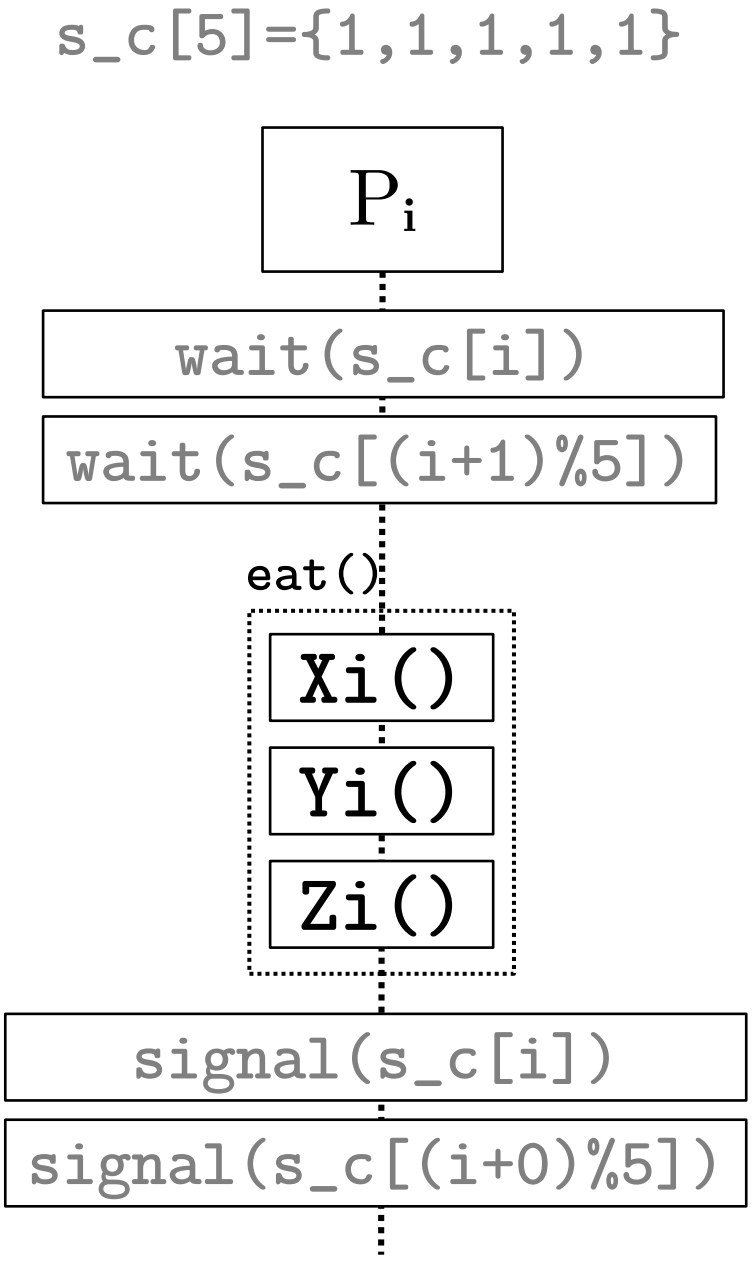

`s_c[5]` is array of semaphores of size 5, one semaphore for each chopstick. Each semaphore `s_c[i]` is initialized to `1`.

Thread $\mathbf{P}_i$ calls `wait()` for the semaphores `s_c[i]and s_c[i+1]` (more accurately `s_c[(i+1)%5]`) associated with $\mathbf{P}_i$'s left and right chopsticks.

After acquire the two chopstick permits, $\mathbf{P}_i$ will eat and then release both permits by calling `signal()` for each semaphore.

Task/Problem 5: Dining Philosophers: Deadlock

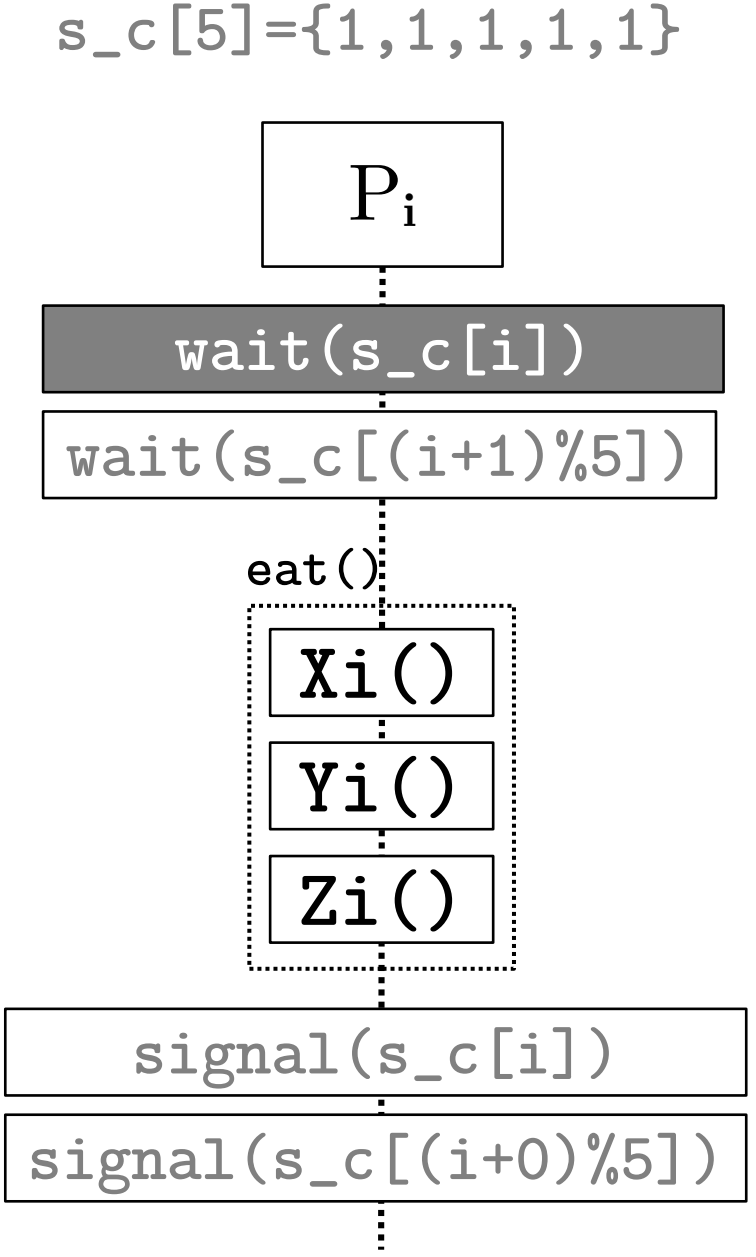

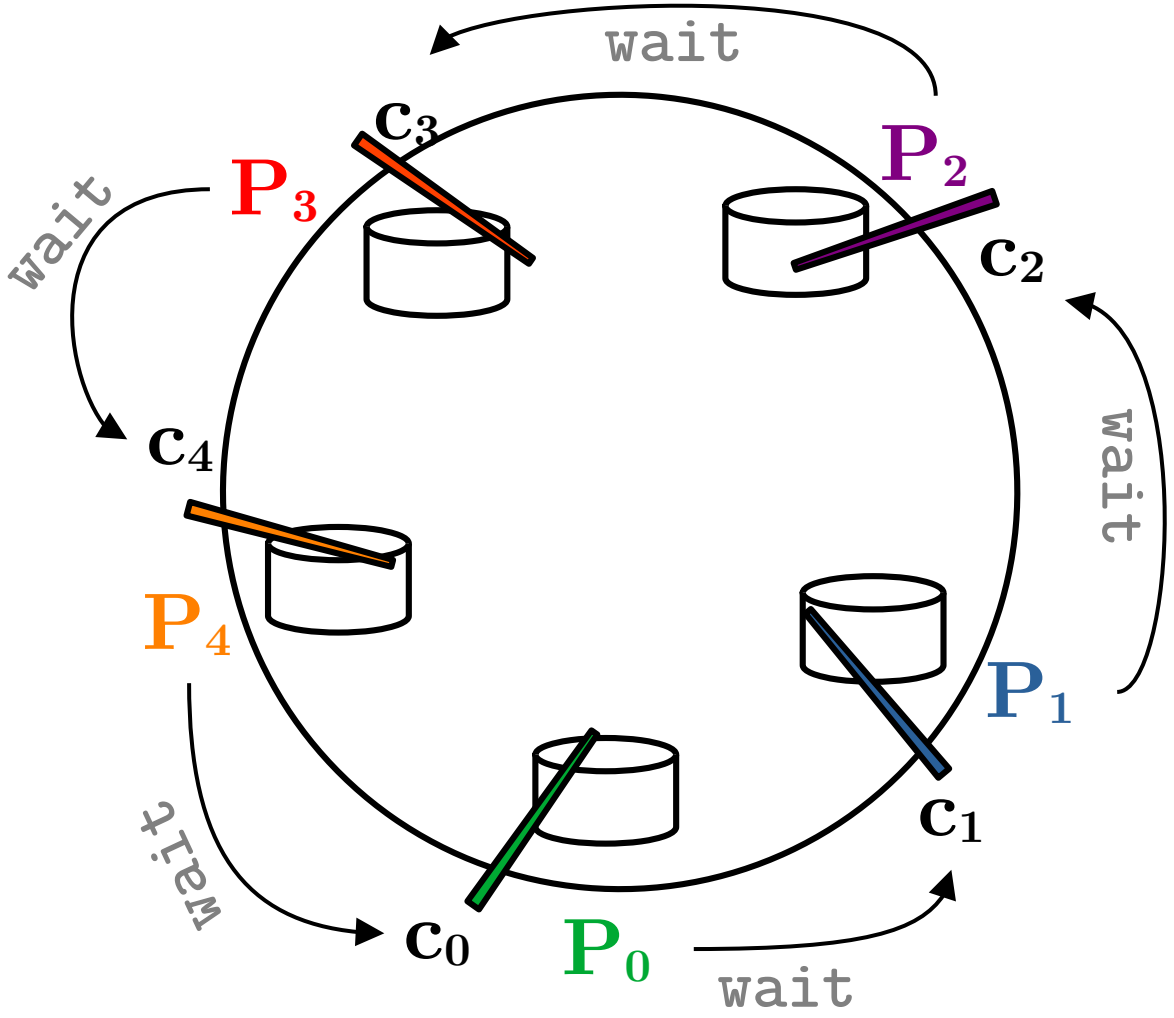

Each $\mathbf{P_i}$ was able to get the permission to use $\mathbf{c_i}$ (via semaphore `s_c[i]`) but each **Pi** is also waiting for permission to use $\mathbf{c_{i+1}}$ which is held by $\mathbf{P_{i+1}}$.

Mutual/Circular waiting = Deadlock.